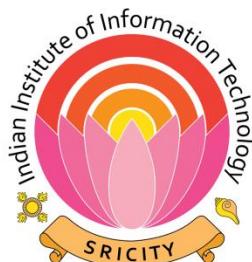


# **Computer Vision**

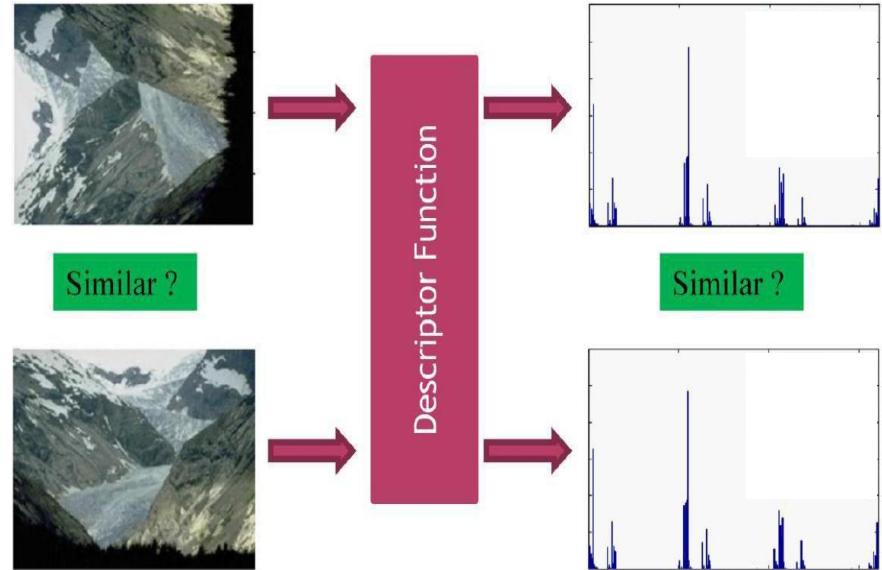
## **Region Detection & Local Descriptors**

**Dr. Mrinmoy Ghorai**

**Indian Institute of Information Technology  
Sri City, Chittoor**

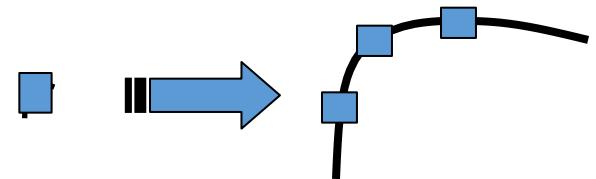


# Region Detection & Local Descriptors



# Previous Class

- Keypoint detection: repeatable and distinctive
  - Corners, Harris
  - Invariant to scale, rotation, etc.
- Harris Corner Detection
  - Rotation Invariant
  - Partial Intensity Change Invariant
  - *Not Invariant to Scale*

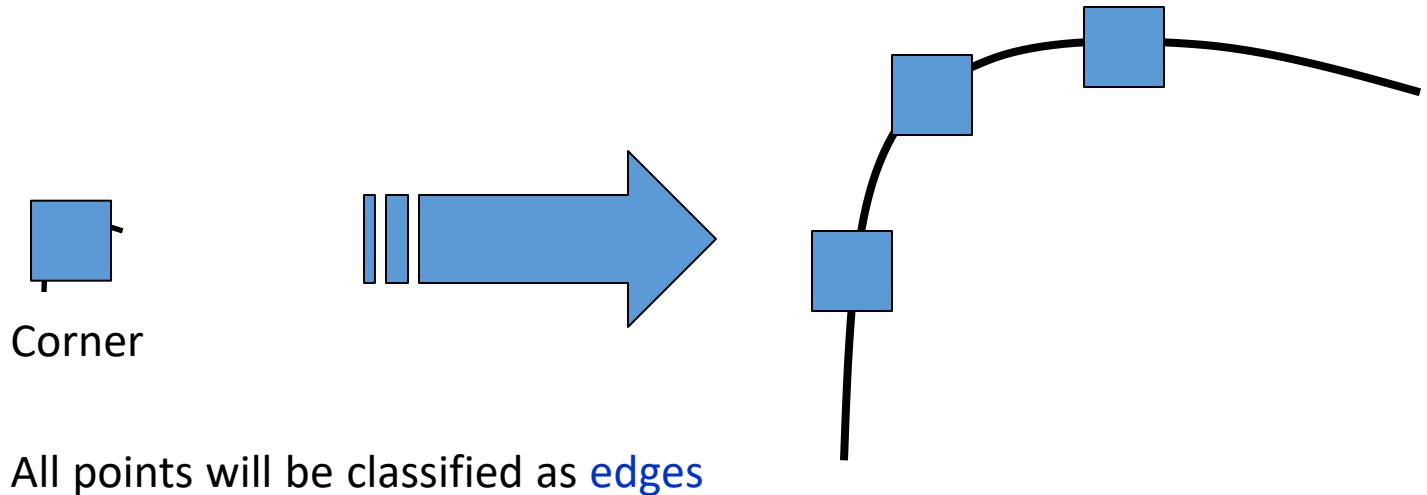


# Today's class

- Scale Invariance
- Region Detection
- Local Descriptors
- Image Matching

# Harris Detector: Invariance Properties

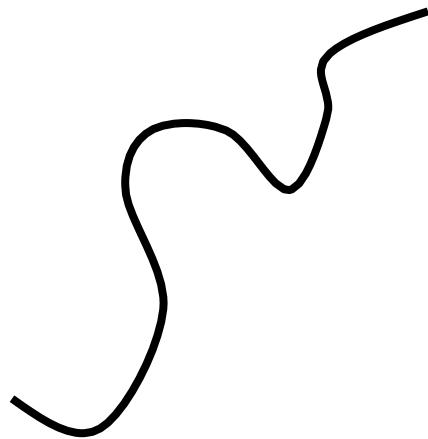
- Scaling



*Not invariant to scaling*

# Scale invariant detection

Suppose you're looking for corners

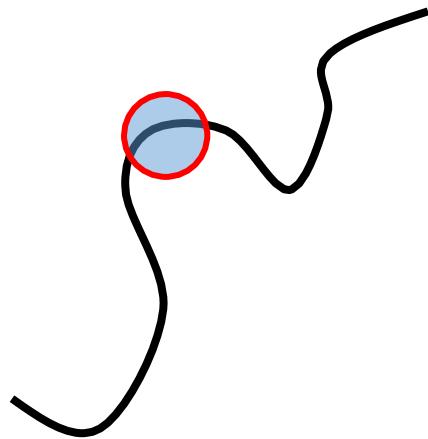


Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Scale invariant detection

Suppose you're looking for corners

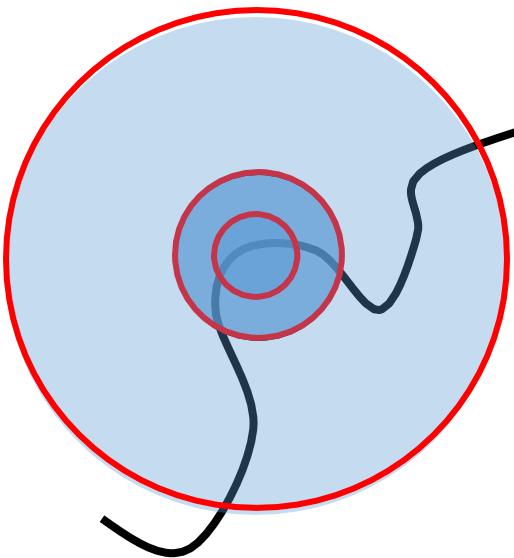


Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Scale invariant detection

Suppose you're looking for corners



Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Automatic Scale Selection

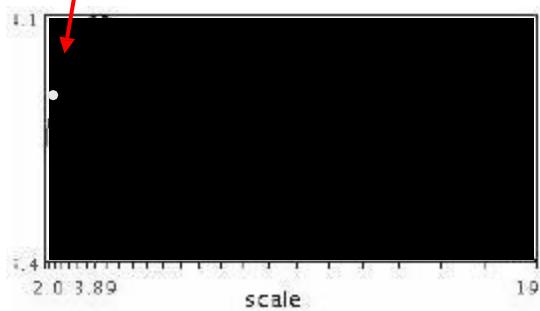


$$f(I_{i_1 \dots i_m}(x, \sigma)) = f(I_{i_1 \dots i_m}(x', \sigma'))$$

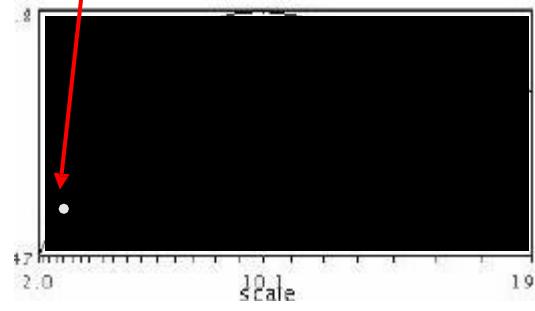
How to find corresponding patch sizes?

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



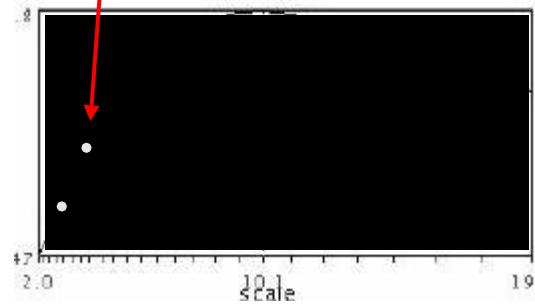
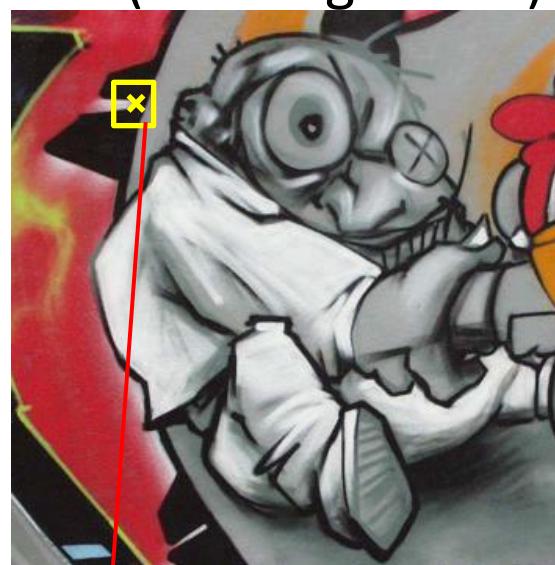
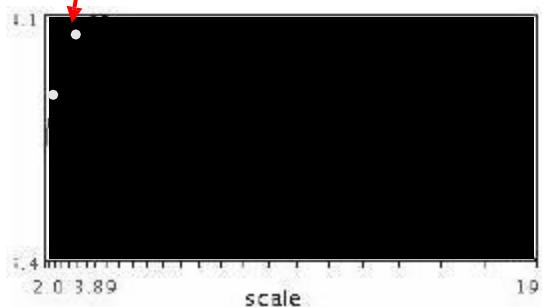
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

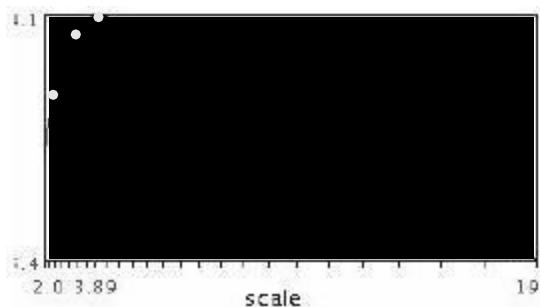
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)

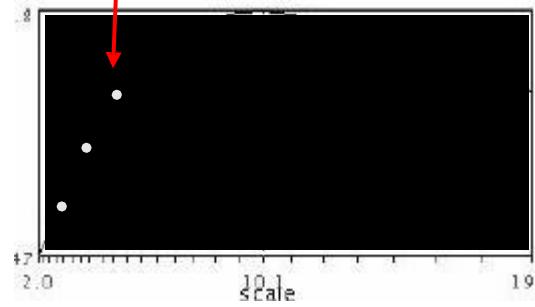


# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



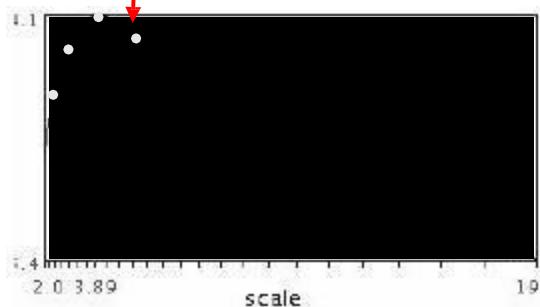
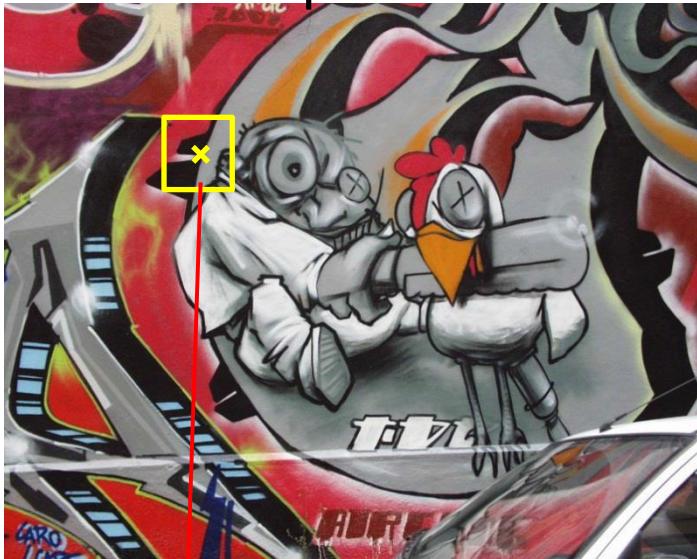
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



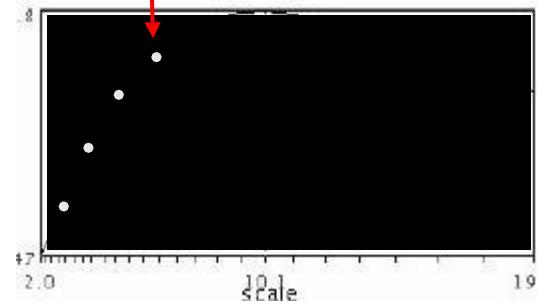
$$f(I_{i_1 \dots i_m}(x', \sigma))$$

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



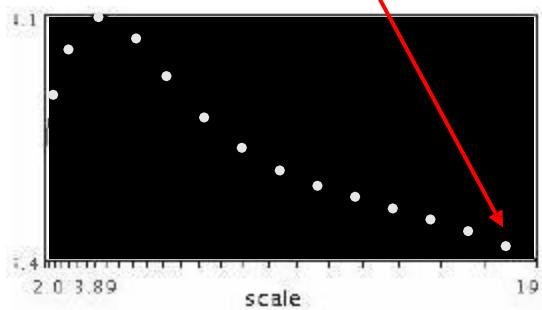
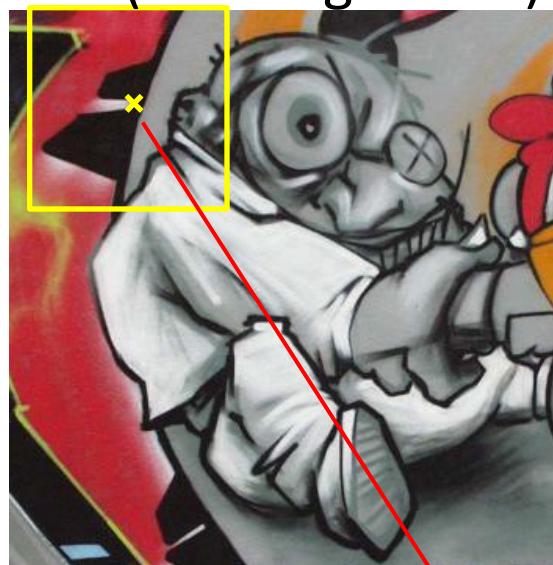
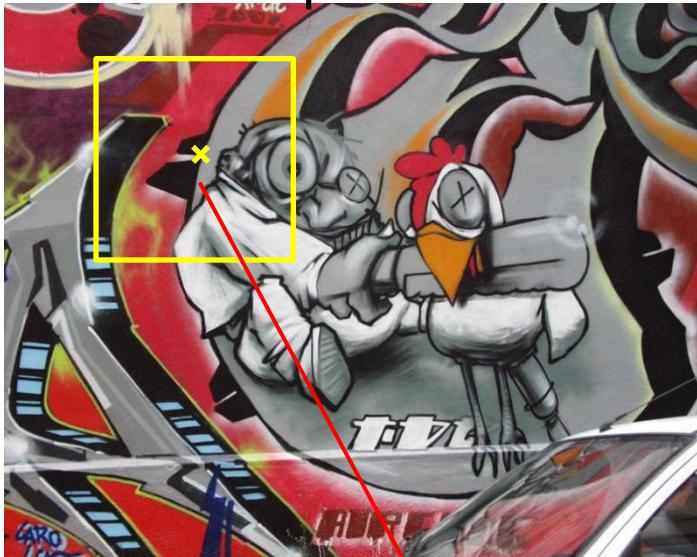
$$f(I_{i_1...i_m}(x, \sigma))$$



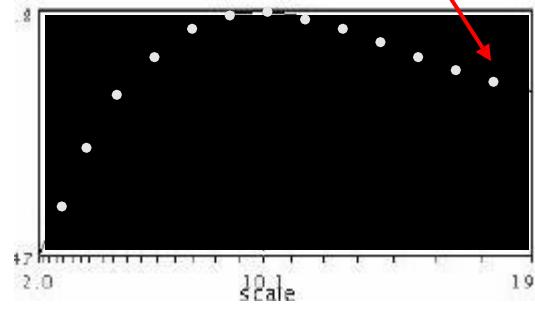
$$f(I_{i_1...i_m}(x', \sigma))$$

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



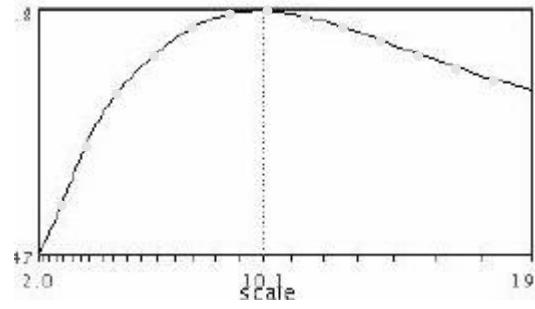
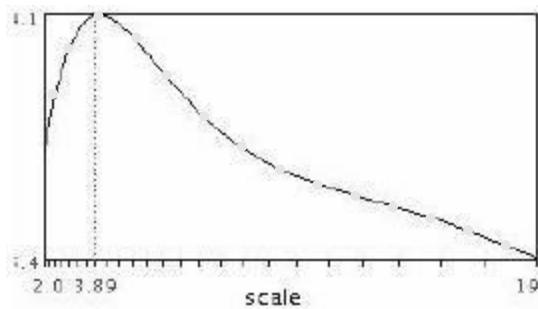
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



# Implementation

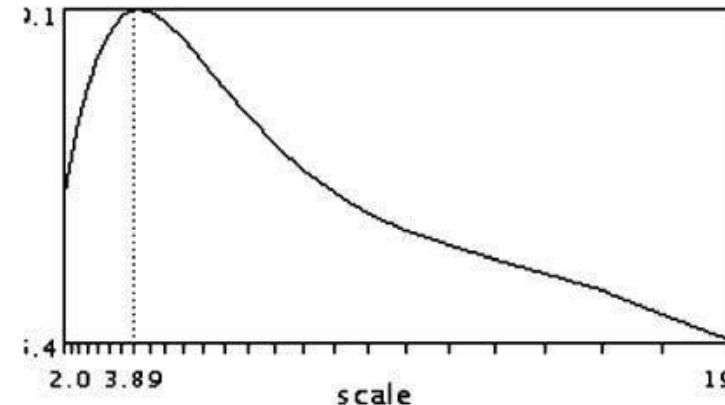
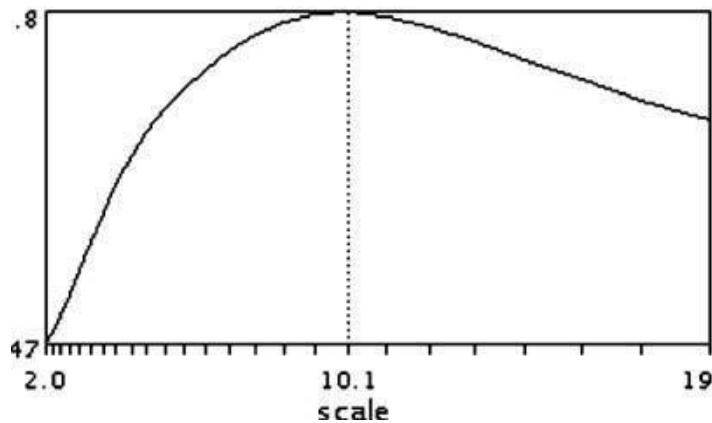
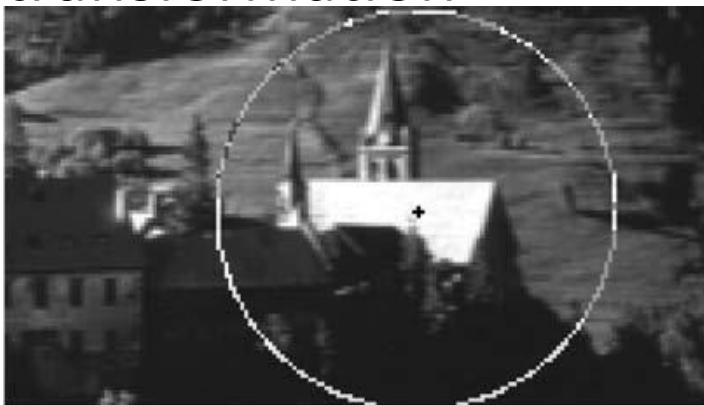
- Instead of computing  $f$  for larger and larger windows, we can implement using a fixed window size with a Gaussian pyramid



(sometimes need to create in-between levels, e.g. a  $\frac{3}{4}$ -size image)

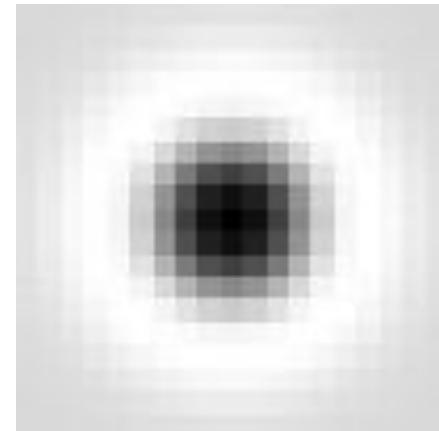
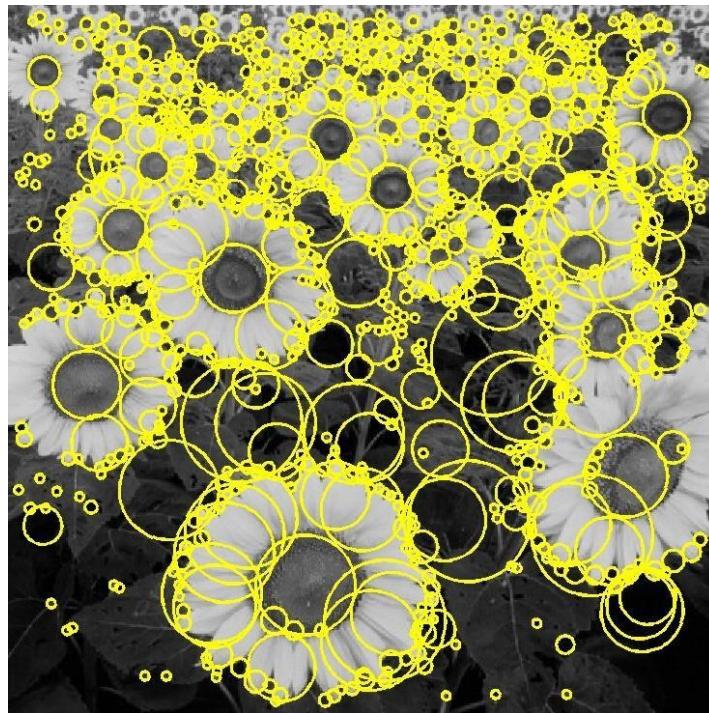
# Keypoint detection with scale selection

- We want to extract keypoints with characteristic scale that is *covariant* with the image transformation



# Basic idea

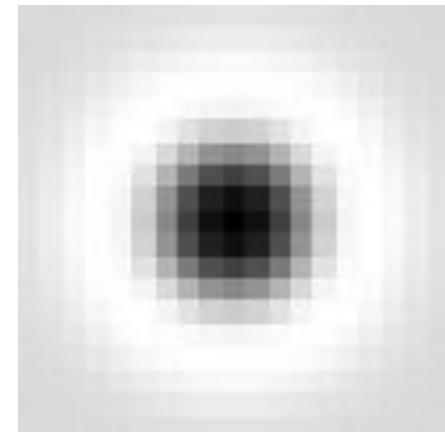
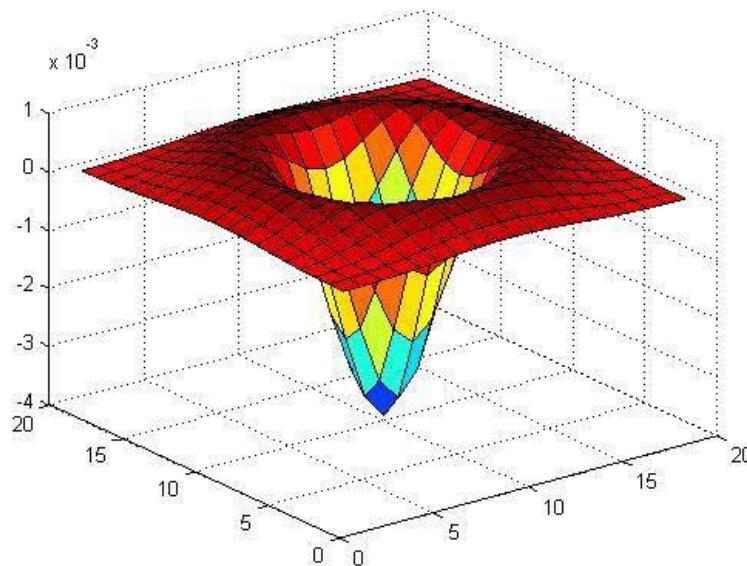
- Convolve the image with a “blob filter” at multiple scales and look for extrema of filter response in the resulting *scale space*



T. Lindeberg. [Feature detection with automatic scale selection.](#)  
*IJCV* 30(2), pp 77-116, 1998.

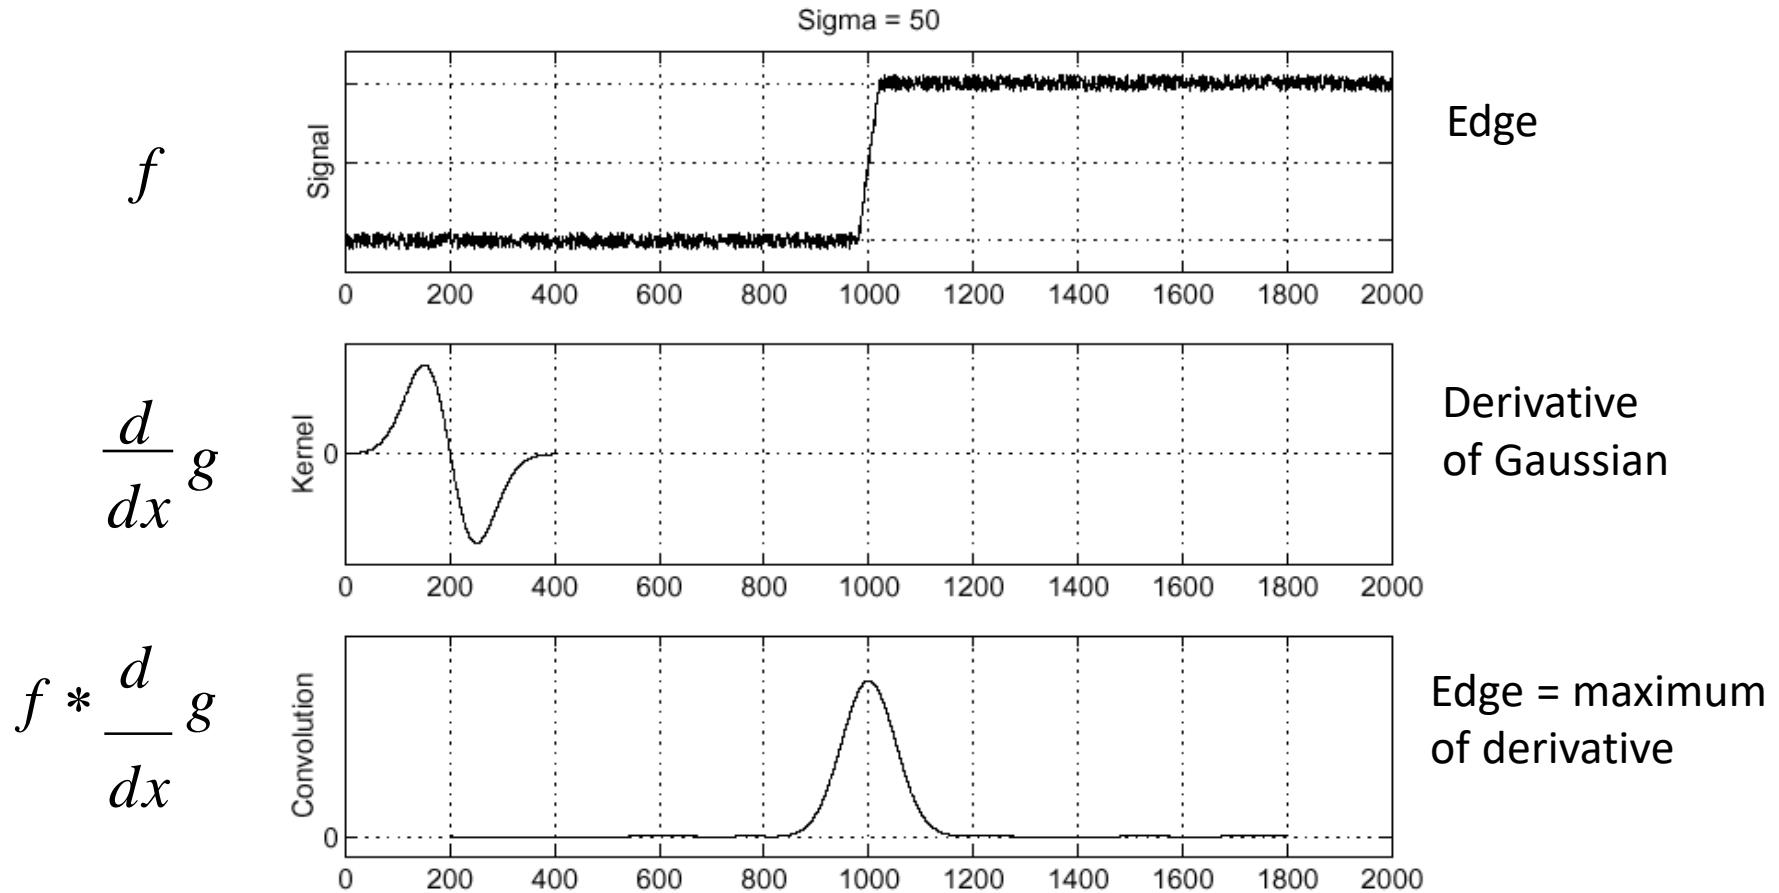
# Blob filter

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D

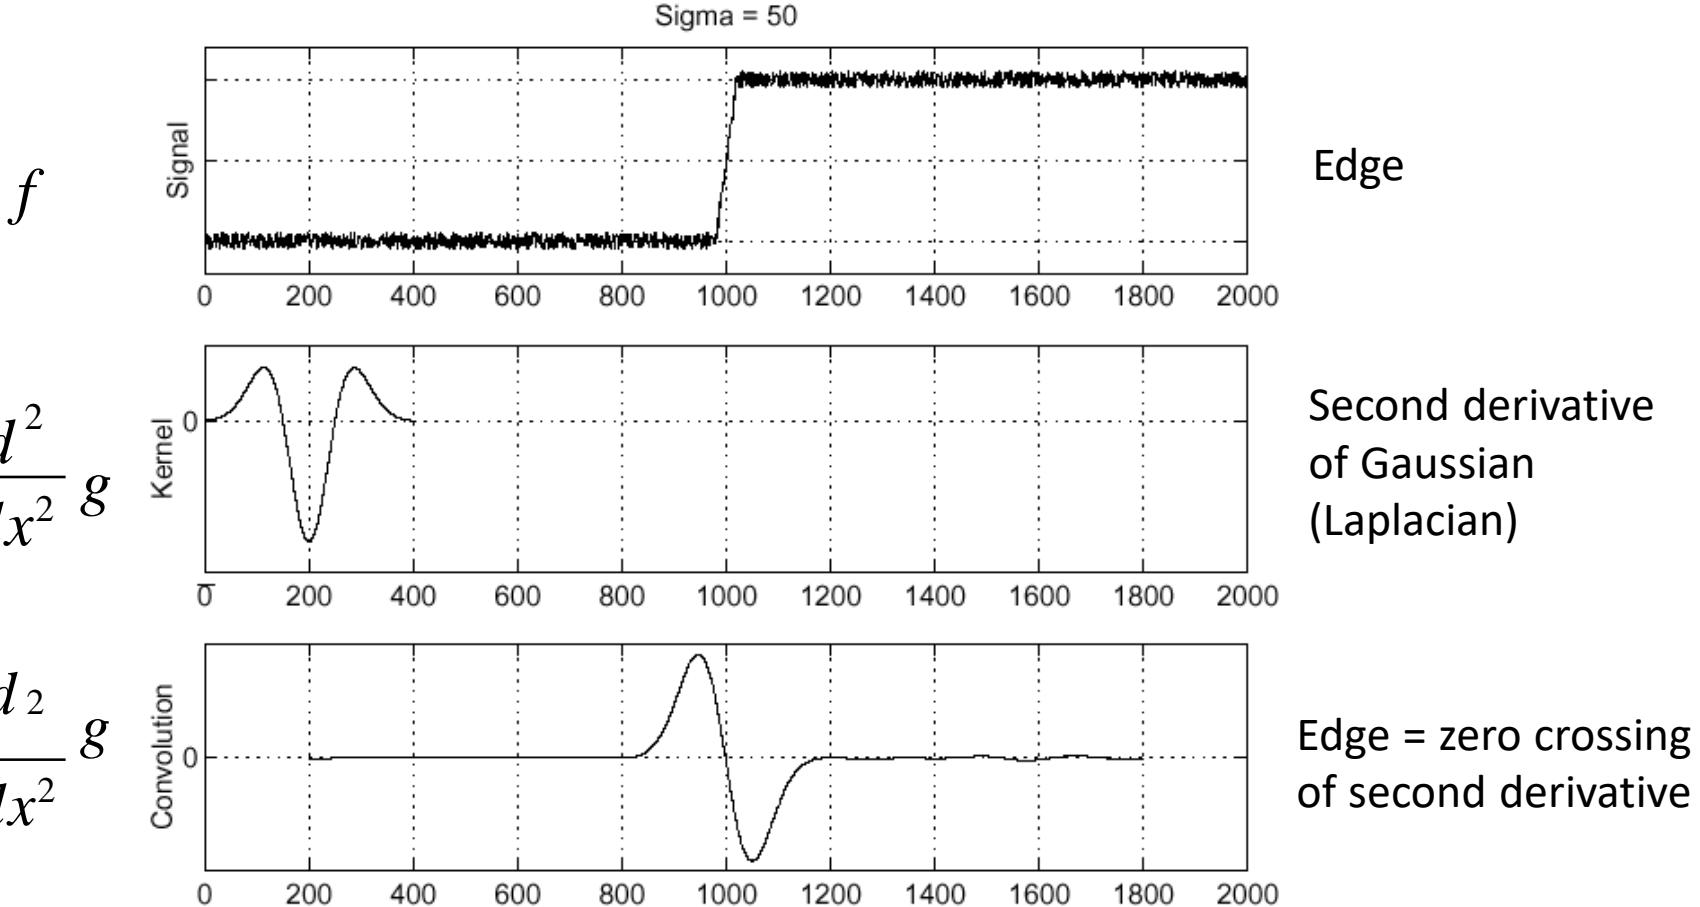


$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Recall: Edge detection

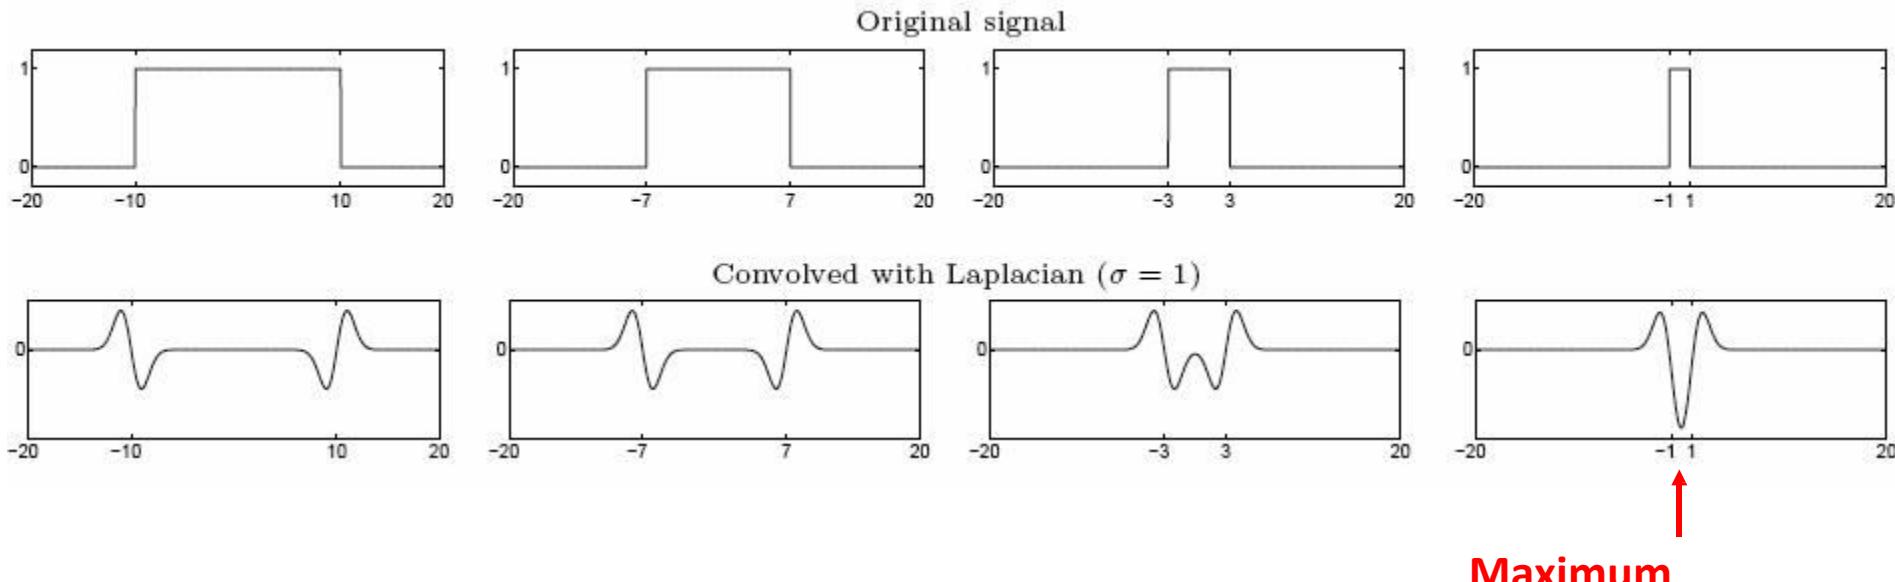


# Edge detection, Take 2



# From edges to blobs

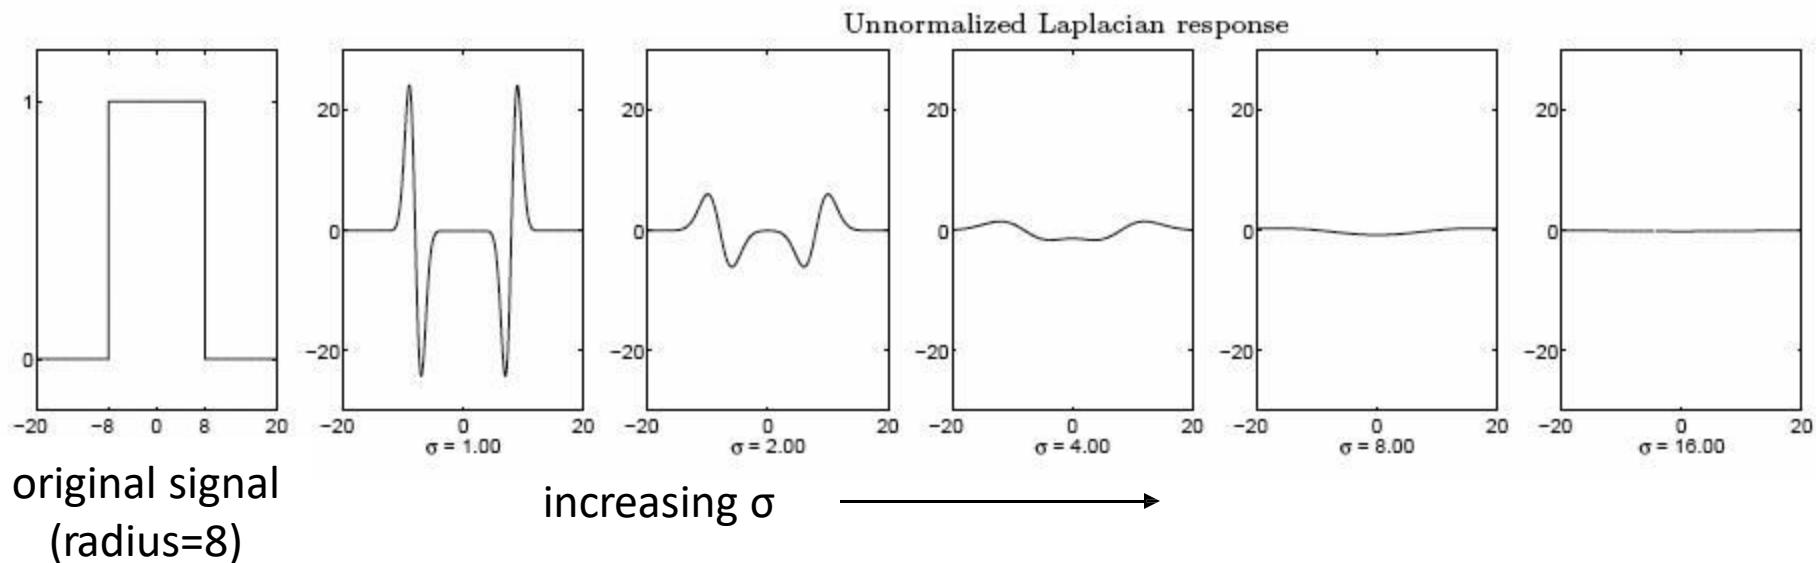
- Edge = ripple
- Blob = superposition of two ripples



**Spatial selection:** the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

# Scale selection

- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases:

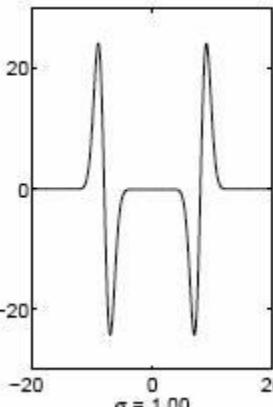
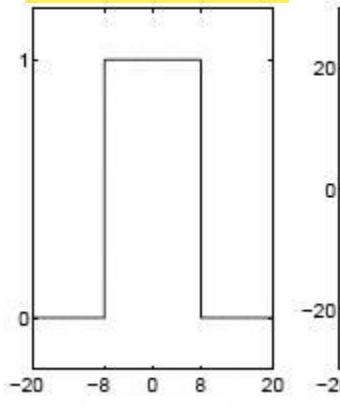


# Scale normalization

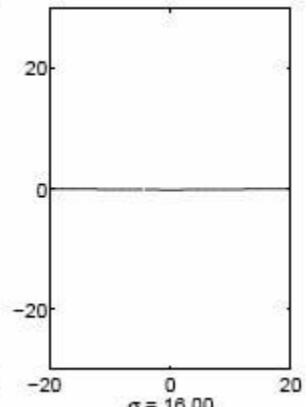
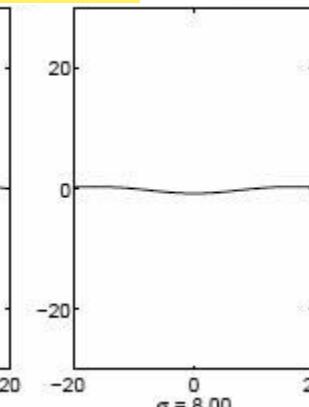
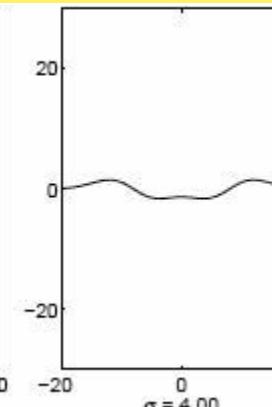
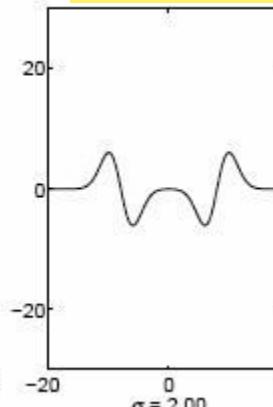
- The response of a derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by  $\sigma$
- Laplacian is the second Gaussian derivative, so it must be multiplied by  $\sigma^2$

# Effect of scale normalization

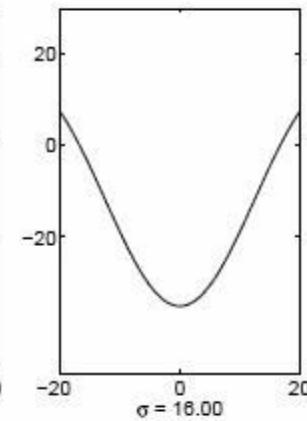
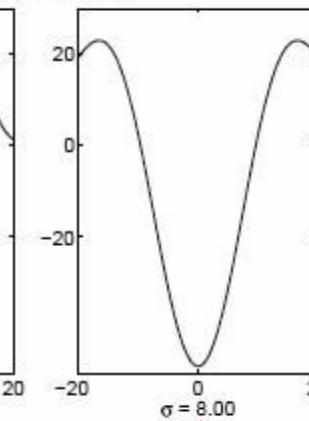
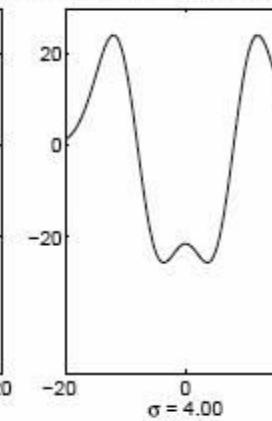
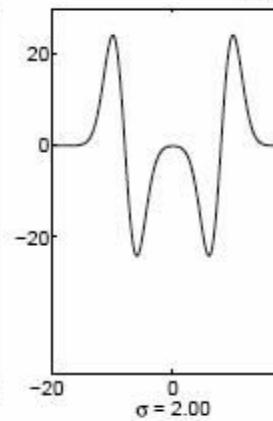
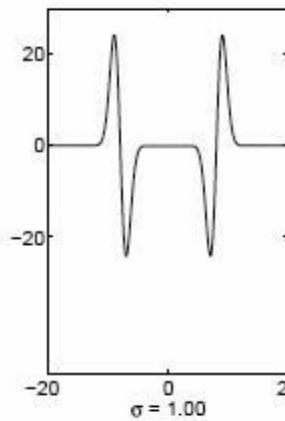
Original signal



Unnormalized Laplacian response



Scale-normalized Laplacian response

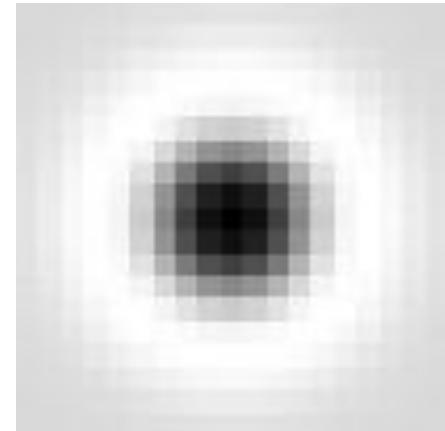
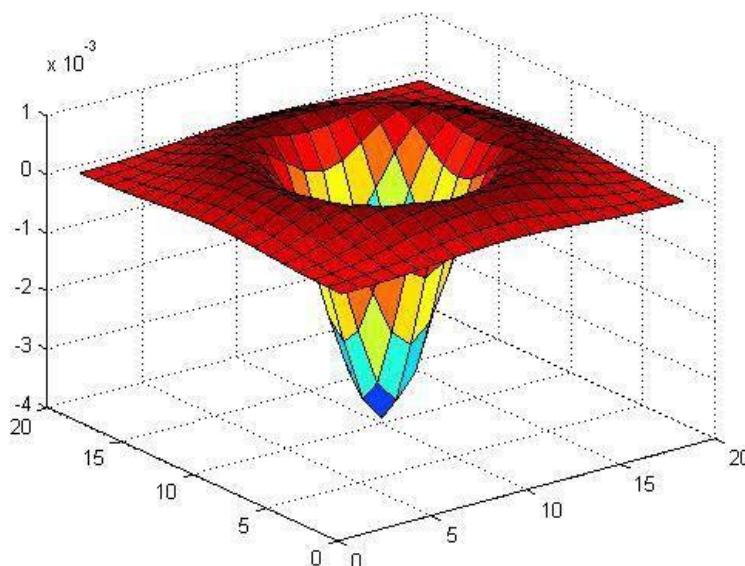


maximum

# Blob detection in 2D

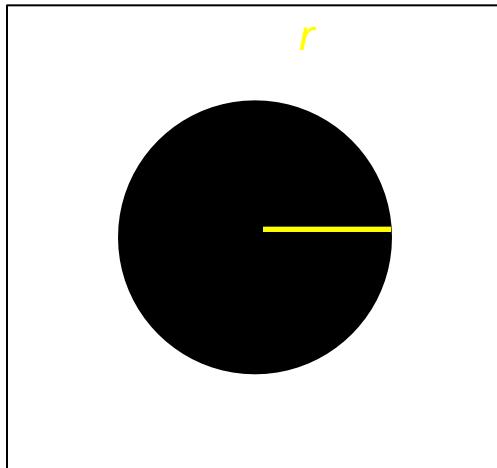
- *Scale-normalized Laplacian of Gaussian:*

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left( \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

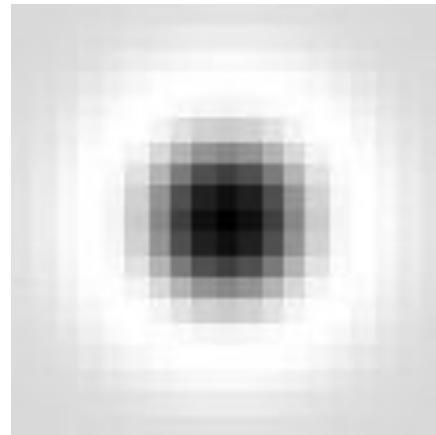


# Blob detection in 2D

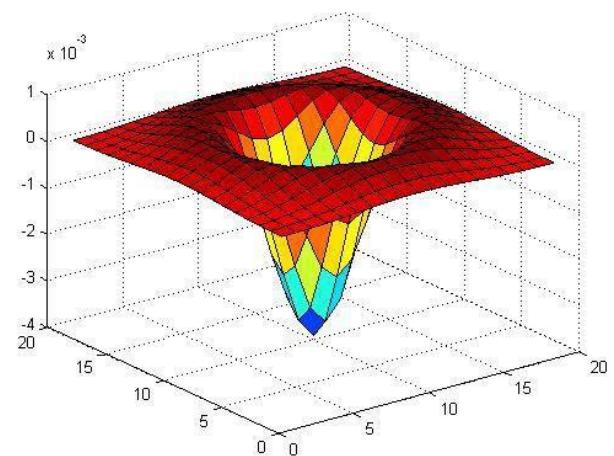
- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?



image

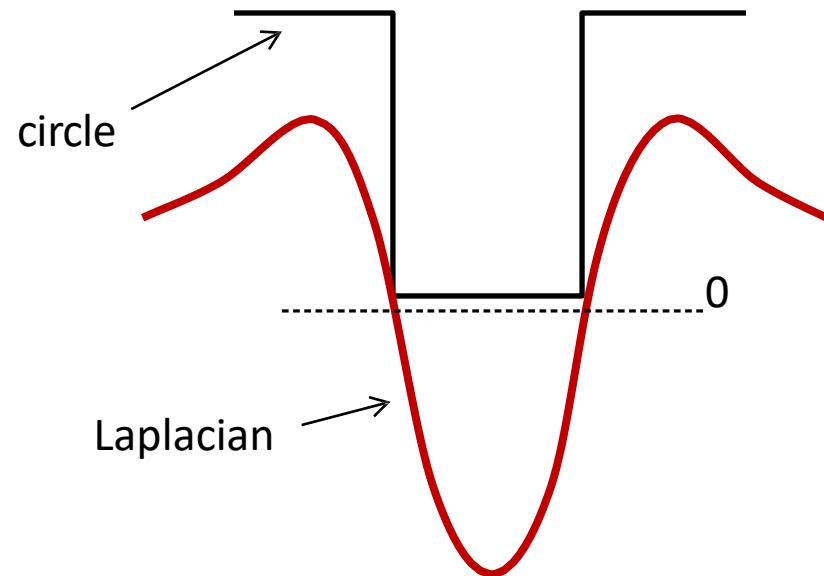
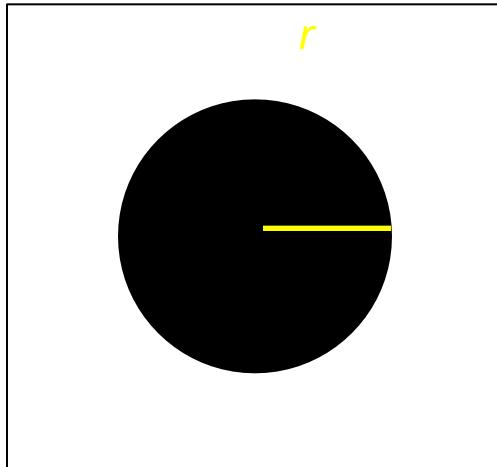


Laplacian



# Blob detection in 2D

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle

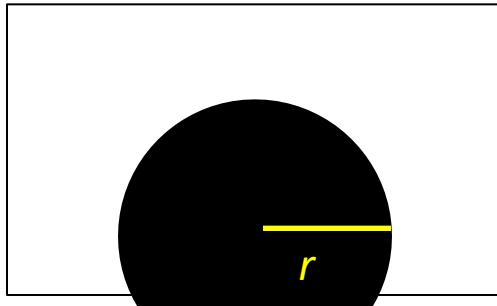


image

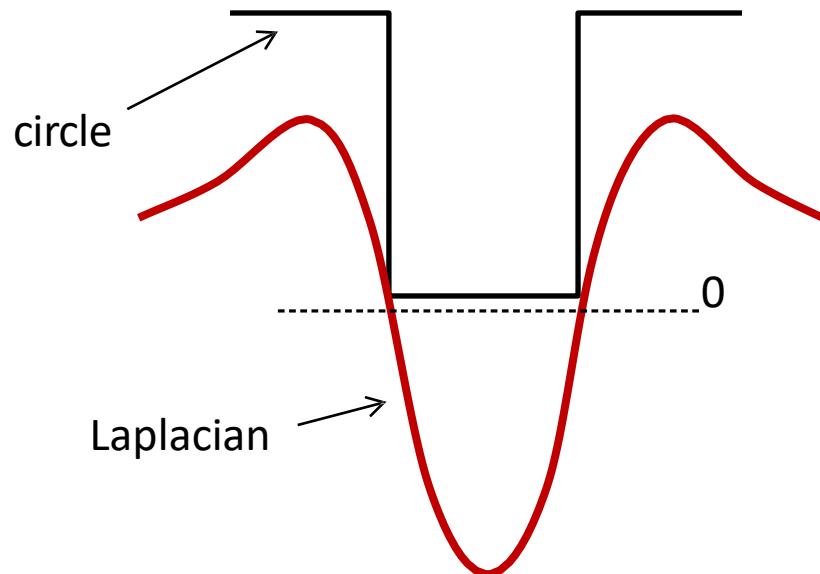
# Blob detection in 2D

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle
- The Laplacian is given by (up to scale):

$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2+y^2)/2\sigma^2}$$



image

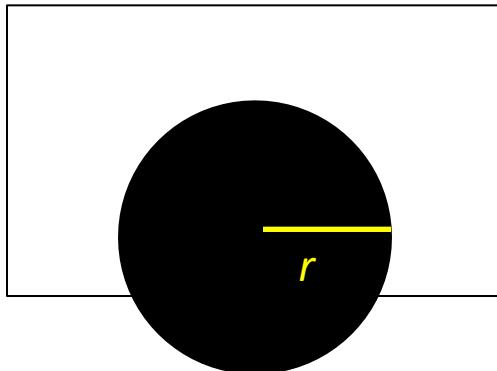


# Blob detection in 2D

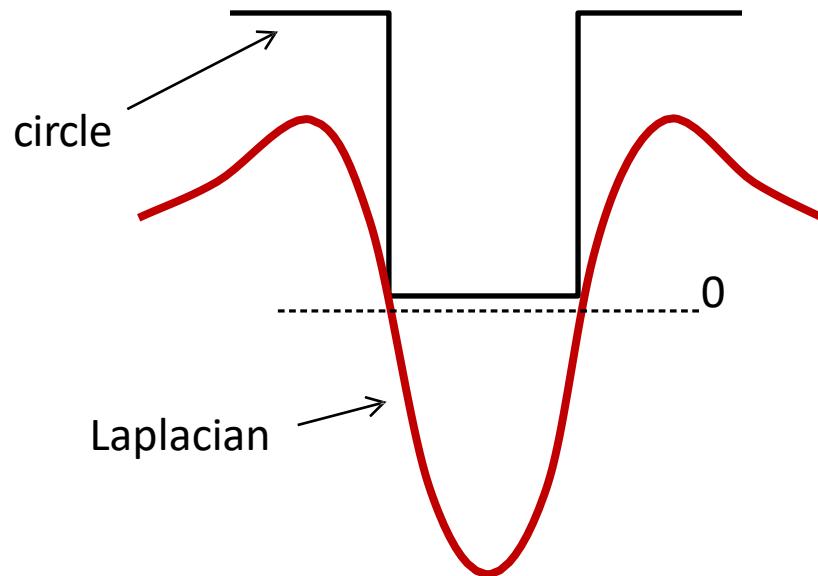
- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle
- The Laplacian is given by (up to scale):

$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2+y^2)/2\sigma^2}$$

- Therefore, the maximum response occurs at  $\sigma = r / \sqrt{2}$ .



image



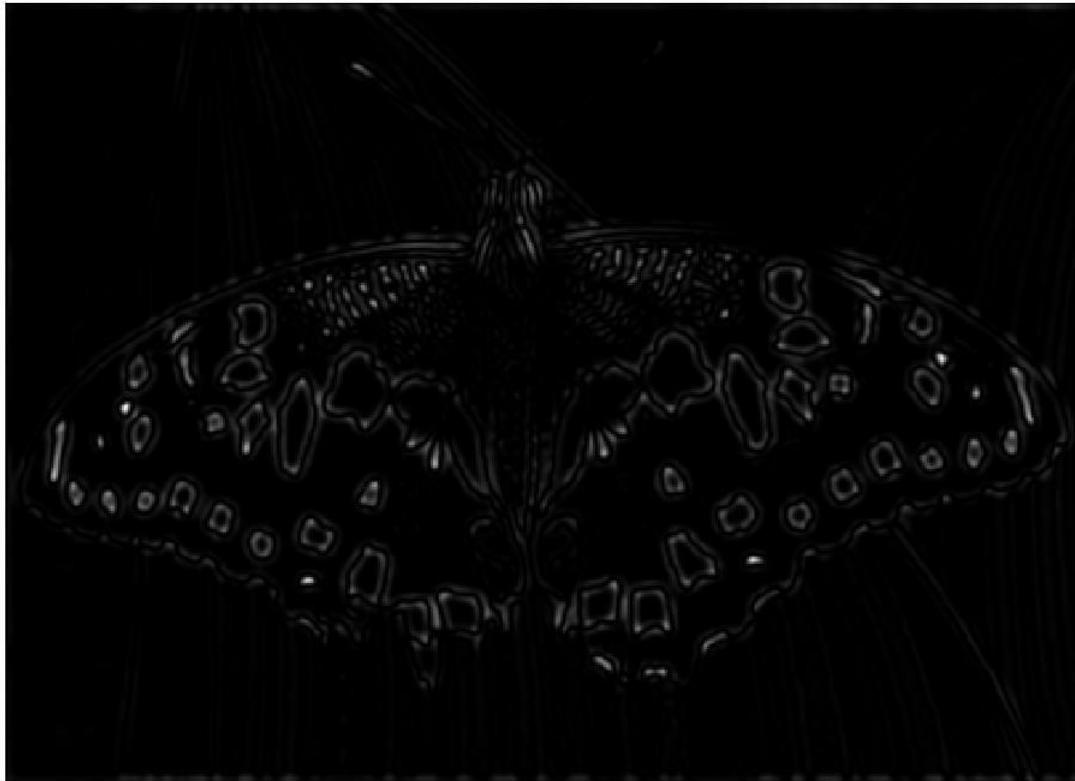
# Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales

# Scale-space blob detector: Example

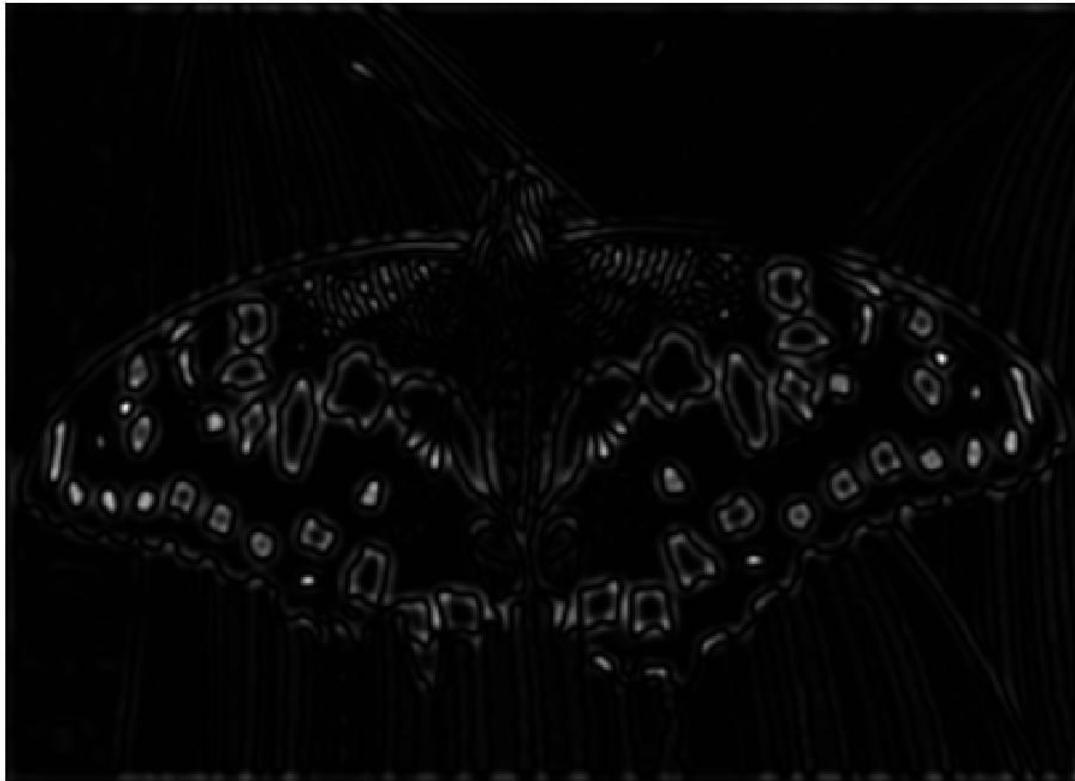


# Scale-space blob detector: Example



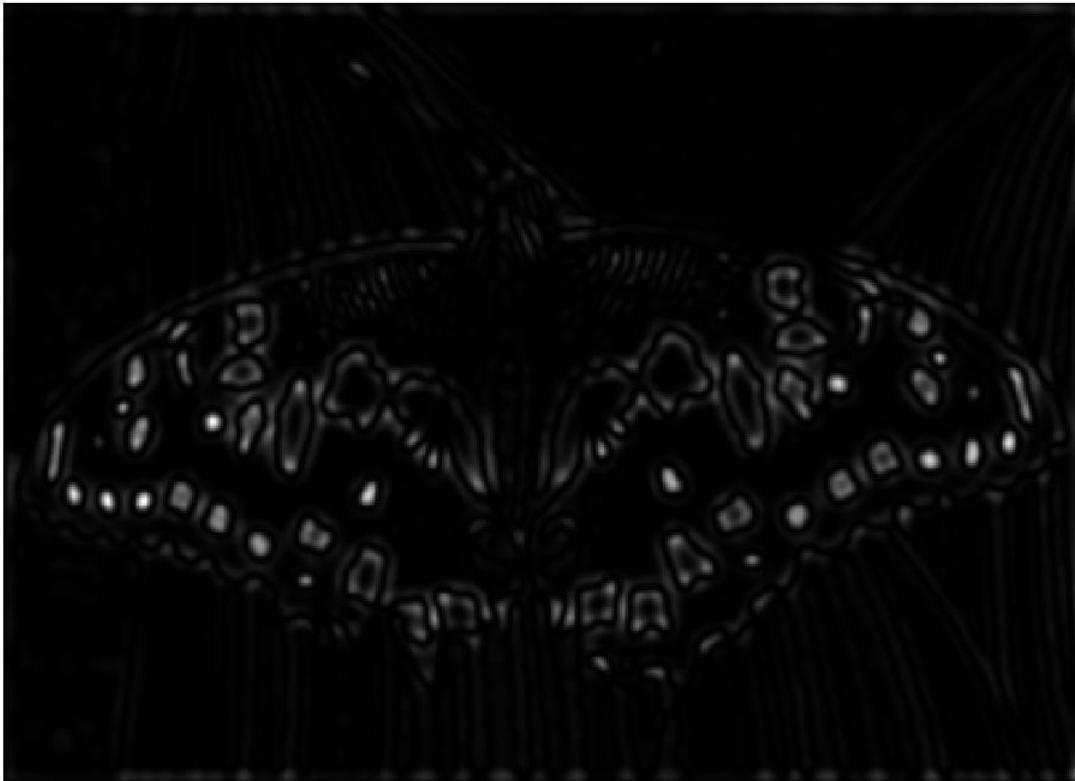
$\sigma = 2$

# Scale-space blob detector: Example



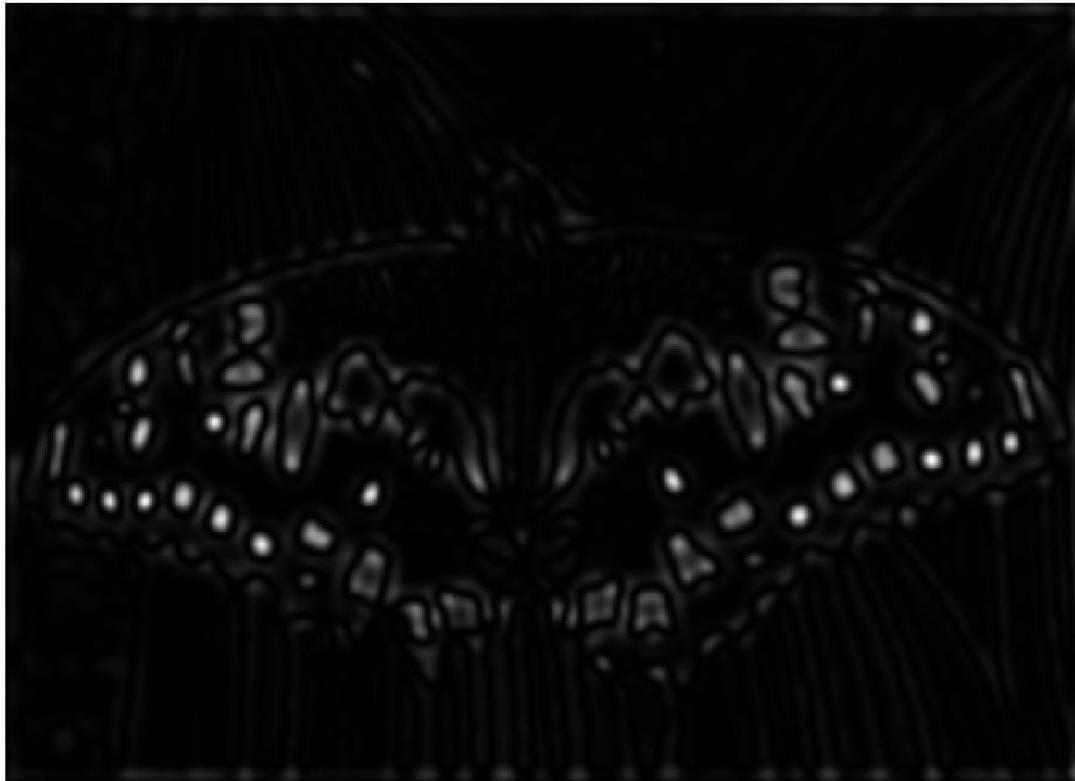
$\sigma = 2.5018$

# Scale-space blob detector: Example



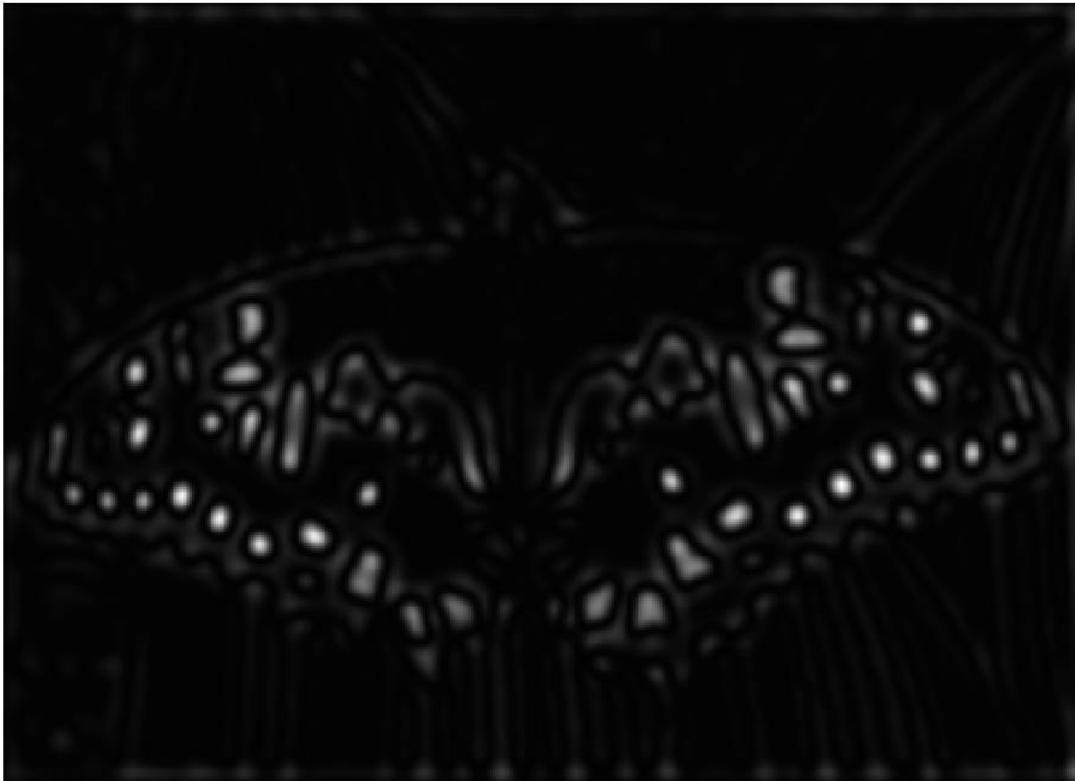
$\sigma = 3.1296$

# Scale-space blob detector: Example



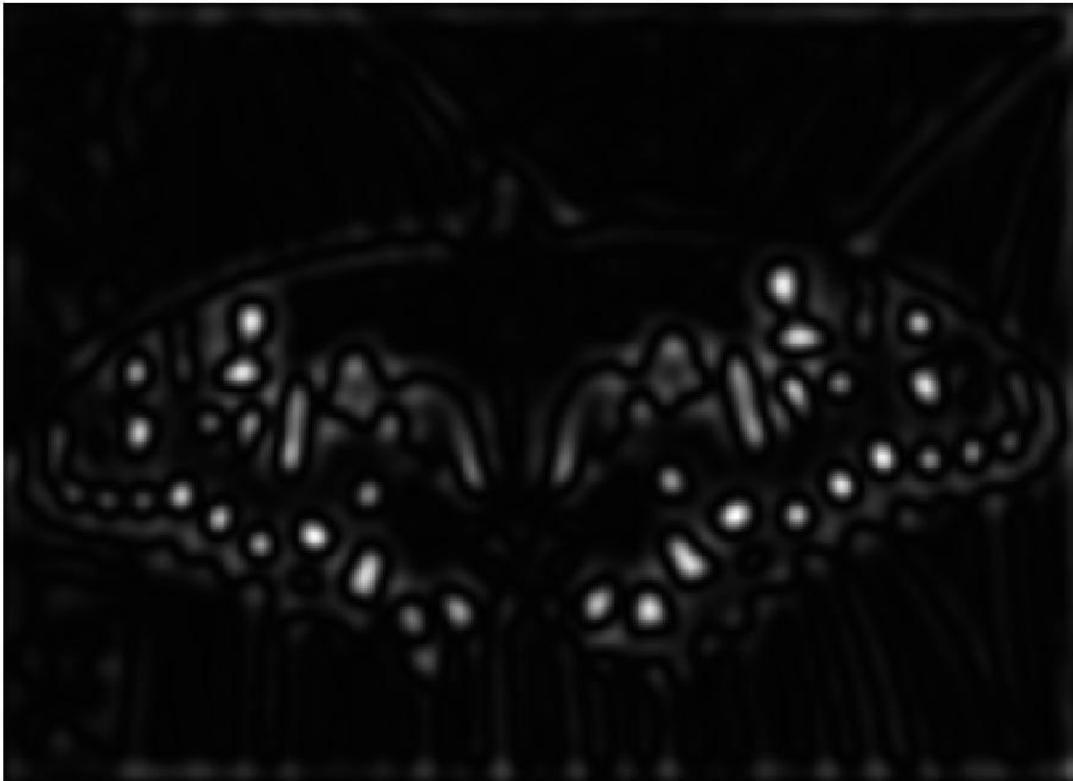
$\sigma = 3.9149$

# Scale-space blob detector: Example



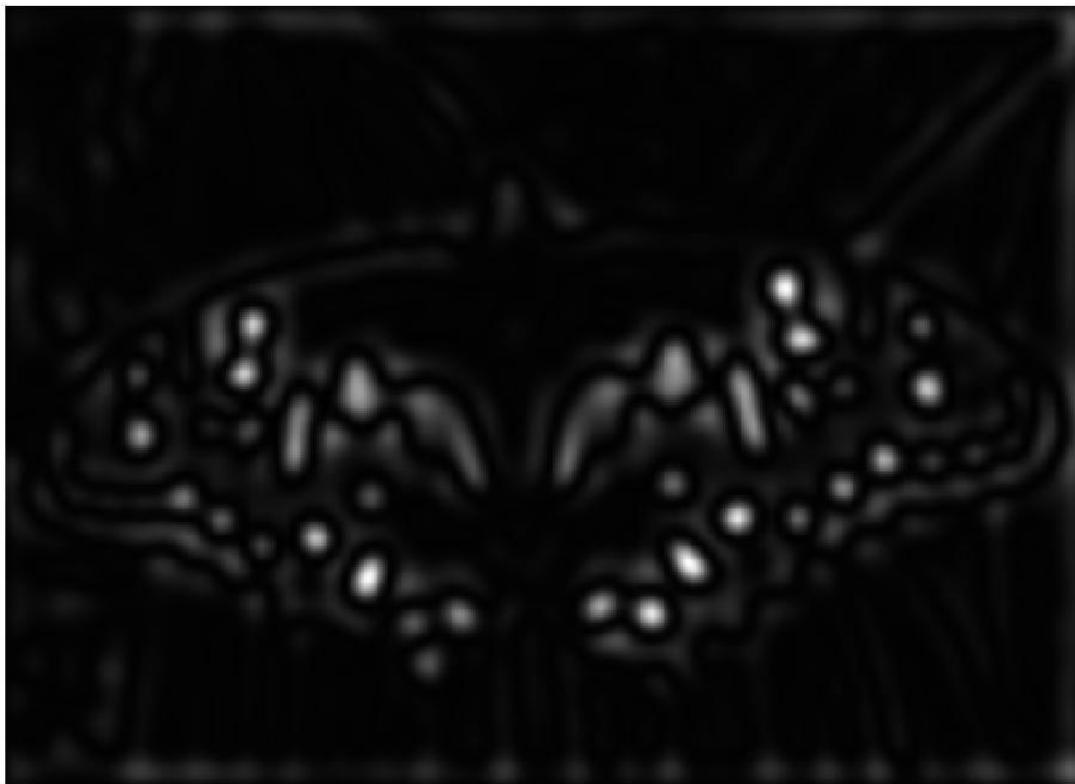
$\sigma = 4.8972$

# Scale-space blob detector: Example



$\sigma = 6.126$

# Scale-space blob detector: Example



$\sigma = 7.6631$

# Scale-space blob detector: Example



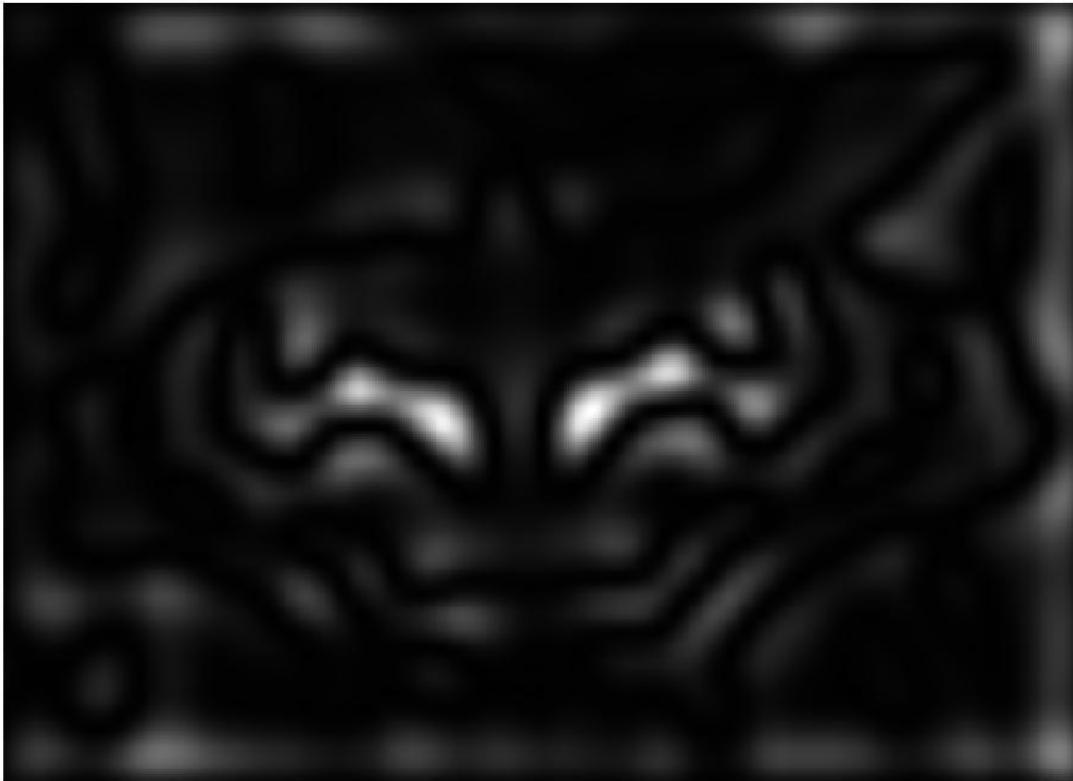
sigma = 9.5859

# Scale-space blob detector: Example



$\sigma = 11.9912$

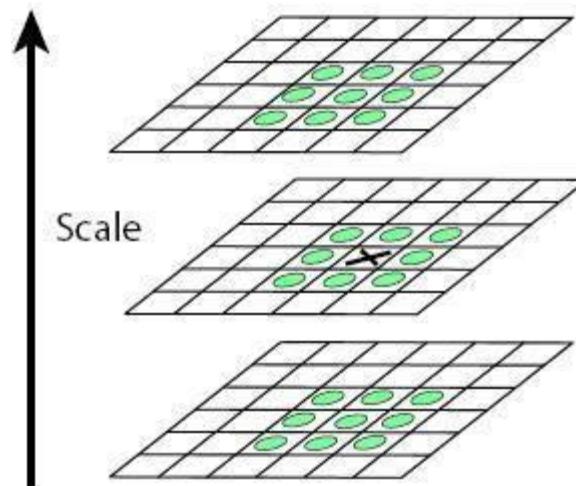
# Scale-space blob detector Example



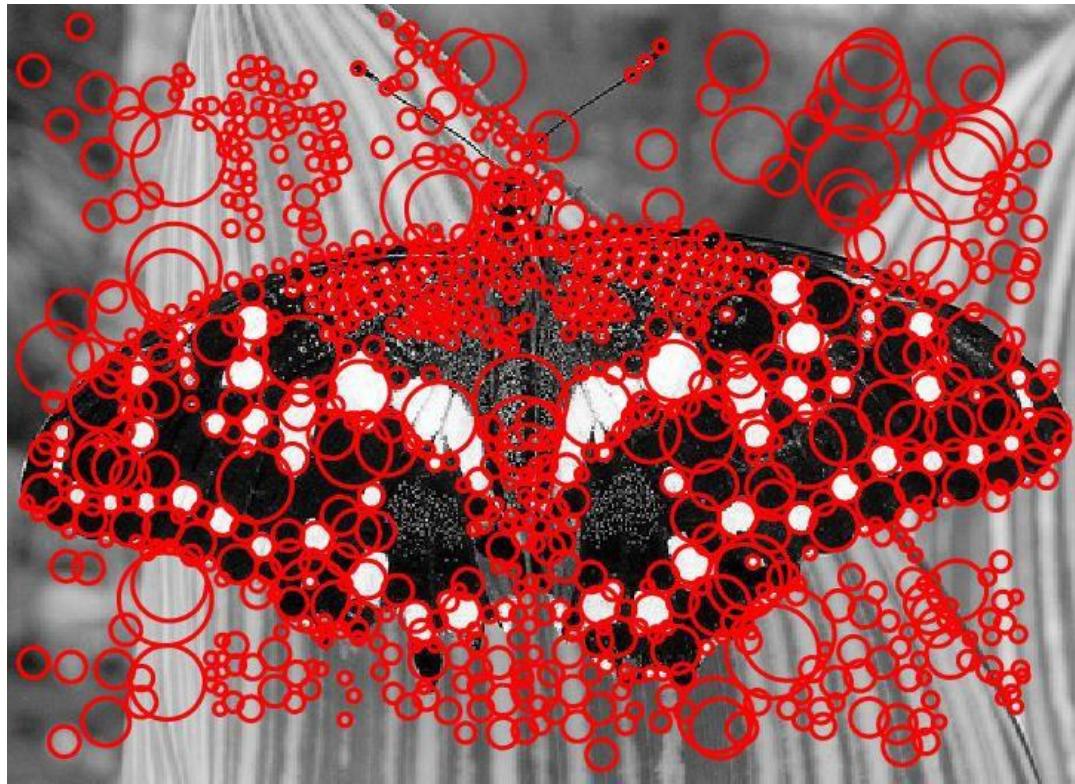
sigma = 15

# Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



# Scale-space blob detector: Example



# Efficient implementation (SIFT)

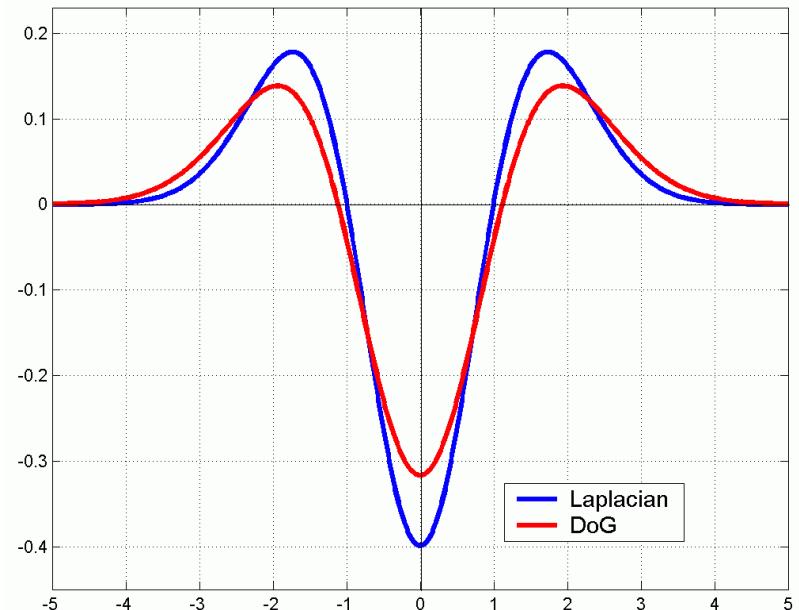
- Approximating the Laplacian with a difference of Gaussians by SIFT detector:

$$L = \sigma^2 \left( G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



# Difference-of-Gaussian (DoG)



-

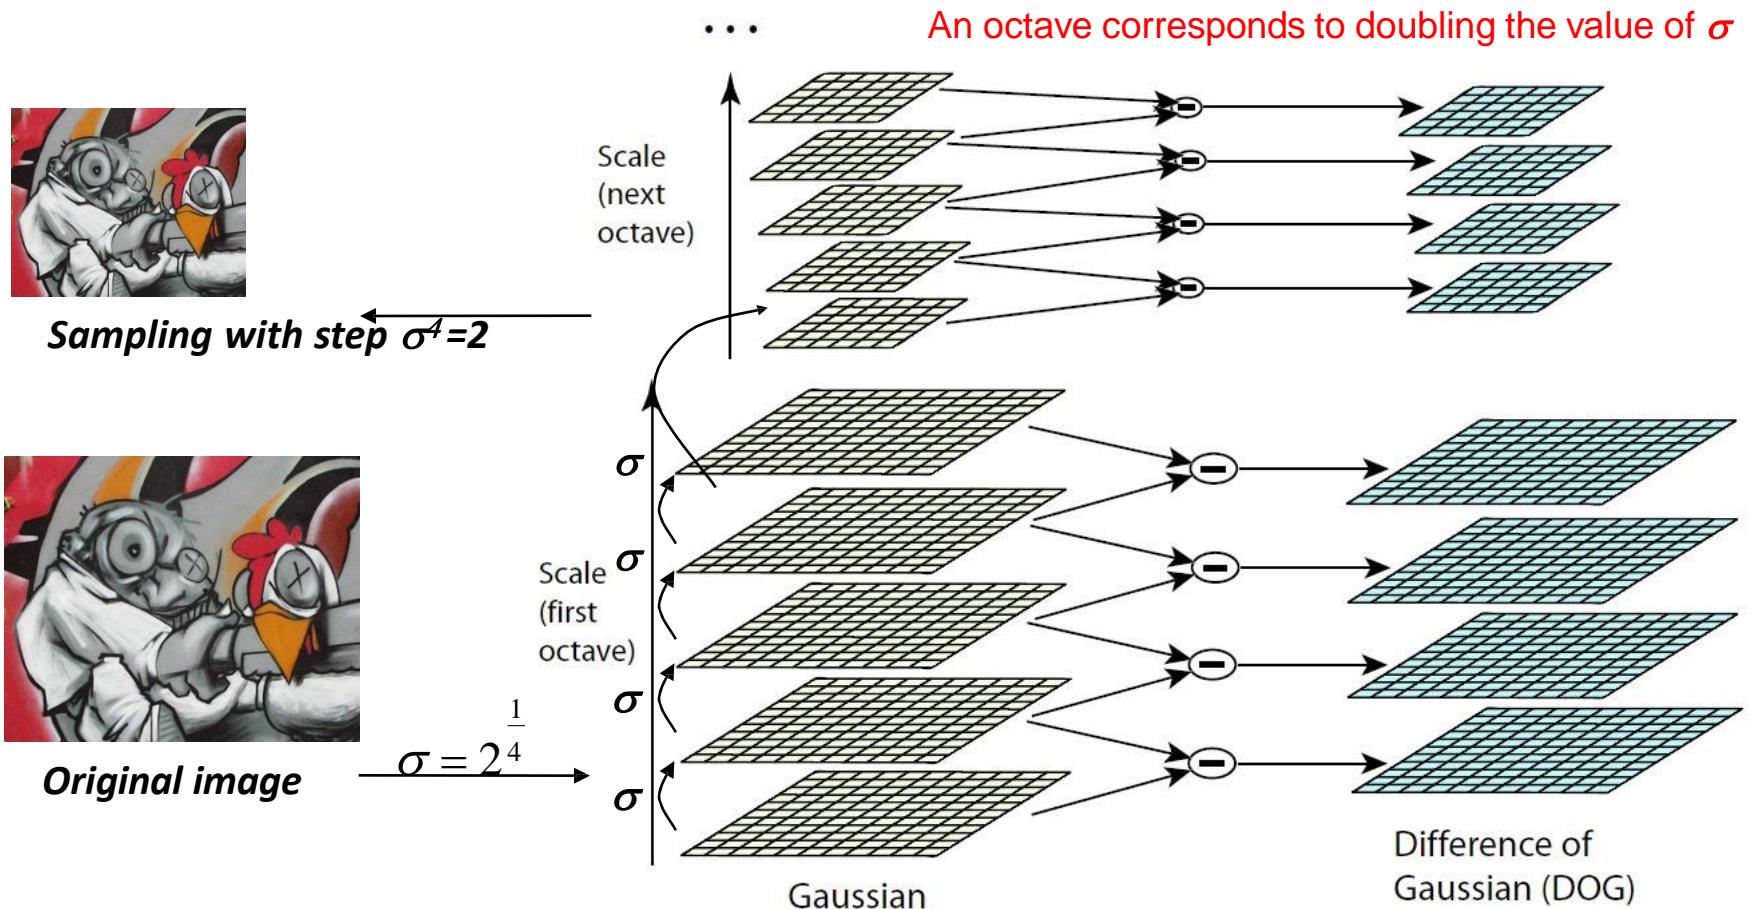


=

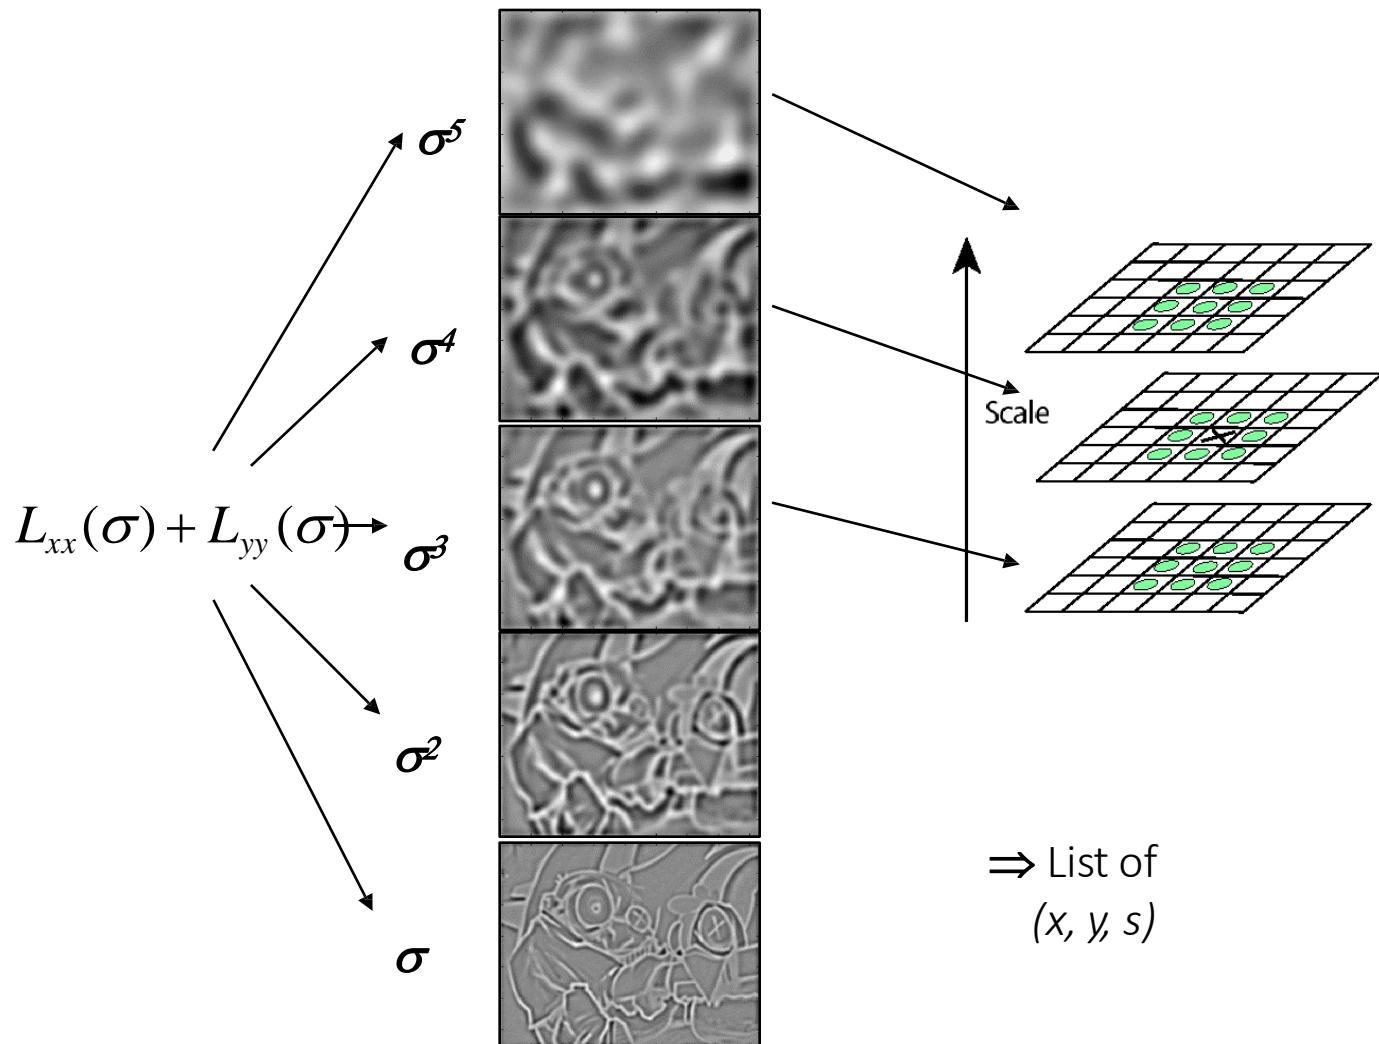
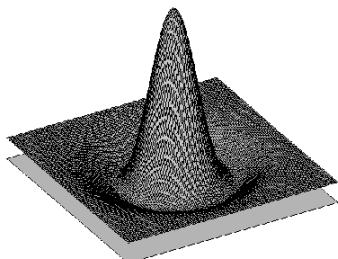


# DoG – Efficient Computation

- Computation in Gaussian scale pyramid

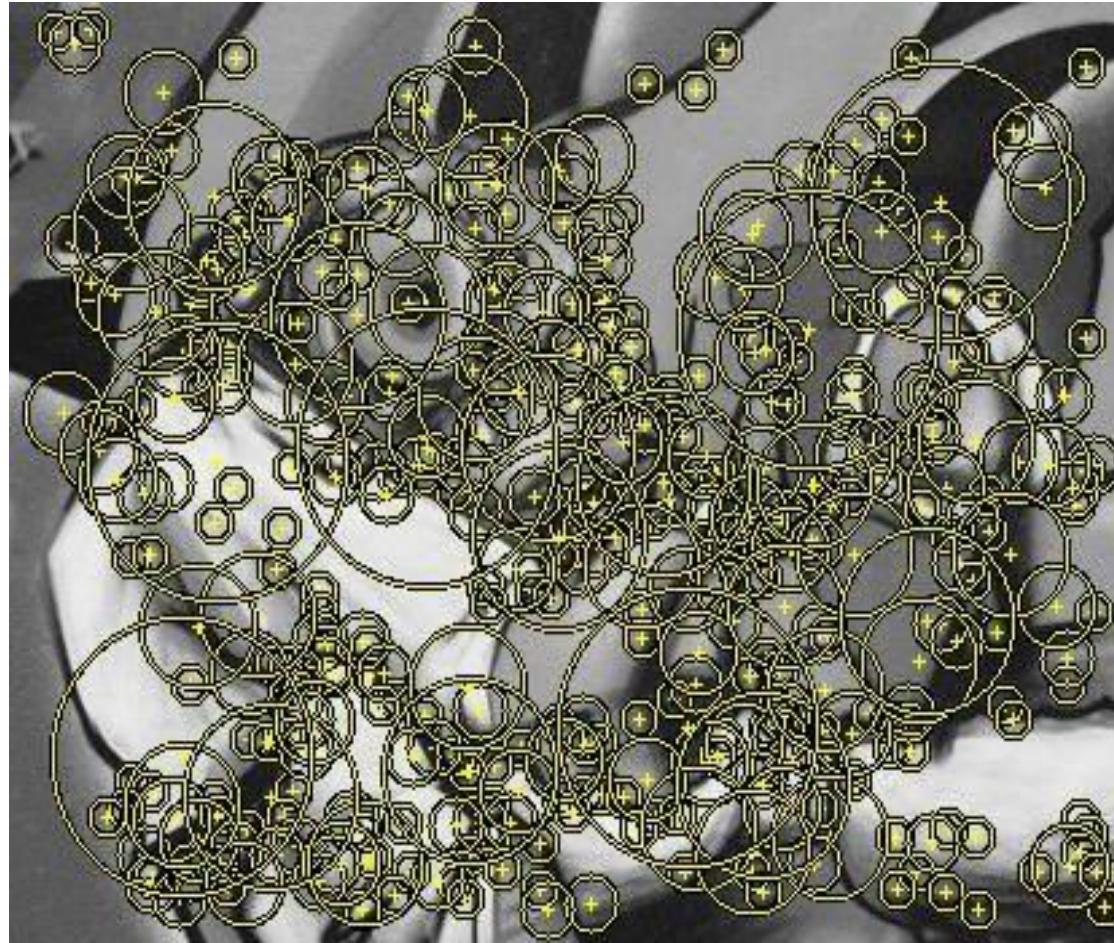


# Find local maxima in position-scale space of Difference-of-Gaussian



⇒ List of  
 $(x, y, s)$

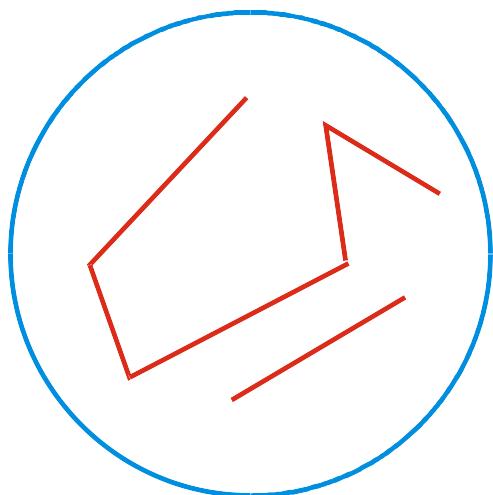
# Results: Difference-of-Gaussian



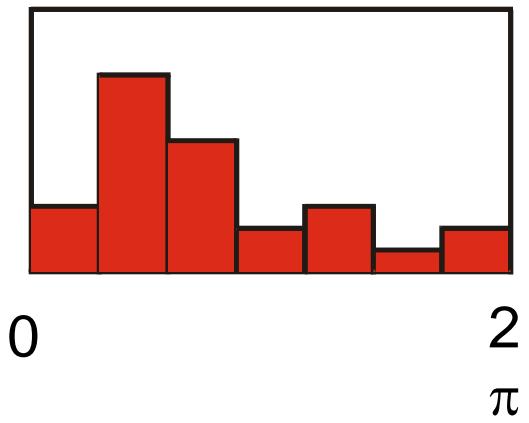
# Orientation Normalization

[Lowe, SIFT, 1999]

- Compute orientation histogram



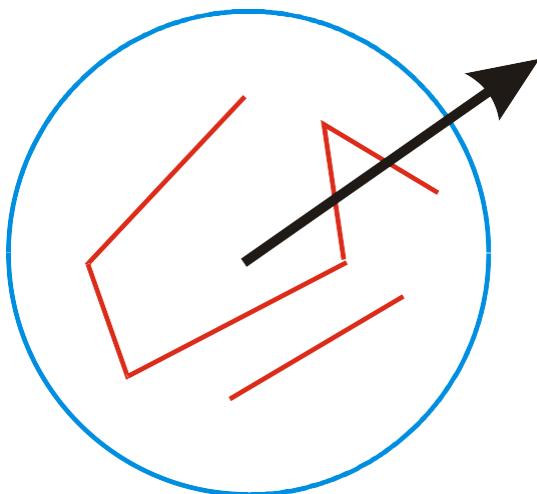
T. Tuytelaars, B. Leibe



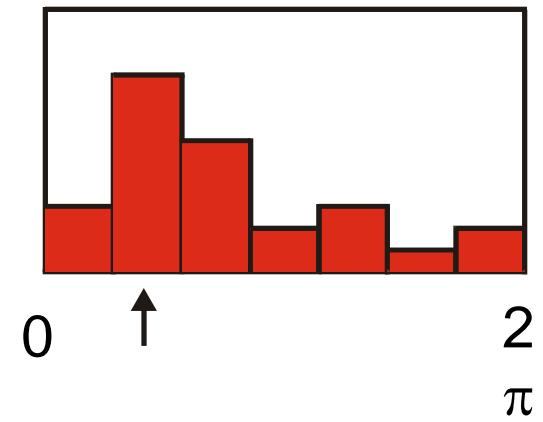
# Orientation Normalization

[Lowe, SIFT, 1999]

- Compute orientation histogram
- Select dominant orientation



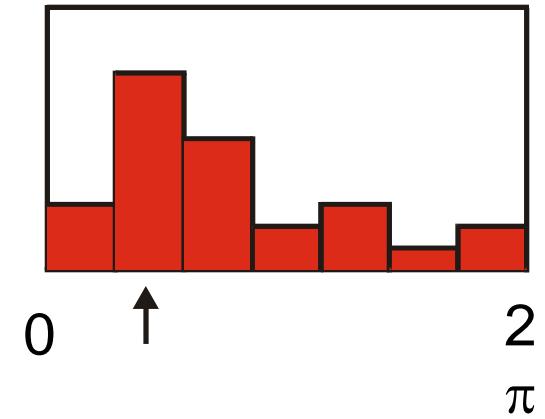
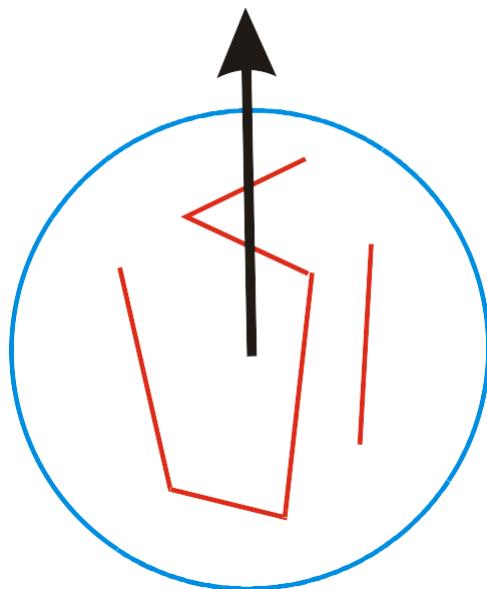
T. Tuytelaars, B. Leibe



# Orientation Normalization

[Lowe, SIFT, 1999]

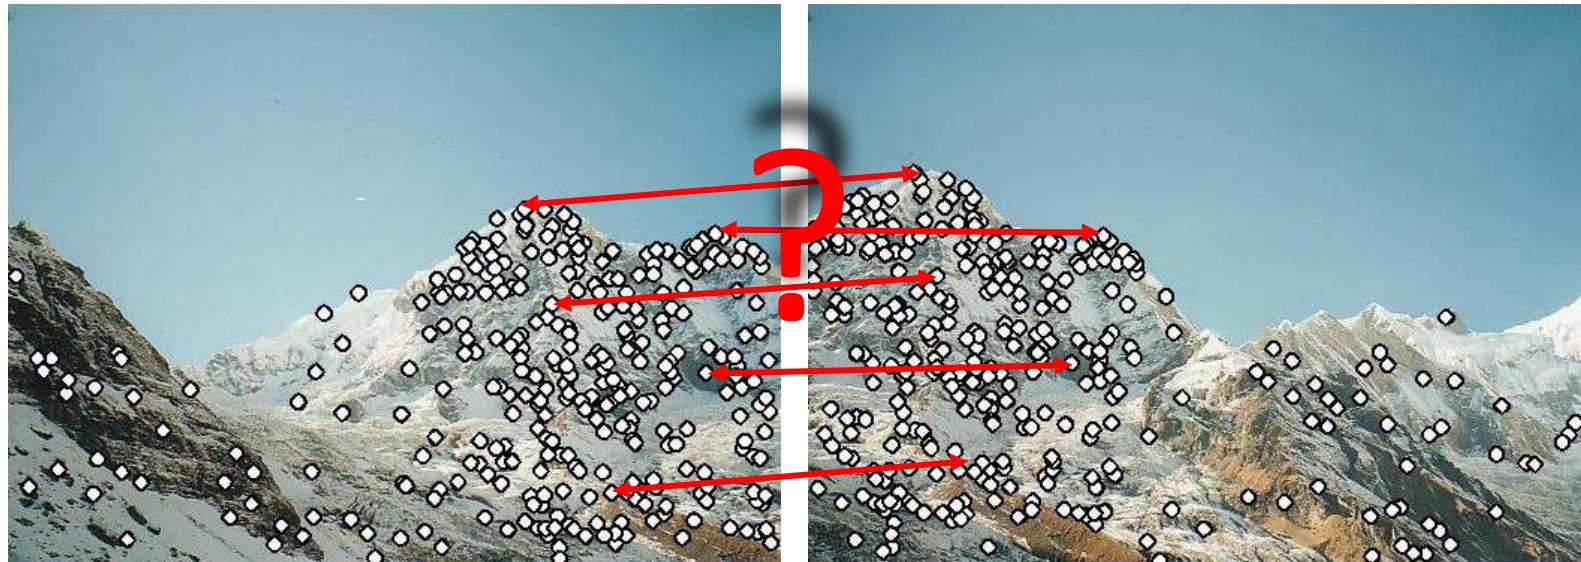
- Compute orientation histogram
- Select dominant orientation
- Normalize: rotate to fixed orientation



# Feature descriptors

We know how to detect good points

Next question: **How to match them?**

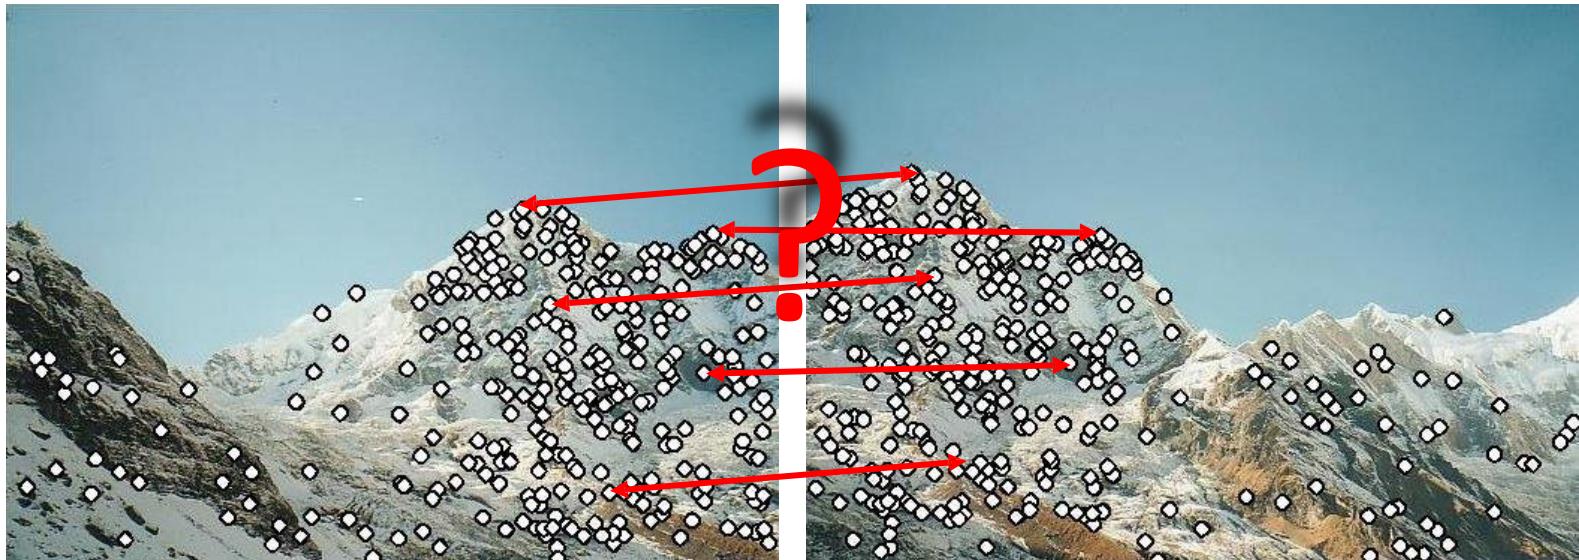


**Answer:** Come up with a *descriptor* for each point,  
find similar descriptors between the two images

# Feature descriptors

We know how to detect good points Next question:

**How to match them?**



Lots of possibilities (this is a popular research area)

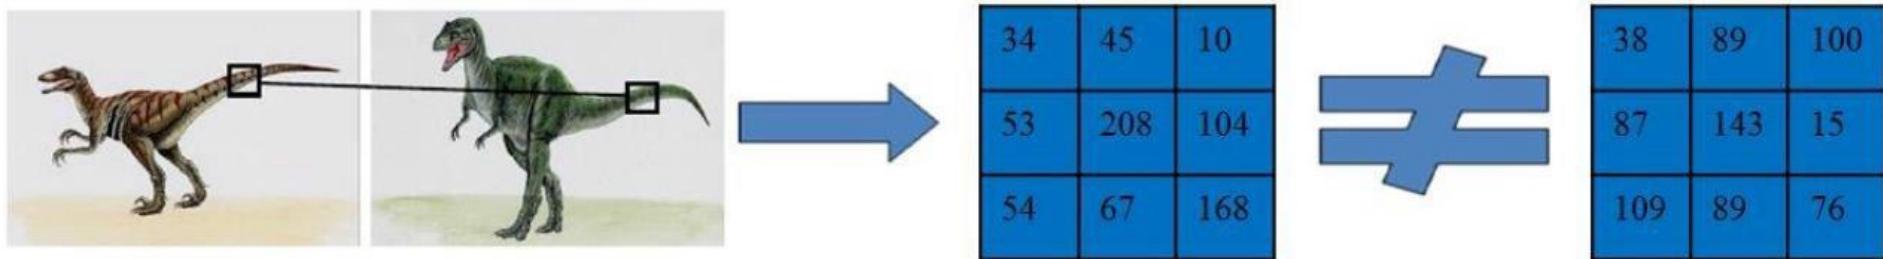
- Simple option: match square windows around the point
- State of the art approach: SIFT
  - David Lowe, UBC <http://www.cs.ubc.ca/~lowe/keypoints/>

# Image/Region Matching

- Automatically recognize whether two images/regions contain the similar content.

# Image/Region Matching

- Automatically recognize whether two images/regions contain the similar content.
- Comparing the image pixels as they are, will not work.



# Image/Region Matching

- Pixel-based distances on high-dimensional data (and images especially) can be very unintuitive.

original



shifted



messed up



darkened



# Challenges

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation

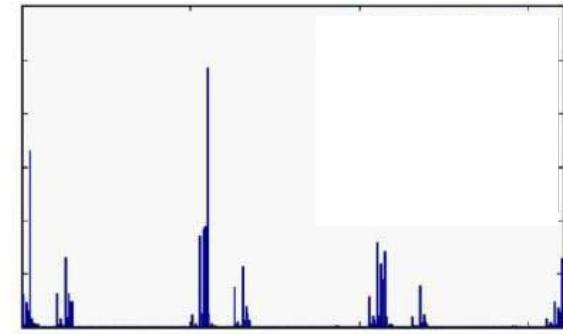


# Solution

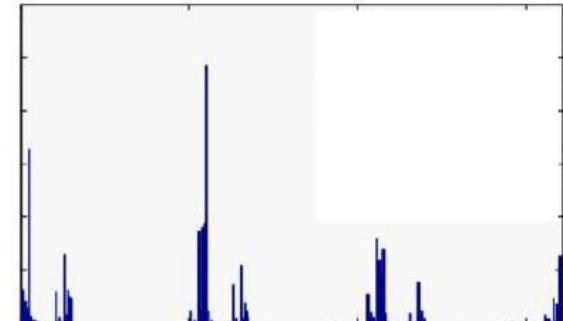
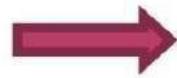
- Descriptors allow certain differences between the images.



Descriptor Function



Similar ?



Similar ?

Comparing using descriptor function

(images are taken from Corel-database and RSHD descriptor is used)

# Applications

- Image Matching
- Image Retrieval
- Biomedical Image Analysis
- Texture Classification
- Image Correspondence
- Face Analysis
- Biometrics
- Building Panorama
- And many more...

## Image descriptor

- Descriptions of the visual features
- Described by appearance based characteristics such as color, shape, etc.

## Image descriptor

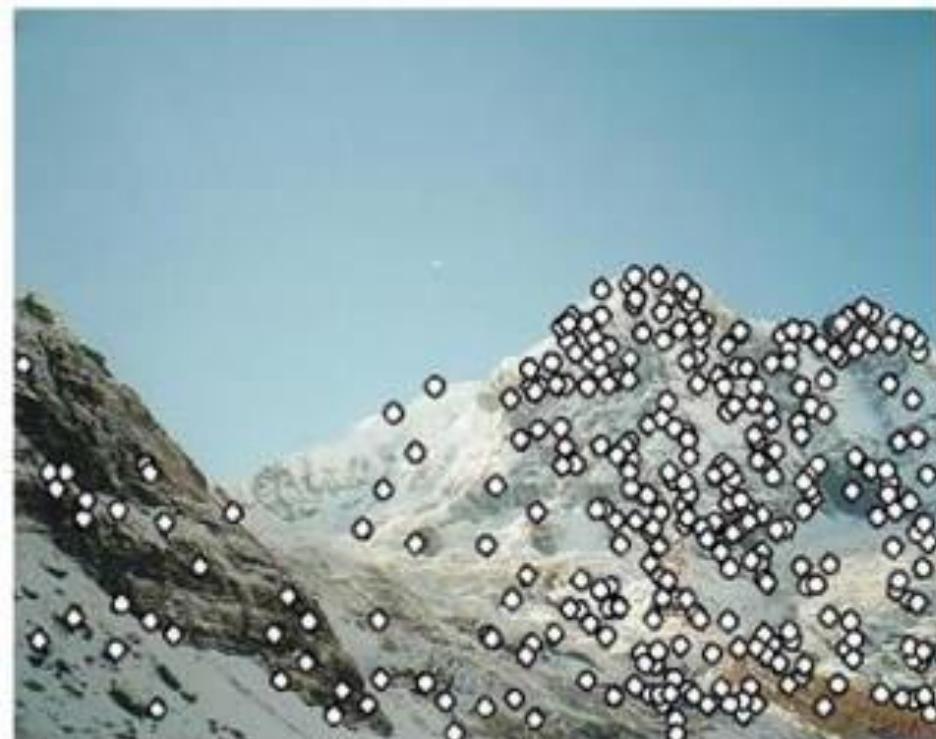
- Descriptions of the visual features
- Described by appearance based characteristics such as color, shape, etc.

## A descriptor must be

- Distinctive
- Robust
- Compact
- Low Dimensional

# Where to compute the descriptors?

- Over interest regions.
- Interest region may be
  - Key-Points or Global based.



# Local Descriptors

- Most available descriptors focus on -
  - Edge/gradient information
  - Capture texture information
  - Exploit local relationship
  - Color also play a vital role
  - Shape features
  - Feature fusion

# Widely Used Local Descriptors

- SIFT – Scale Invariant Feature Transform

Distinctive image features from scale-invariant keypoints

DG Lowe - International journal of computer vision, 2004 - Springer

... the assigned orientation, **scale**, and location for each **feature**, thereby providing **invariance** to these ... **Invariant Feature Transform** (SIFT), as it **transforms** image data into **scale-invariant** coordinates relative ... that densely cover the image over the full range of **scales** and locations ...

☆ 99 [Cited by 50252](#) Related articles All 179 versions

- LBP – Local Binary Pattern

Multiresolution gray-scale and rotation invariant texture classification with **local binary patterns**

T Ojala, M Pietikainen... - ... Transactions on pattern ..., 2002 - ieeexplore.ieee.org

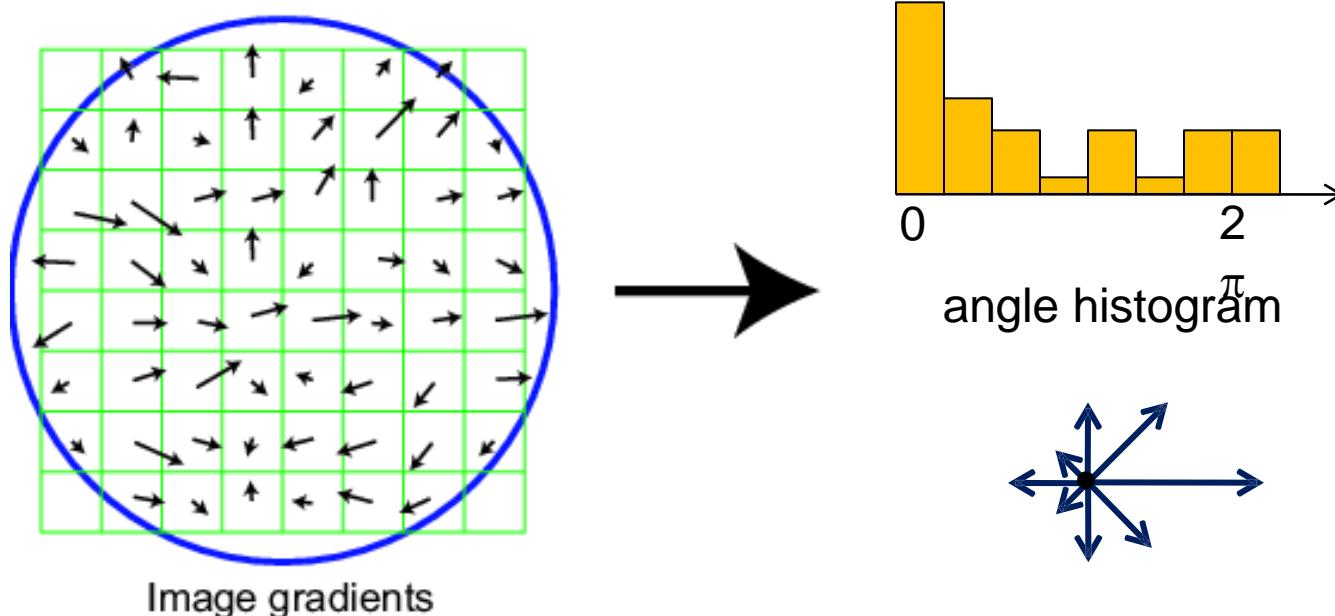
Presents a theoretically very simple, yet efficient, multiresolution approach to gray-scale and rotation invariant texture classification based on **local binary patterns** and nonparametric discrimination of sample and prototype distributions. The method is based on recognizing ...

☆ 99 [Cited by 12251](#) Related articles All 16 versions

# Scale Invariant Feature Transform

Basic idea:

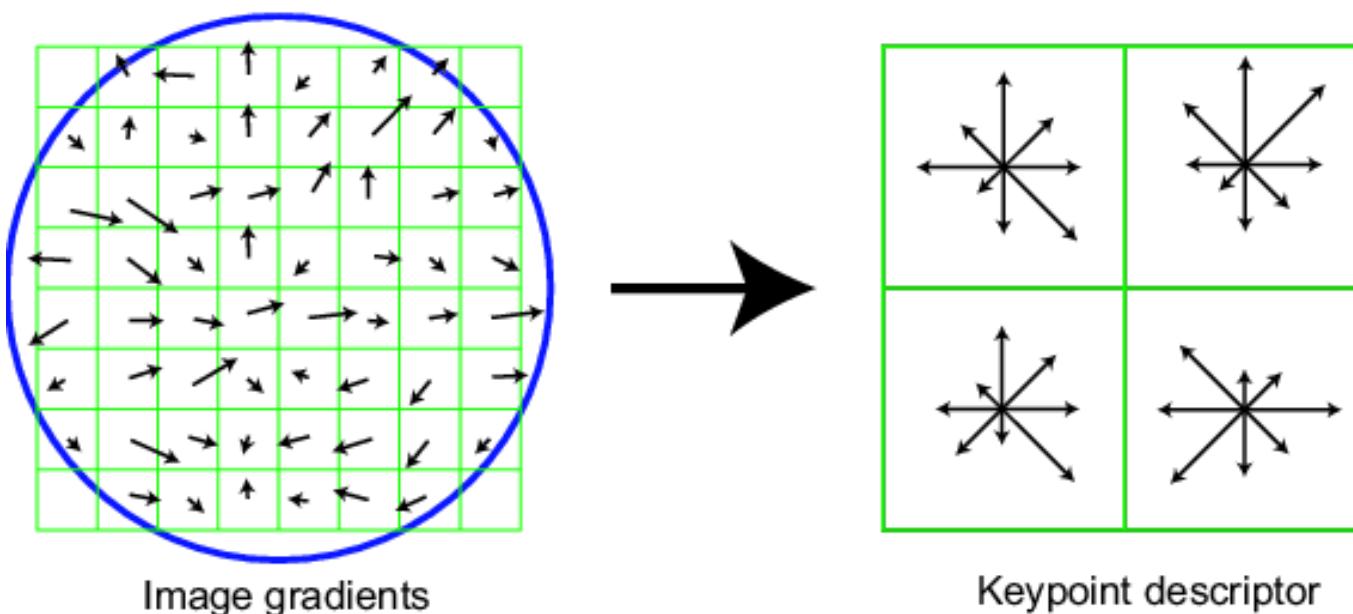
- Take 16x16 square window around detected feature
- Compute edge orientation for each pixel
- Throw out weak edges (threshold gradient magnitude)
- Create histogram of surviving edge orientations



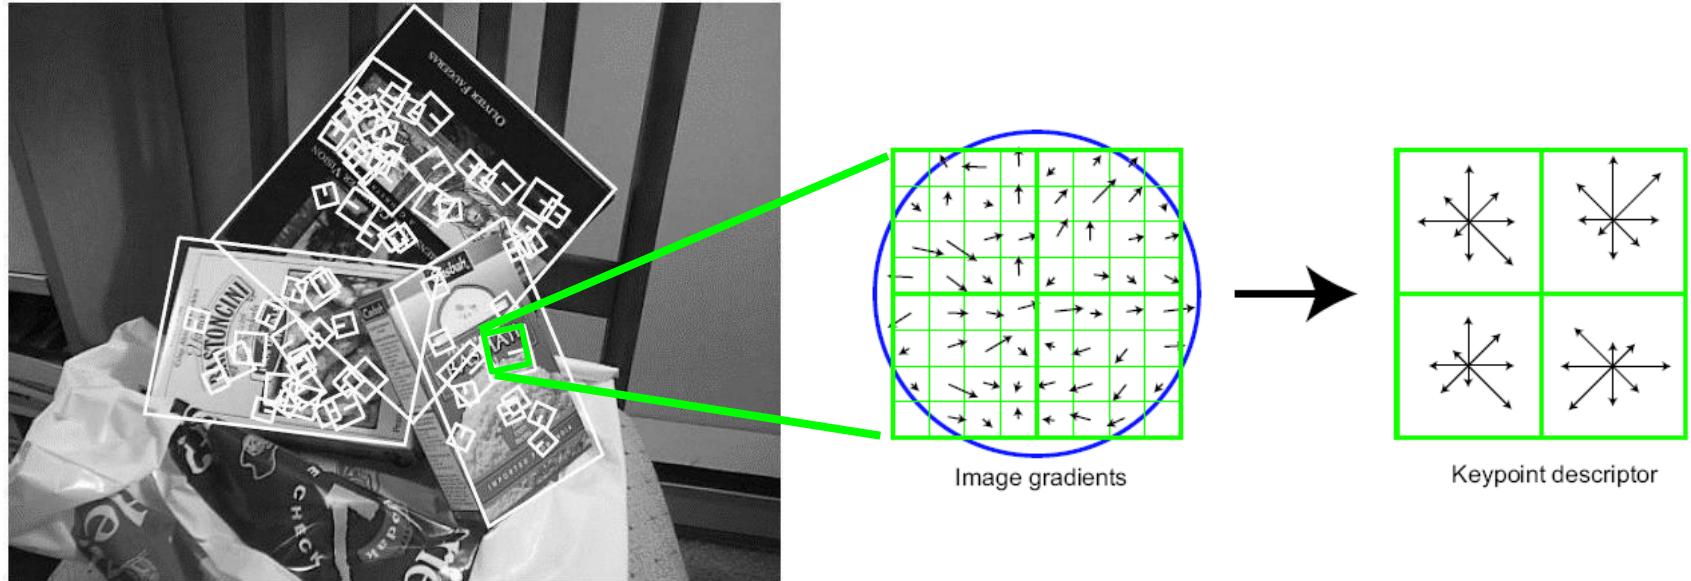
# SIFT descriptor

Full version

- Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
- Compute an orientation histogram for each cell
- 16 cells \* 8 orientations = 128 dimensional descriptor



# Local Descriptors: SIFT Descriptor



**Histogram of oriented  
gradients**

- Captures important texture information
- Robust to small translations / affine deformations

[Lowe, ICCV 1999]

# Feature matching

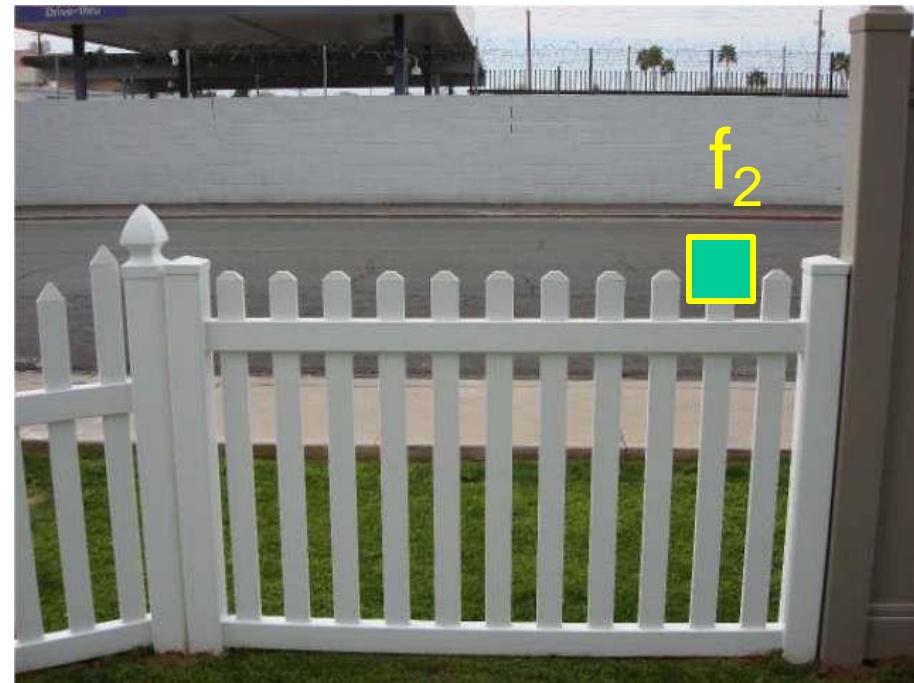
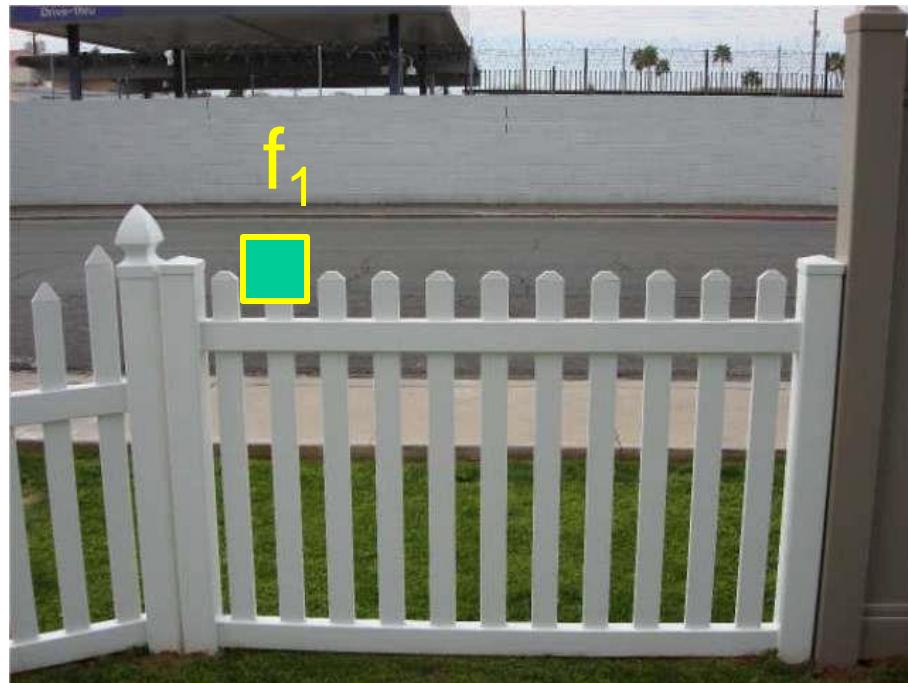
Given a feature in  $I_1$ , how to find the best match in  $I_2$ ?

1. Define distance function that compares two descriptors
2. Test all the features in  $I_2$ , find the one with min distance

# Feature distance

How to define the difference between two features  $f_1, f_2$ ?

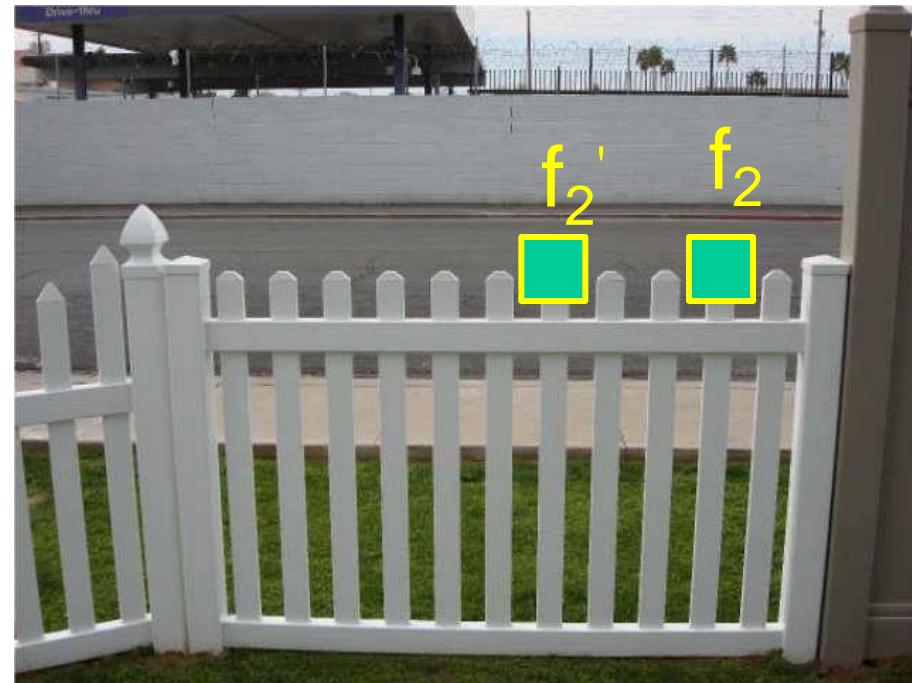
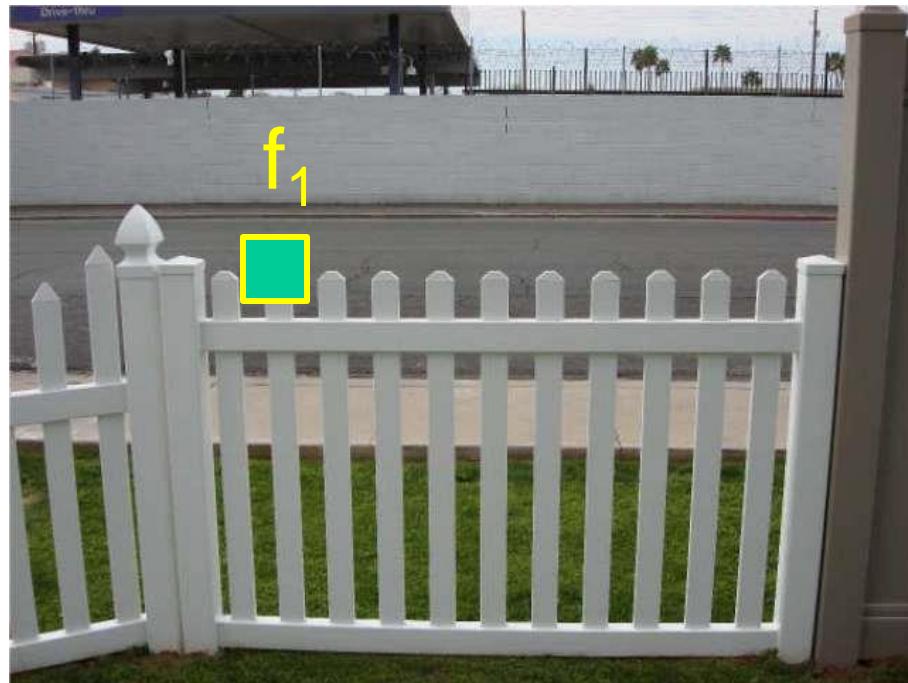
- Simple approach:  $L_2$  distance,  $\|f_1 - f_2\|$
- can give good scores to ambiguous (incorrect) matches



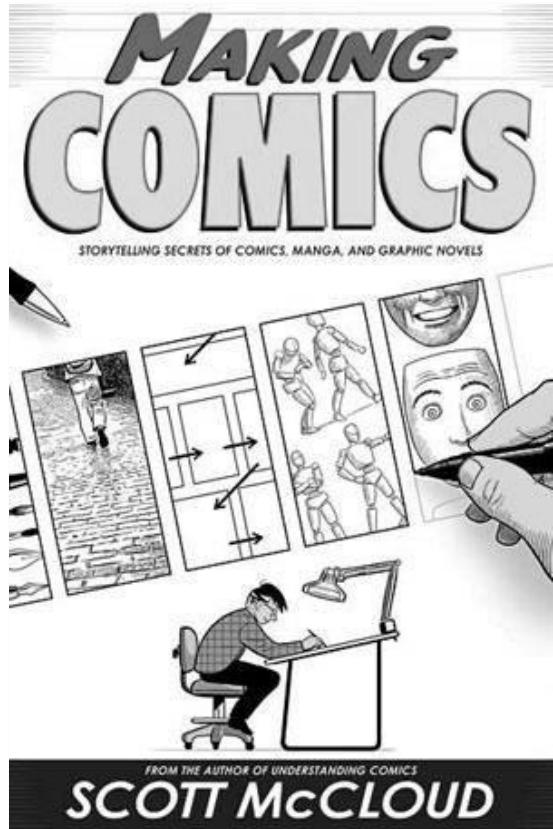
# Feature distance

How to define the difference between two features  $f_1, f_2$ ?

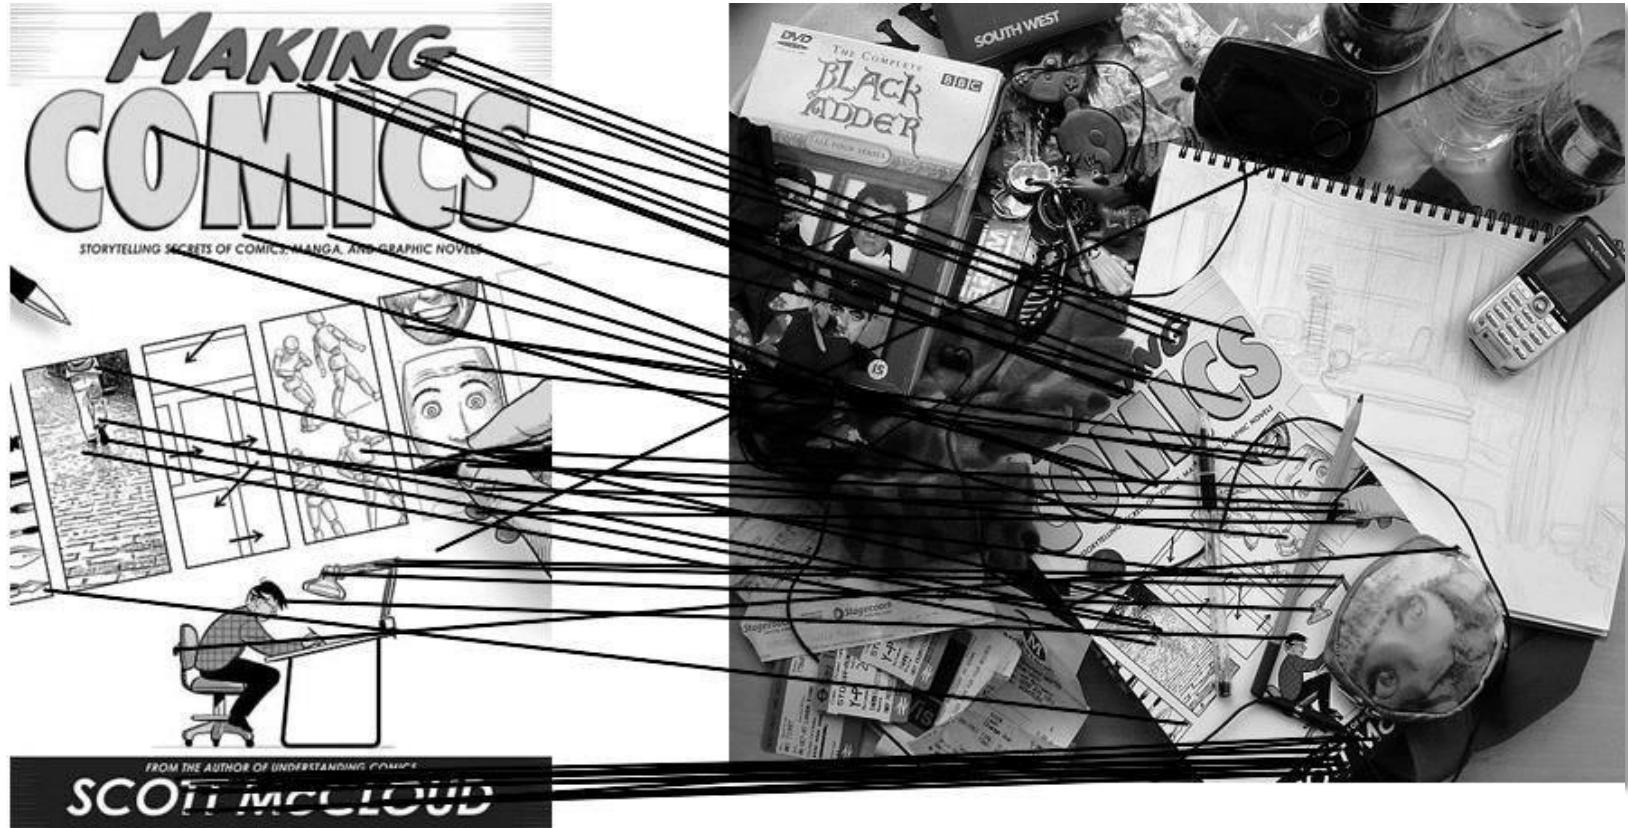
- Better approach: ratio distance =  $\|f_1 - f_2\| / \|f_1 - f_2'\|$ 
  - $f_2$  is best SSD match to  $f_1$  in  $I_2$
  - $f_2'$  is 2<sup>nd</sup> best SSD match to  $f_1$  in  $I_2$
  - gives large values for ambiguous matches



# Feature matching example

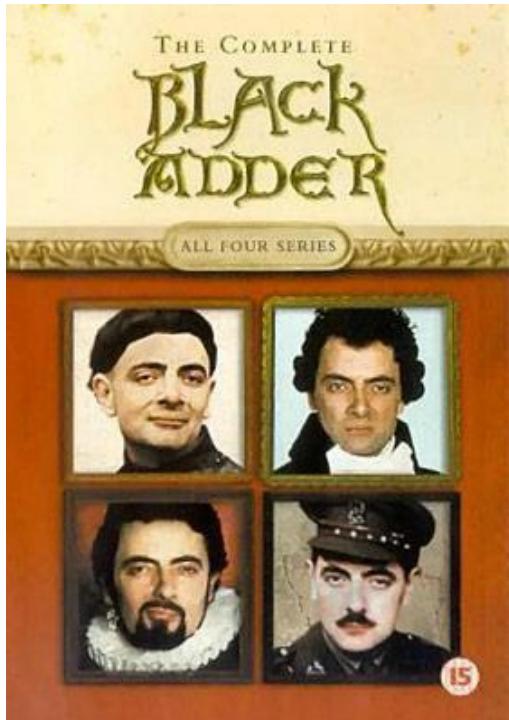


# Feature matching example

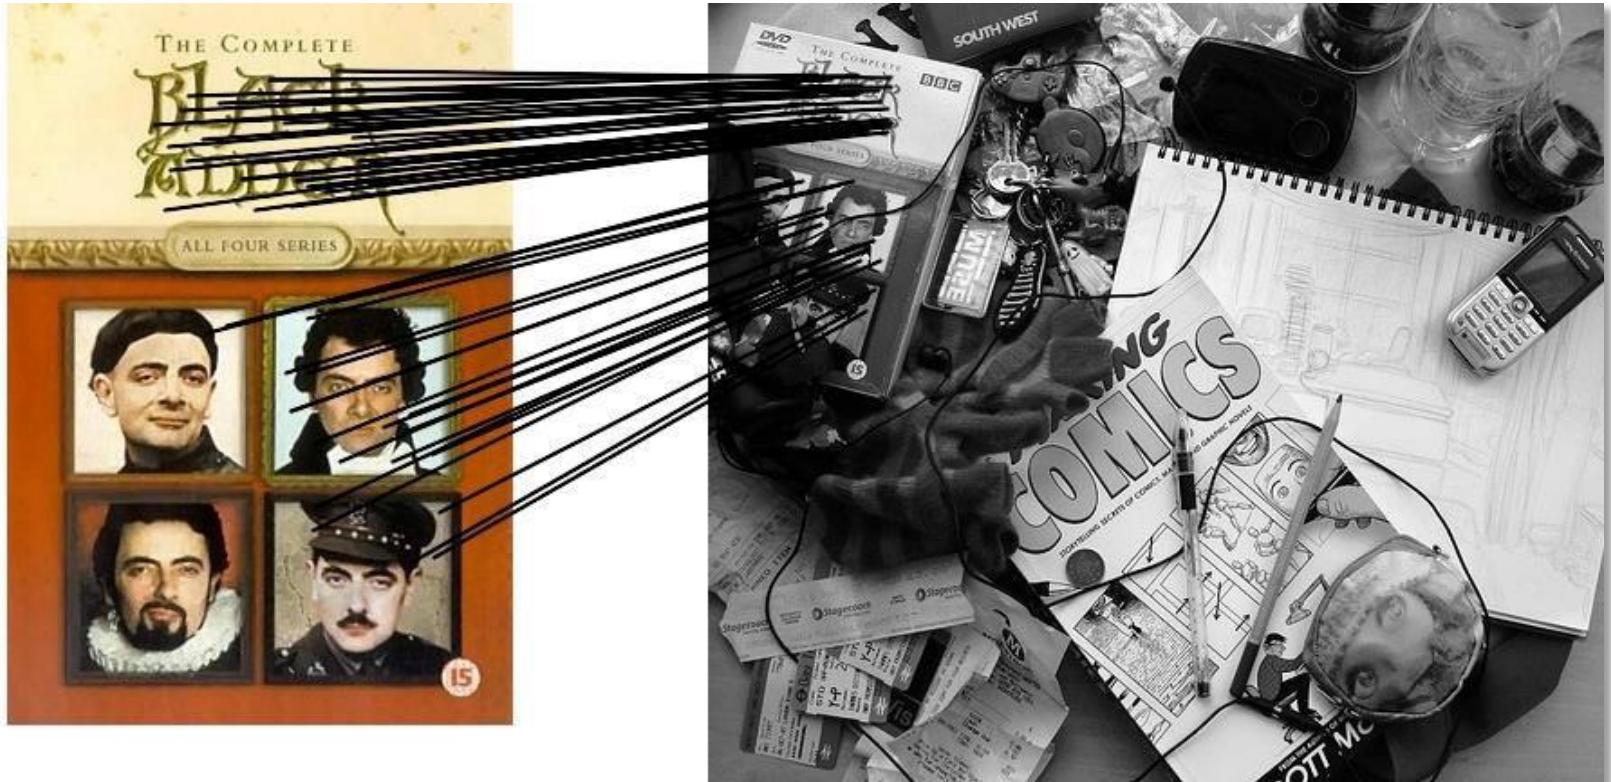


51 matches

# Feature matching example



# Feature matching example



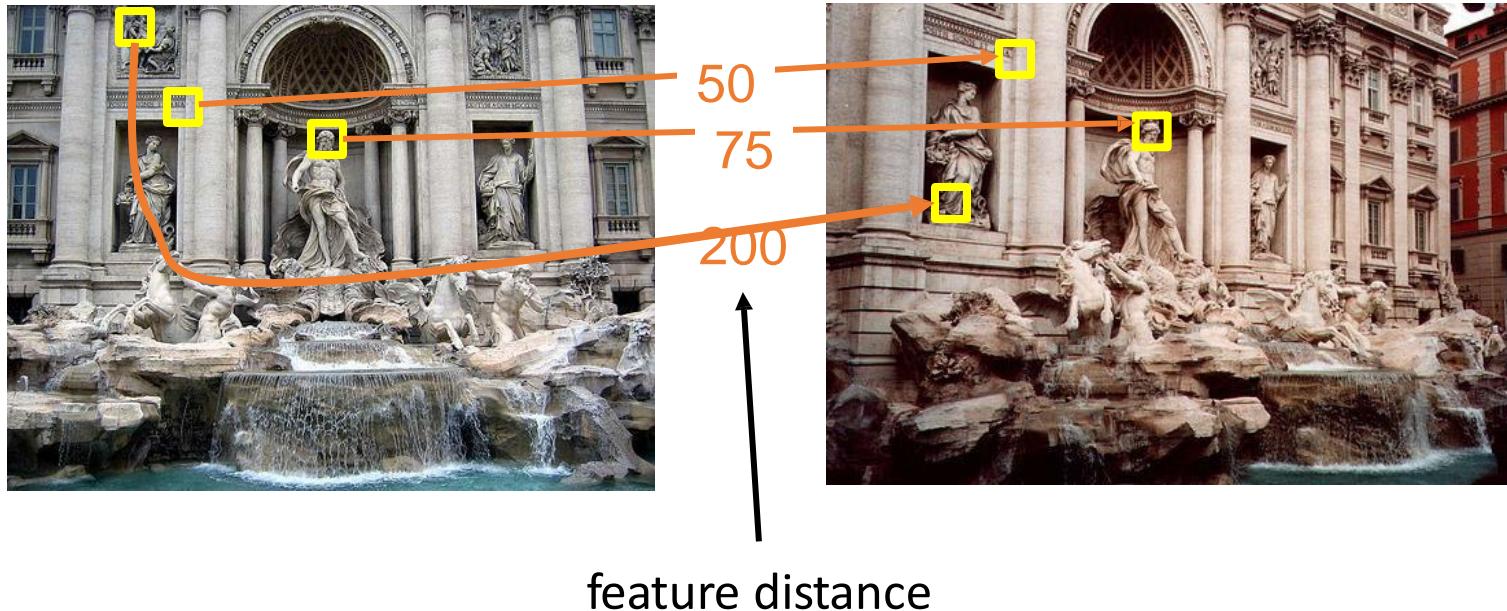
58 matches

# Evaluating the results

How can we measure the performance of a feature matcher?

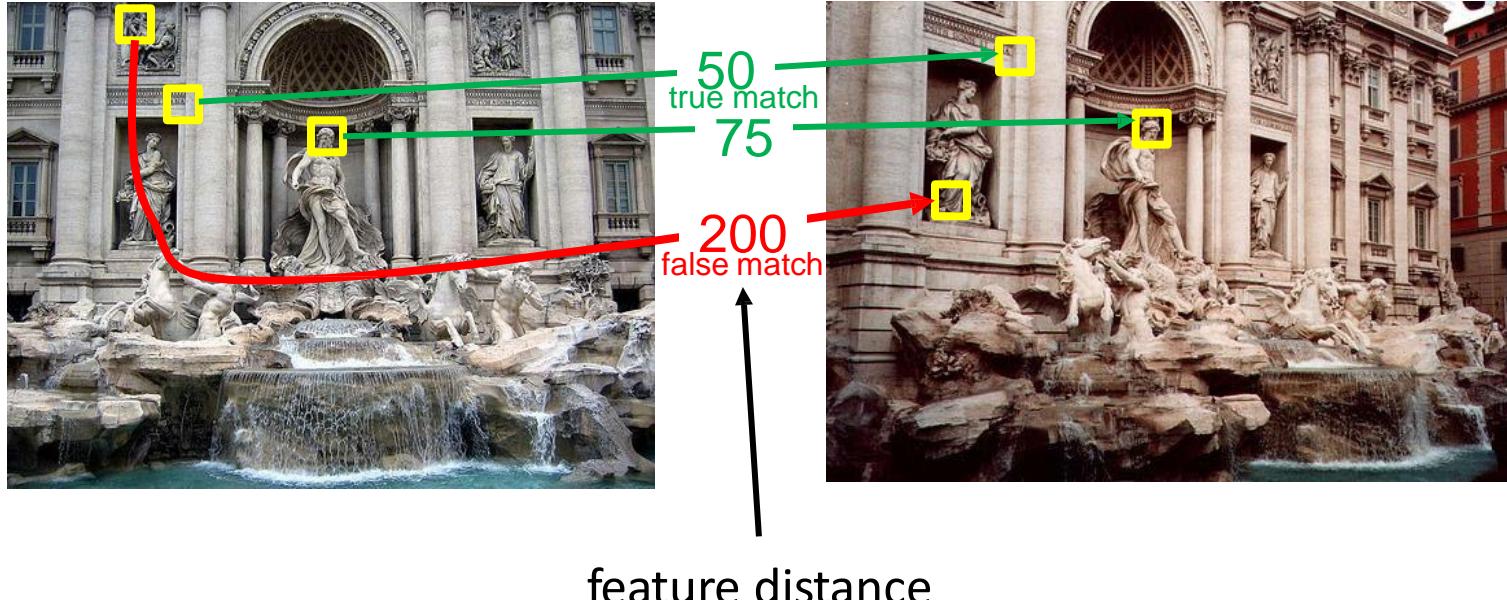
# Evaluating the results

How can we measure the performance of a feature matcher?



# True/false positives

How can we measure the performance of a feature matcher?



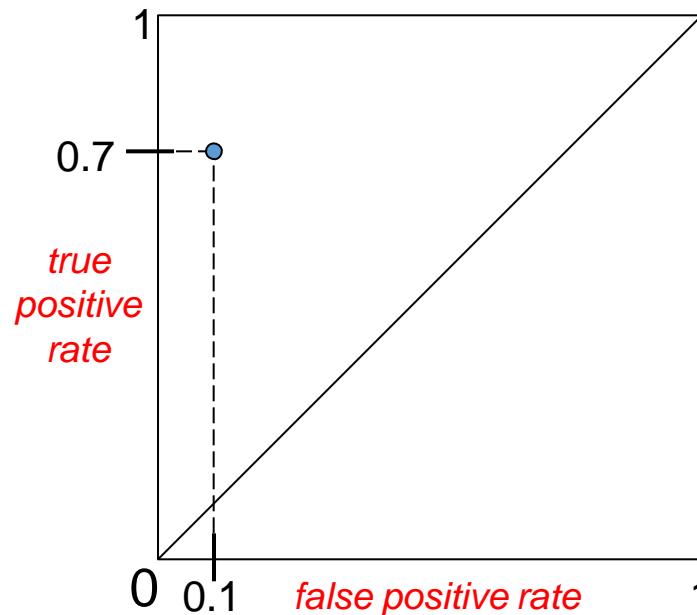
The distance threshold affects performance

- True positives = # of detected matches that are correct
  - Suppose we want to maximize these—how to choose threshold?
- False positives = # of detected matches that are incorrect
  - Suppose we want to minimize these—how to choose threshold?

# Evaluating the results

How can we measure the performance of a feature matcher?

$$\frac{\# \text{ true positives}}{\# \text{ correctly matched features} \text{ (positives)}} \text{ "recall"}$$

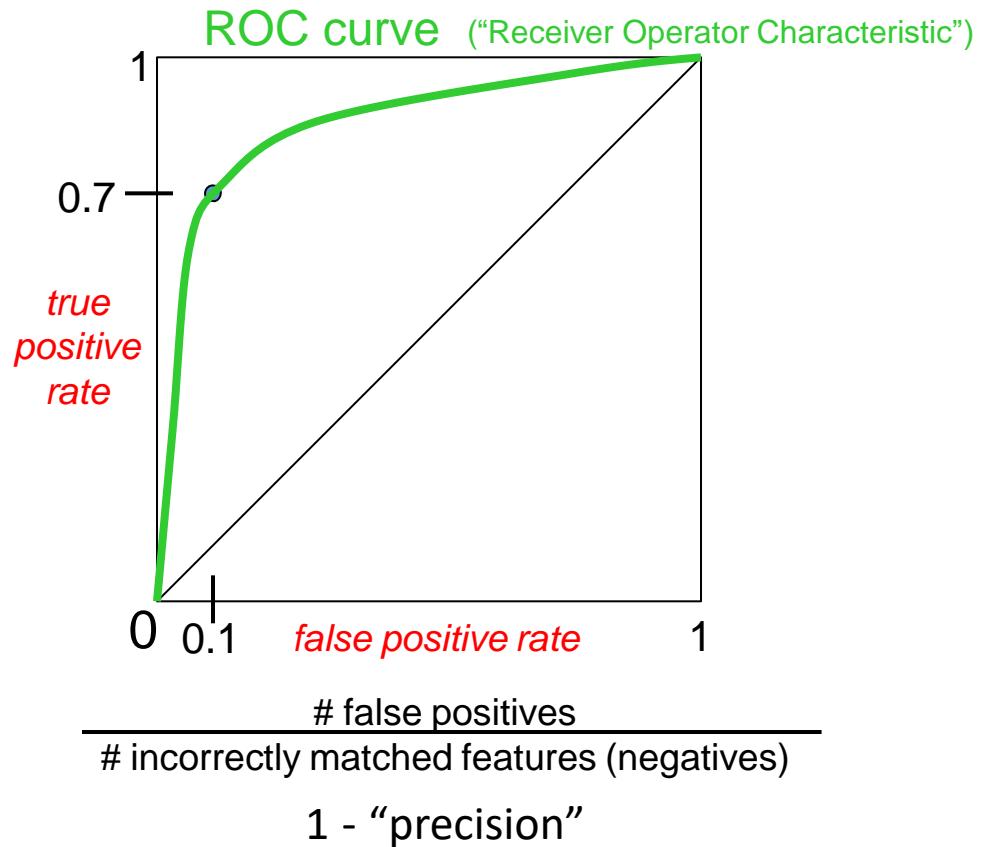


$$\frac{\# \text{ false positives}}{\# \text{ incorrectly matched features} \text{ (negatives)}} \text{ } 1 - \text{"precision"}$$

# Evaluating the results

How can we measure the performance of a feature matcher?

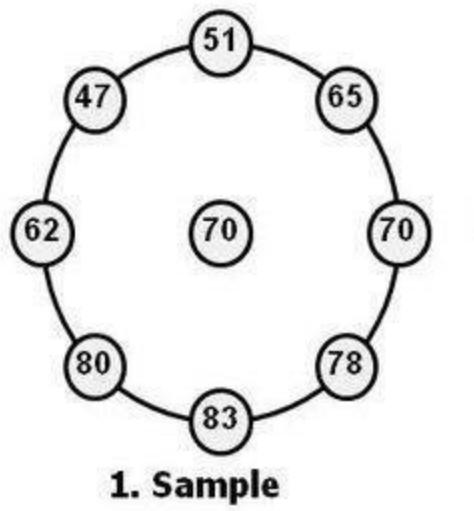
$$\frac{\# \text{ true positives}}{\# \text{ correctly matched features} \text{ (positives)}} \text{ "recall"}$$



# Evaluating the results

- ROC Curves summarize the trade-off between the true positive rate and false positive rate for a predictive model using different probability thresholds.
- Precision-Recall curves summarize the trade-off between the true positive rate and the positive predictive value for a predictive model using different probability thresholds.
- ROC curves are appropriate when the observations are balanced between each class, whereas precision-recall curves are appropriate for imbalanced datasets.

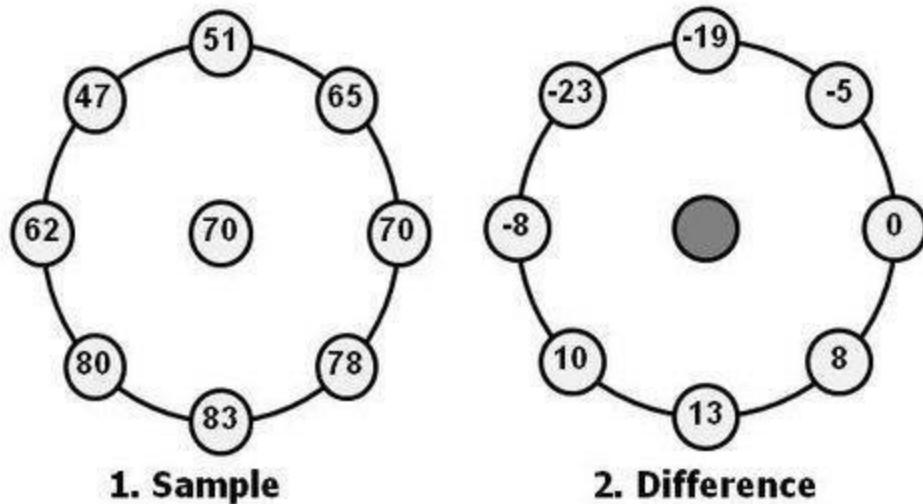
# Local Binary Pattern (LBP)



*Image Source: scholarpedia*

[IEEE TPAMI 2002]

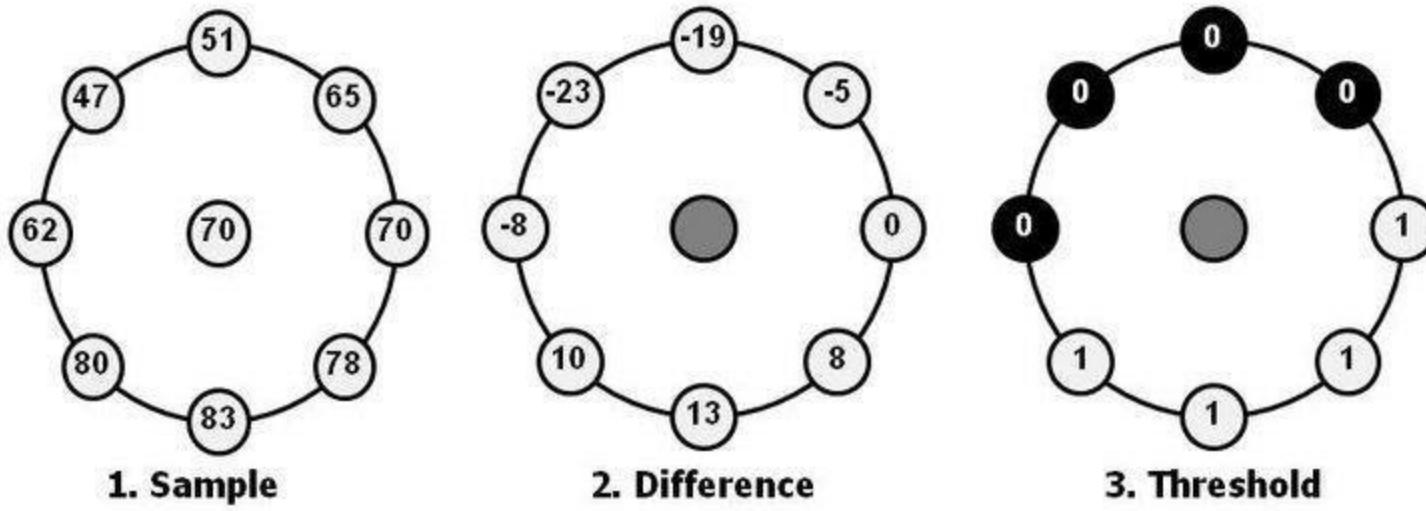
# Local Binary Pattern (LBP)



*Image Source: scholarpedia*

[IEEE TPAMI 2002]

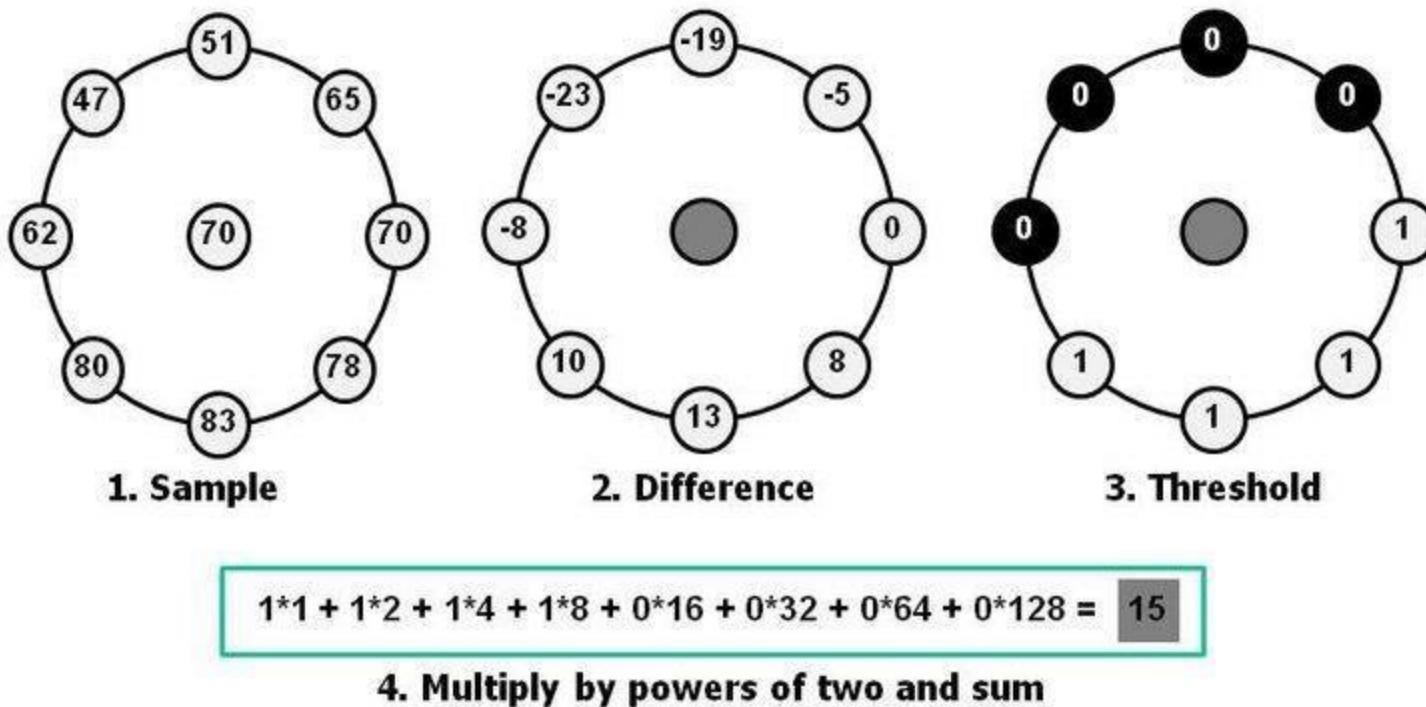
# Local Binary Pattern (LBP)



*Image Source: scholarpedia*

[IEEE TPAMI 2002]

# Local Binary Pattern (LBP)



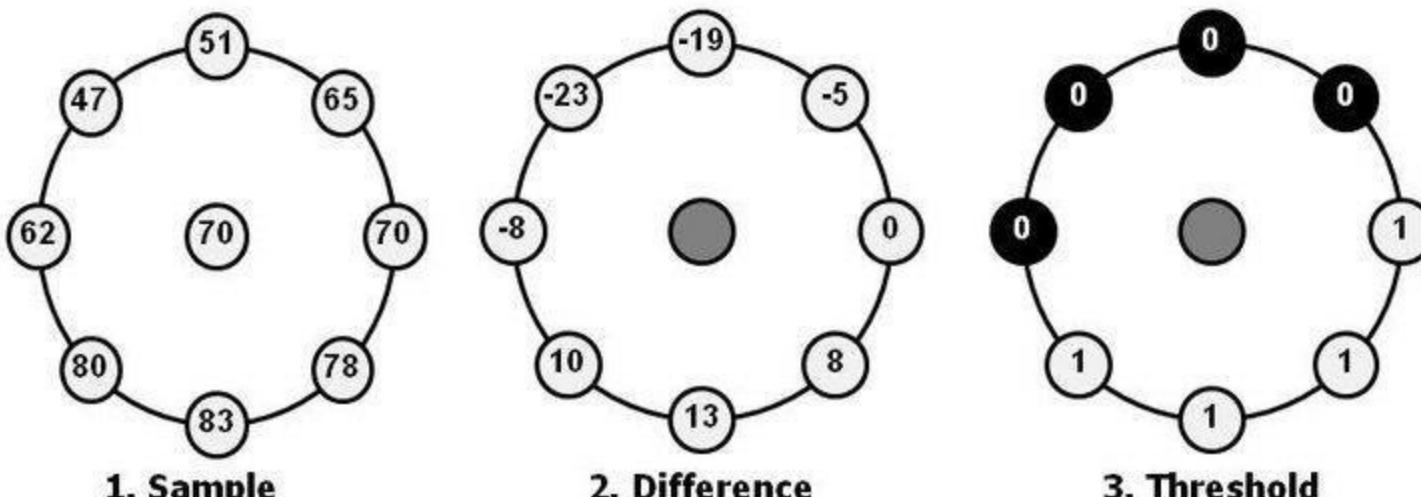
*Image Source: scholarpedia*

[IEEE TPAMI 2002]

# Local Binary Pattern (LBP)

The value of the LBP code of a pixel  $(x_c, y_c)$  is given by:

$$LBP_{P,R} = \sum_{p=0}^{P-1} s(g_p - g_c)2^p \quad s(x) = \begin{cases} 1, & \text{if } x \geq 0; \\ 0, & \text{otherwise.} \end{cases}$$



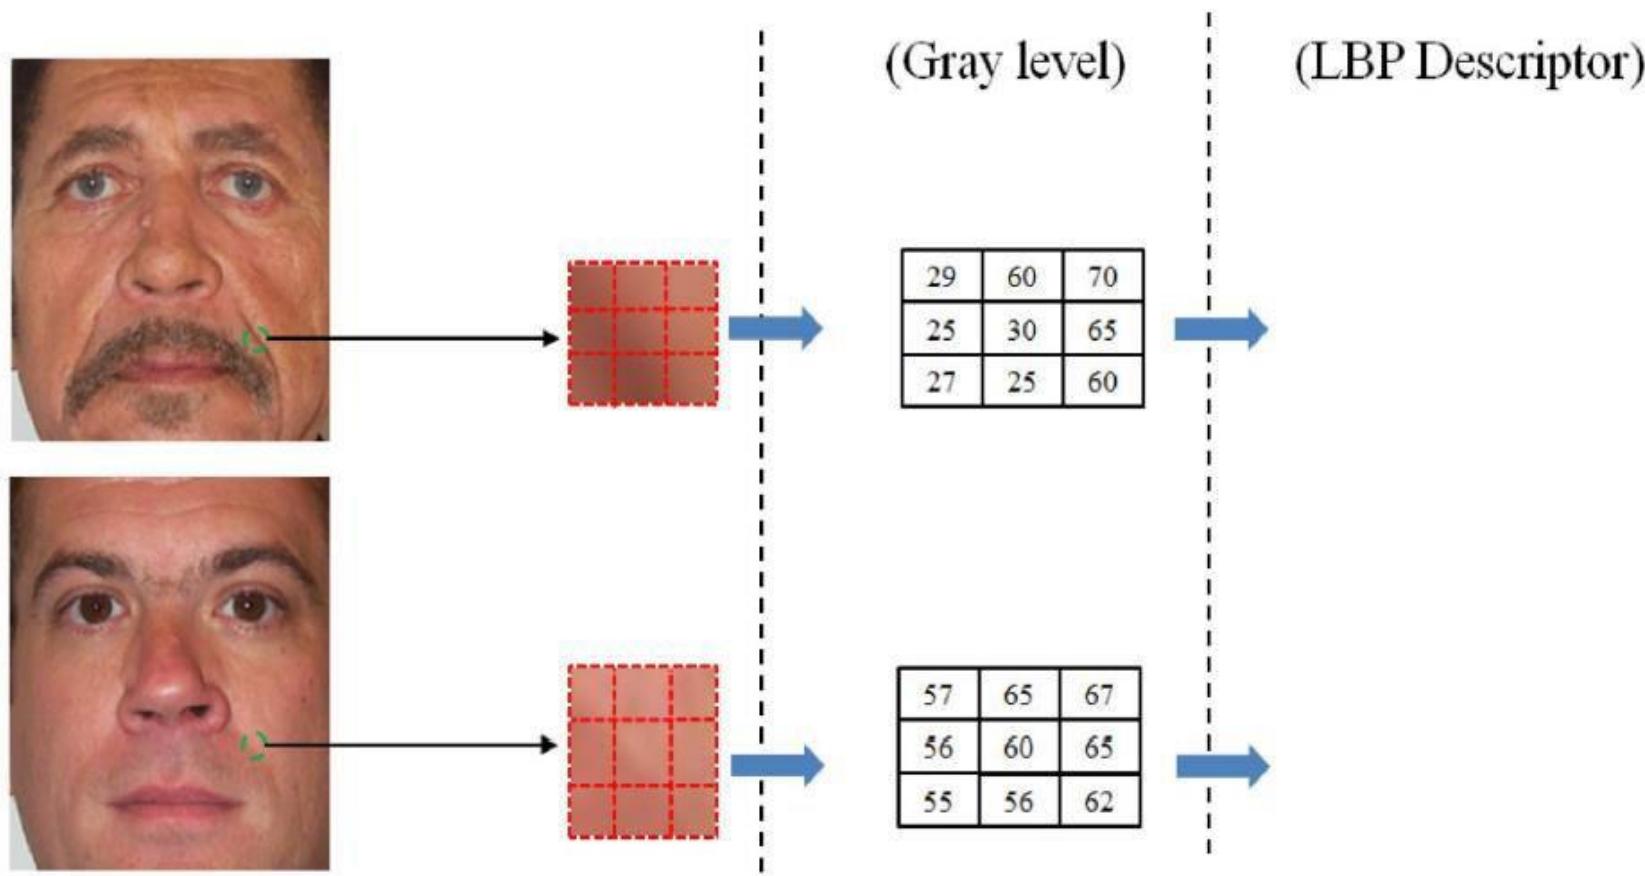
$$1*1 + 1*2 + 1*4 + 1*8 + 0*16 + 0*32 + 0*64 + 0*128 = 15$$

**4. Multiply by powers of two and sum**

*Image Source: scholarpedia*

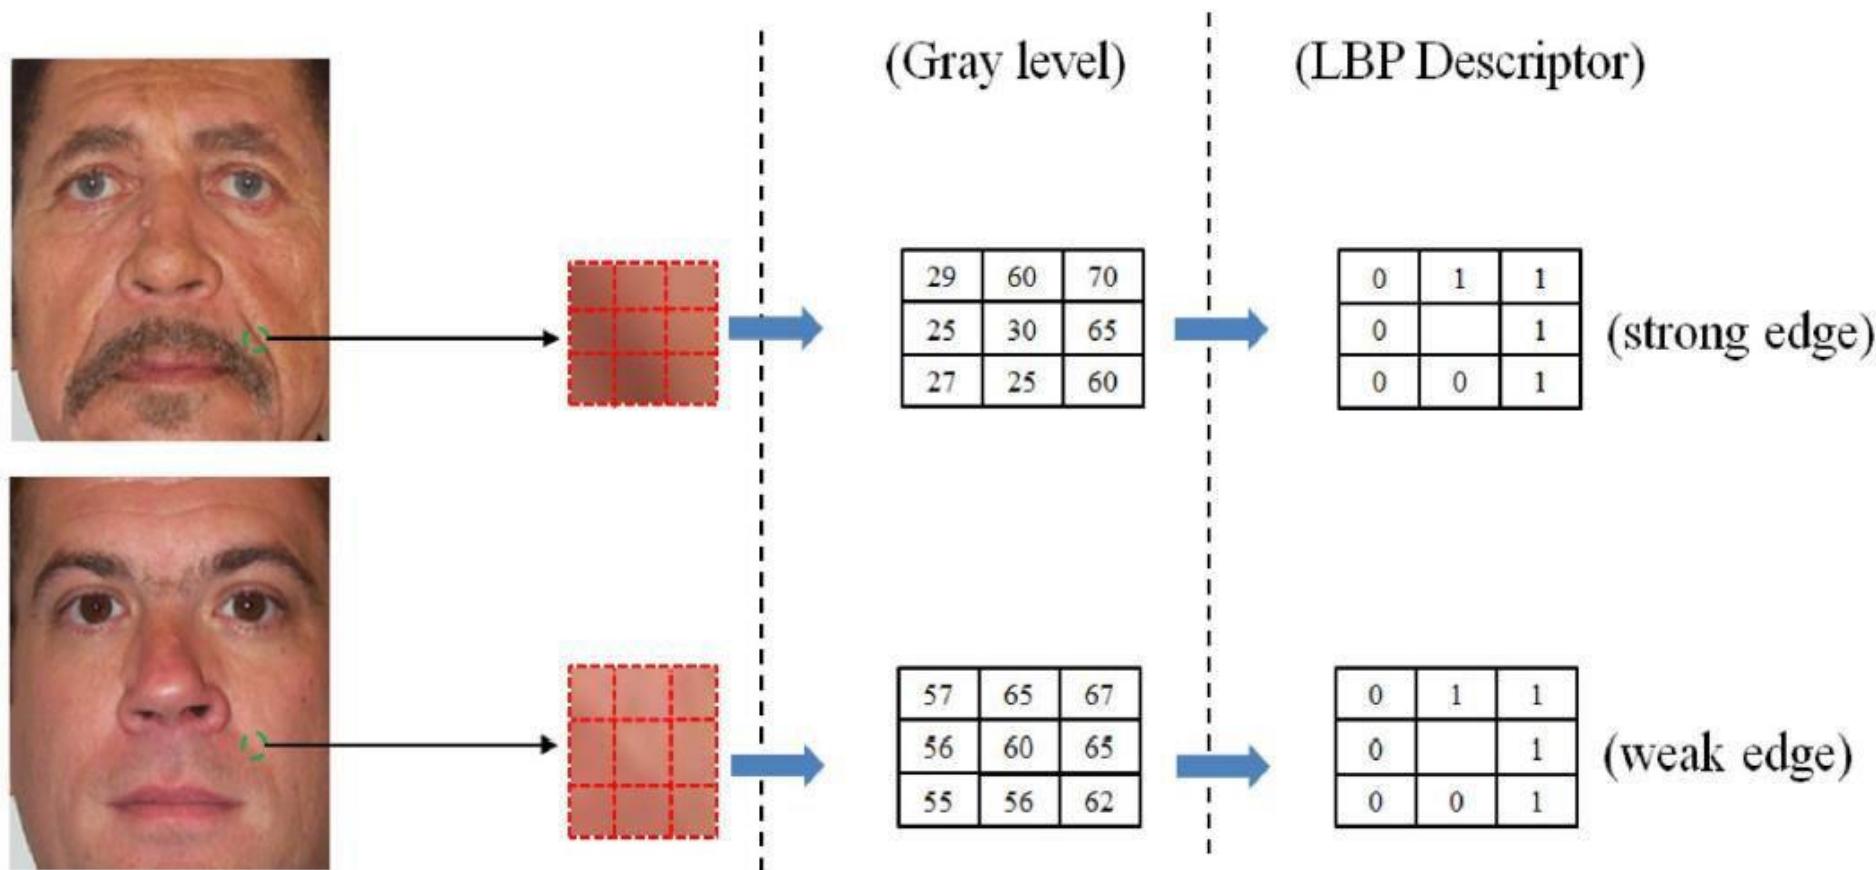
[IEEE TPAMI 2002]

# Local Binary Pattern (LBP)



*Image Source:* Nguyen, D. T., Cho, S. R., & Park, K. R. (2015). Age Estimation-Based Soft Biometrics Considering Optical Blurring Based on Symmetrical Sub-Blocks for MLBP. *Symmetry*, 7(4), 1882-1913.

# Local Binary Pattern (LBP)



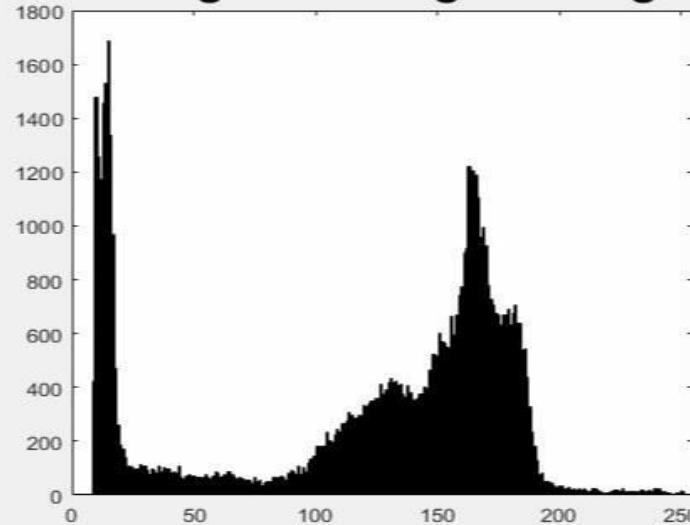
*Image Source: Nguyen, D. T., Cho, S. R., & Park, K. R. (2015). Age Estimation-Based Soft Biometrics Considering Optical Blurring Based on Symmetrical Sub-Blocks for MLBP. *Symmetry*, 7(4), 1882-1913.*

# Local Binary Pattern (LBP)

Original Grayscale Image



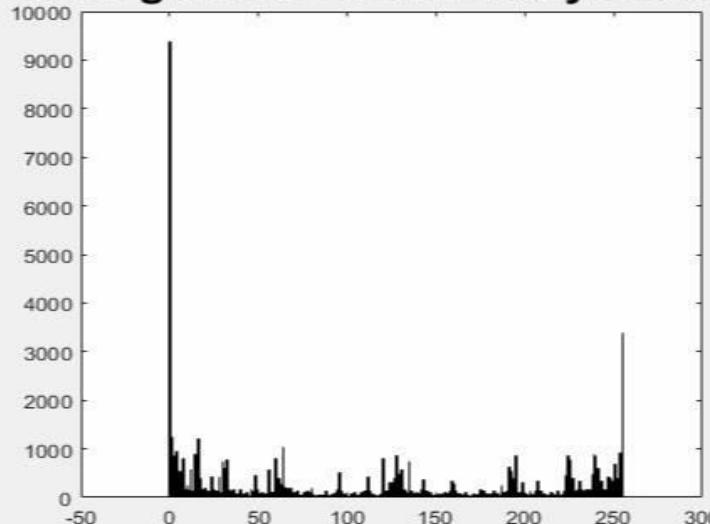
Histogram of original image



Local Binary Pattern



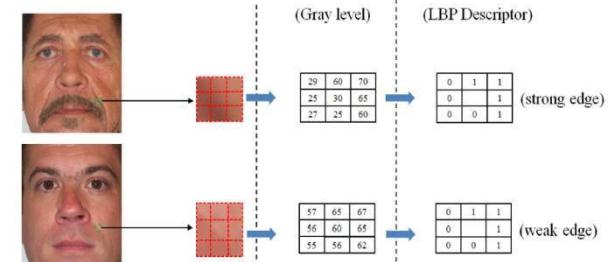
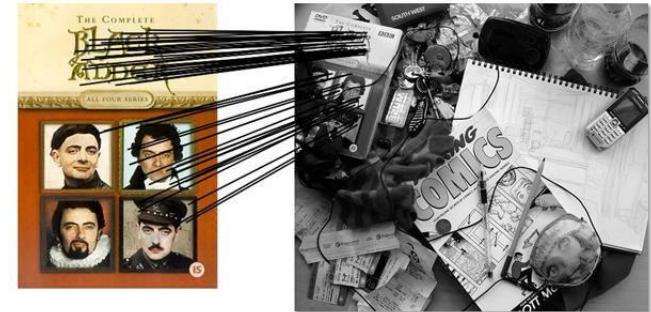
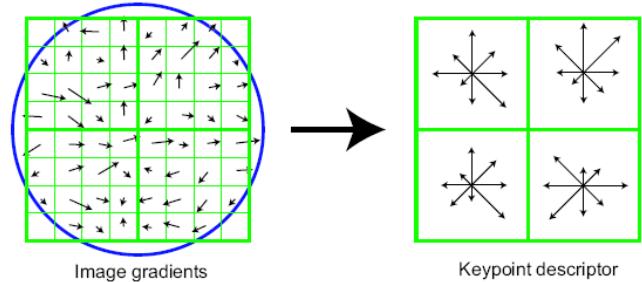
Histogram of Local Binary Pattern



*Image Source:  
Mathworks*

# Things to remember

- Region detection: repeatable and distinctive
  - Blobs, DoG, stable regions
  - Invariant to affine transformation
- Local Descriptors:
  - Discriminative
  - Robust
  - Compact
  - Efficient
- Scale Invariant Feature Transform
  - Utilizes gradient information
  - Gradient direction binning
- Local Binary Pattern
  - Exploits local relationship
  - Captures edge, spot, corner, etc.

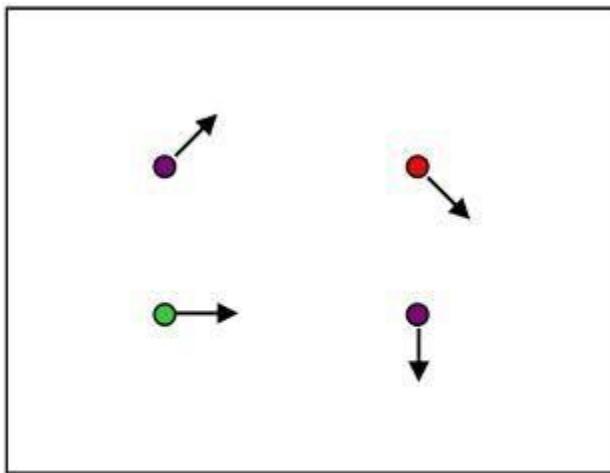


# Acknowledgements

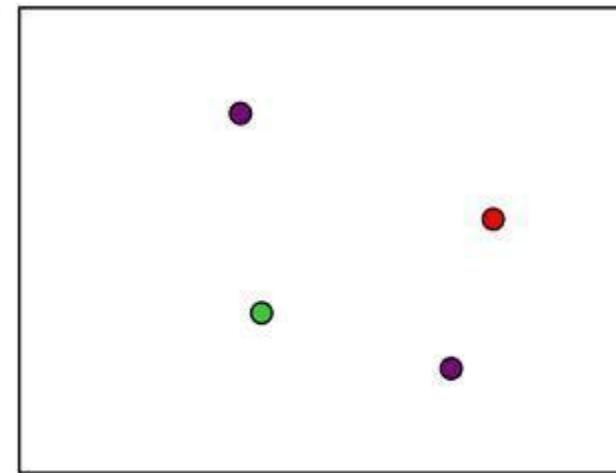
- Thanks to the following researchers for making their teaching/research material online
  - Forsyth
  - Steve Seitz
  - Noah Snavely
  - J.B. Huang
  - Derek Hoiem
  - D. Lowe
  - A. Bobick
  - S. Lazebnik
  - K. Grauman
  - R. Zaleski
  - Leibe

Thank you

- Next class: feature tracking and optical flow



$$I(x,y,t)$$



$$I(x,y,t+1)$$