

Spring 2023



Distributed Computing

- Token-Based Mutual Exclusion algos



Dr. Rajendra Prasath

Indian Institute of Information Technology Sri City, Chittoor



23rd March 2023 (<http://rajendra.2power3.com>)

> Distributed Computing?

- How will you design a Distributed Algorithm?



- Learn to Solve using Distributed Algorithms

Recap: Distributed Systems

A Distributed System:

- A collection of independent systems that appears to its users as a single coherent system
- A system in which hardware and software components of networked computers communicate and coordinate their activity only by passing messages
- A computing platform built with many computers that:
 - Operate concurrently
 - Are physically distributed (have their own failure modes)
 - Are linked by a network
 - Have independent clocks

Recap: Characteristics

- **Concurrent execution of processes:**
 - Non-determinism, Race Conditions, Synchronization, Deadlocks, and so on
- **No global clock**
 - Coordination is done by message exchange
 - No Single Global notion of the correct time
- **No global state**
 - No Process has a knowledge of the current global state of the system
- **Units may fail independently**
 - Network Faults may isolate computers that are still running
 - System Failures may not be immediately known

What did you learn so far?

- Goals / Challenges in Message Passing systems
- Distributed Sorting / Space-Time diagram
- Partial Ordering / Total Ordering
- Concurrent Events / Causal Ordering
- Logical Clocks vs Physical Clocks
- Global Snapshot Detection
- Termination Detection Algorithm
- Leader Election in Rings
- Topology Abstraction and Overlays
- Message Ordering and Group Communication
- Mutual Exclusion Algorithm

[Now] → → →

> About this **Lecture**

What do we learn today?

- **Mutual Exclusion Algorithms**
 - Centralized Algorithm
 - Token-Based / Permission-Based Algorithms
 - Quorum-Based Algorithm
 - Tree-Based Algorithm

Let us **explore** these topics ➔ ➔ ➔

Distributed Mutual Exclusion – Token-Based MutEx Algorithms



Recap: The need for Mutex?

→ Mutual Exclusion

→ Operating systems: Semaphores

→ In a single machine, you could use semaphores to implement mutual exclusion

→ How to implement semaphores?

→ Inhibit interrupts

→ Use clever instructions (e.g. test-and-set)

→ On a multiprocessor shared memory machine, only the latter works



Characteristics

- Processes communicate only through messages
 - no shared memory or no global clocks
- Processes must expect unpredictable message but finite delays
- Processes coordinate access to shared resources that should only be used in a mutually exclusive manner.

Recap: Distributed MutEx

- **No Deadlocks** - no set of sites should be permanently blocked, waiting for messages from other sites in that set
- **No starvation** - no site should have to wait indefinitely to enter its critical section, while other sites are executing the CS more than once
- **Fairness** - requests honored in the order they are made. This means processes have to be able to agree on the order of events. (Fairness prevents starvation.)
- **Fault Tolerance** - the algorithm is able to survive a failure at one or more sites

Recap: Performance Metrics



Token-Based MutEx Algorithms

Token-based solution: Processes share a special message known as a **token**

- **Token holder has right to access shared resource**
- **Wait for/ask for (depending on algorithm) token; enter Critical Section (CS) when it is obtained, pass to another process on exit or hold until requested (depending on algorithm)**
- **If a process receives the token and doesn't need it, just pass it on**

Distributed MutEx - A Few Issues

- Who can access the resource?
- When does a process to be privileged to access the resource?
- How long does a process access the resource?
Any finite duration?
- How long can a process wait to be privileged?
- Computation complexity of the solution

Types of Distributed MutEx

- ➔ Token-based distributed mutual exclusion algorithms
 - ➔ Suzuki - Kasami's Algorithm
 - ➔ Feuerstein et al's Algorithm
- ➔ Non-token based distributed mutual exclusion algorithms
 - ➔ Lamport's Algorithm
 - ➔ Ricart-Agrawala's Algorithm

Token Based Methods

Advantages:

- Starvation can be avoided by efficient organization of the processes
- Deadlock is also avoidable

Disadvantage: Token Loss

- Must initiate a cooperative procedure to recreate the token
- Must ensure that only one token is created!

Token Ring Approach

- Processes are organized in a logical ring: P_i has a communication channel to $P_{(i+1) \bmod N}$

Operations:


- Only the process holding the token T can enter the CS
- To enter the critical section, wait passively for T
When in CS, hold on to T and don't release it
- To exit the CS, send T onto your neighbor
- If a process does not want to enter the CS when it receives T, it simply forwards T to the next neighbor

Token Ring - Idea

- Single token circulates, enter CS when token is present
- Mutual exclusion obvious
- Algorithms differ in how to find and get the token
- Uses **sequence numbers** rather than timestamps to differentiate between old and current requests

Token Rings - Illustration

➔ Request movements in an unidirectional ring network



Suzuki - Kasami's Algorithm

- Broadcast a request for the token
- Process with the token sends it to the requestor if it does not need it
- Issues:
 - Current versus outdated requests
 - Determining sites with pending requests
 - Deciding which site to give the token to

Data Structures

The token:

- Queue (FIFO) Q of requesting processes
- $LN[1..n]$: sequence number of request that j executed most recently

The request message:

- $REQUEST(i, k)$: request message from node i for its k^{th} critical section execution

Other data structures:

- $RN_i[1..n]$ for each node i , where $RN_i[j]$ is the largest sequence number received so far by i in a $REQUEST$ message from j

Suzuki-Kasami's algorithm

To request critical section:

- If i does not have token, increment $RN_i[i]$ and send $REQUEST(i, RN_i[i])$ to all nodes
- If i has token already, enter critical section if the token is idle (no pending requests), else follow rule to release critical section

On receiving $REQUEST(i, s_n)$ at j :

- Set $RN_j[i] = \max(RN_j[i], s_n)$
- If j has the token and the token is idle then
 - send it to i if $RN_j[i] = LN[i] + 1$
 - If token is not idle, follow rule to release critical section

Suzuki-Kasami's algorithm

To enter critical section:

→ Enter CS if token is present

To release critical section:

→ Set $LN[i] = RN_i[i]$

→ For every node j which is not in Q (in token),
add node j to Q if $RN_i[j] = LN[j] + 1$

→ If Q is non empty after the above, delete first
node from Q and send the token to that node

Complexity

→ No. of messages:

- 0 if node holds the token already,
- n otherwise

→ Synchronization delay:

- 0 (node has the token) or
- max. message delay (token is elsewhere)

→ No starvation

Token Based Control in Rings

- Requests move in either Clockwise or Anticlockwise
- Proposed by Feuerstein et al. (1996)
- There are 3 steps:
 - P needs the resource
 - P has T: enter CS
 - P has no T: send the request to the next P
 - P receives a request
 - P has T: increase TC by 1 and send T to the next P
 - P has no T: send request to the next P
 - P receives Token
 - P has a pending request: enter CS and decrease TC by 1 and send T to next P if $TC > 0$
 - P has no pending request: send T to the next P

Refer to: Feuerstein et al., Efficient token-based control in rings, Information Processing Letters, 66 (4) (1998) pp. 175-180



Token Based Control In Rings

- When p needs the resource then
 - If p has T then it enters the critical section.
 - If p has not T then it sends a request message to the next processor
- When p receives a request message then:
 - If p has T then it increases the counter and passes T to the next processor.
 - If p has not T then the request message is passed to the next processor.
- When p receives T then:
 - If p has a pending request then it enters the critical section in which the token counter is decreased by 1; at the exit of the critical section p passes T to the next processor if the token counter is greater than 0
 - If p has not a pending request then T is passed to the next processor in the ring.

Read from: E. Feuerstein et al. Efficient token-based control in rings, Information Processing Letters, 66 (1998) pp. 175-180, doi: 10.1016/S0020-0190(98)00054-4

Token Rings - Features

- Safety & Liveness are guaranteed
- Ordering is not guaranteed
- Bandwidth: 1 message per exit
- Client delay: 0 to N message transmissions
- Synchronization delay between one process's exit from the CS and the next process's entry is between 1 and $N-1$ message transmissions

Non-Token Based Methods

- **Permission-based solutions:** a process that wishes to access a shared resource must first get permission from one or more other processes.
- Avoids the problems of token-based solutions, but is more complicated to implement

Non-Token Based Algorithms

→ Notations:

→ P_i : i^{th} Process

→ R_i : Request set, containing IDs of all P_i s from which permission must be received before accessing CS

→ Non-token based approaches use **time stamps** to order requests for CS

→ **Smaller time stamps get priority over larger ones**

→ Lamport's Algorithm

→ $R_i = \{P_1, P_2, \dots, P_n\}$, i.e., all processes.

→ Request queue: maintained at each P_i ordered by time stamps.

→ **Assumption:** message delivered in FIFO

Lamport's Algorithm

→ Requesting CS:

- Send REQUEST(ts_i, i) where (ts_i, i) - Request time stamp; Place REQUEST in $request_queue_i$
- On receiving the message; P_j sends time-stamped REPLY message to P_i ; P_i 's request placed in $request_queue_j$

→ Executing CS:

- P_i has received a message with time stamp larger than (ts_i, i) from all other sites
- P_i 's request is the top most one in $request_queue_i$

→ Releasing CS:

- Exiting CS: send a time stamped RELEASE message to all sites in its request set
- Receiving RELEASE message: P_j removes P_i 's request from its queue

Notable Points

- Purpose of REPLY messages from i to j is to ensure that j knows of all requests of i prior to sending the REPLY (possibly any request of i with timestamp lower than j 's request)
- Requires FIFO channels
- $3(n - 1)$ messages per critical section invocation
- Synchronization delay = max msg transmission time
- Requests are granted in order of increasing timestamps

Performance Improvements

- $3(n-1)$ messages per Critical Section invocation
 - $(n - 1)$ REQUEST messages
 - $(n - 1)$ REPLY messages
 - $(n - 1)$ RELEASE messages
- Synchronization delay: T
- **Optimization:**
 - **Suppress reply messages:** For example, P_j receives a REQUEST message from P_i after sending its own REQUEST message with time stamp higher than that of P_i 's then Do NOT send a REPLY message
 - Messages reduced to between $2(n-1)$ and $3(n-1)$

Ricart & Agrawala's Algorithm

- A time-stamp based approach
- Originally proposed by Lamport using logical clocks
- Modified by Ricart & Agrawala

Ricart & Agrawala's Algorithm

Main Idea:

- Process j need not send a REPLY to Process i if j has a request with timestamp lower than the request of i (since i cannot enter before j here)
- Does not require FIFO
- $2(n - 1)$ messages per critical section invocation
- Synchronization delay = maximum message transmission time
- Requests granted in order of increasing timestamps

Ricart & Agrawala (contd)

- Processes need entry to critical section multicast a request, and can enter it only when all other processes have replied positively
- Messages requesting entry are of the form $\langle T, P_i \rangle$
 - T - sender's timestamp (Lamport clock)
 - P_i the sender's identity

Ricart & Agrawala - Algorithm

To enter the Critical Section (CS):

- Set state = wanted
- multicast "request" to all processes (including timestamp)
- wait until all processes send back "reply"
- change state to held and enter the CS

On receipt of a request $\langle T_j, P_j \rangle$ at P_i :

- if (state == held) or (state == wanted & $(T_i, P_i) < (T_j, P_j)$) then enqueue the request
- else "reply" to P_j

On exiting the CS:

- change state to release and "reply" to all queued requests

Ricart & Agrawala - Simplified

To request Critical Section:

- send timestamped REQUEST message (ts_i, i)

On receiving request (ts_i, i) at j :

- if j is neither requesting nor executing critical section then send REPLY to i
- if j is requesting and i 's request timestamp is smaller than j 's request timestamp then
 - enqueue the request; Otherwise, defer the request

To enter Critical Section:

- Process i enters critical section on receiving REPLY messages from all processes

To release Critical Section:

- send REPLY to all deferred requests

Summary

- Recap: Distributed Mutual Exclusion Algorithms
 - Mutual Exclusion Problem
 - Basics of MutEx algorithms
 - Types of MutEx algorithms
 - Token-based Algorithm
 - Suzuki-Kasami's Algorithm
 - Feuerstein et al's Algorithm
 - Non-Token based Algorithm
 - Lamport's Algorithm
 - Ricart - Agrawala's Algorithm
 - Performance Metrics

Many more to come up ... ! Stay tuned in !!

Penalties



- Every Student is expected to strictly follow a fair Academic Code of Conduct to avoid penalties
- Penalties is heavy for those who involve in:
 - Copy and Pasting the code
 - Plagiarism (copied from your neighbor or friend - in this case, both will get "0" marks for that specific take home assignments)
 - If the candidate is unable to explain his own solution, it would be considered as a "copied case"!!
 - Any other unfair means of completing the assignments

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <https://www.iiits.ac.in/people/regular-faculty/dr-rajendra-prasath/>

(OR)

→ <http://rajendra.2power3.com>

Assistance

- You may post your questions to me at any time
- You may meet me in person on available time or with an appointment
- You may ask for one-to-one meeting

Best Approach

- You may leave me an email any time
(email is the best way to reach me faster)



Questions It's Your Time

