

Spring 2023



Distributed Computing

- Non-Token-Based MutEx algorithms



Dr. Rajendra Prasath

Indian Institute of Information Technology Sri City, Chittoor



27th March 2023 (<http://rajendra.2power3.com>)

> Distributed Computing?

- How will you design a Distributed Algorithm?



- Learn to Solve using Distributed Algorithms

Recap: Distributed Systems

A Distributed System:

- A collection of independent systems that appears to its users as a single coherent system
- A system in which hardware and software components of networked computers communicate and coordinate their activity only by passing messages
- A computing platform built with many computers that:
 - Operate concurrently
 - Are physically distributed (have their own failure modes)
 - Are linked by a network
 - Have independent clocks

Recap: Characteristics

- **Concurrent execution of processes:**
 - Non-determinism, Race Conditions, Synchronization, Deadlocks, and so on
- **No global clock**
 - Coordination is done by message exchange
 - No Single Global notion of the correct time
- **No global state**
 - No Process has a knowledge of the current global state of the system
- **Units may fail independently**
 - Network Faults may isolate computers that are still running
 - System Failures may not be immediately known

What did you learn so far?

- Goals / Challenges in Message Passing systems
- Distributed Sorting / Space-Time diagram
- Partial Ordering / Total Ordering
- Concurrent Events / Causal Ordering
- Logical Clocks vs Physical Clocks
- Global Snapshot Detection
- Termination Detection Algorithm
- Leader Election in Rings
- Topology Abstraction and Overlays
- Message Ordering and Group Communication
- Mutual Exclusion Algorithm

[Now] → → →

> About this **Lecture**

What do we learn today?

- **Mutual Exclusion Algorithms**
 - Centralized Algorithm
 - Token-Based / Permission-Based Algorithms
 - Quorum-Based Algorithm
 - Tree-Based Algorithm

Let us **explore** these topics ➔ ➔ ➔



Distributed Mutual Exclusion – Token / Non-Token-Based Mutex Algorithms



Recap: The need for Mutex?

→ Mutual Exclusion

→ Operating systems: Semaphores

→ In a single machine, you could use semaphores to implement mutual exclusion

→ How to implement semaphores?

→ Inhibit interrupts

→ Use clever instructions (e.g. test-and-set)

→ On a multiprocessor shared memory machine, only the latter works

Characteristics

- Processes communicate only through messages
 - no shared memory or no global clocks
- Processes must expect unpredictable message but finite delays
- Processes coordinate access to shared resources that should only be used in a mutually exclusive manner.

Recap: Distributed MutEx

- **No Deadlocks** - no set of sites should be permanently blocked, waiting for messages from other sites in that set
- **No starvation** - no site should have to wait indefinitely to enter its critical section, while other sites are executing the CS more than once
- **Fairness** - requests honored in the order they are made. This means processes have to be able to agree on the order of events. (Fairness prevents starvation.)
- **Fault Tolerance** - the algorithm is able to survive a failure at one or more sites

Quorum Based algorithms

Why Quorum based algorithm?

- Lamports and Ricard-Agrawala' algorithm requires permission from all processes to enter into the critical section.

Modifications:

- Is it necessary to obtain permission from all processes before entering into the CS?
- How to reduce the message exchanges and increase the performance of Mutex algorithm?

Quorum Based algorithms

What is a Quorum?


- There are n requesting processes in a distributed system and any process may request for CS.
- Can we form such a subset of processes who request for Critical Section? YES !!
- Such a set is said to be a Request Set or Quorum
- In fact, we will have a separate Request set for each process P_i

Quorum - Definition

- A quorum system is a collection of subsets of processes, called quorums, such that each pair of quorums have a non-empty intersection
- How do we formally define a quorum of processes in a distributed system?
- Let us look at some examples


Quorum - Why?

- Process may not respond or may go down (any kind of failure)
- The requesting process can not get REPLY from all remaining processes
- It would infinitely wait for CS !!



Quorum - Why?

→ Can the requesting process get permission from a quorum of processes to enter into CS?



Quorum - Definition

More Formally,

→ Given a set of processes

$$P = \{P_1, P_2, \dots, P_n\}$$

→ A quorum system $Q \subseteq 2^P$ is a set of subsets of P such that

for all Q_1, Q_2 in Q : $Q_1 \cap Q_2 \neq \text{empty}$

→ Each Q_i in Q is called a quorum

Maekawa's Algorithm

- Permission obtained from only a subset of other processes, called the Request Set (or Quorum)
- Separate Request Set R_i , for each process i

Maekawa's Algorithm

Requirements

- For all $i, j: R_i \cap R_j \neq \Phi$
- For all $i: i \in R_i$
- For all $i: |R_i| = K$, for some K
- Any node i is contained in exactly D Request Sets, for some Request set D
- $K = D = \text{sqrt}(N)$ for Maekawa's algorithm

Maekawa's Algorithm - Steps

To Request Critical Section:

→ P_i sends REQUEST message to all process in R_i

On receiving a REQUEST message:

→ Send a REPLY message if no REPLY message has been sent since the last RELEASE message is received.

→ Update status to indicate that a REPLY has been sent.

→ Otherwise, queue up the REQUEST

To enter critical section:

→ P_i enters critical section after receiving REPLY from all nodes in R_i

Maekawa's Algorithm - Steps (contd)

To release critical section:

- Send RELEASE message to all nodes in R_i
- On receiving a RELEASE message, send REPLY to next node in queue and delete the node from the queue.
- If queue is empty, update status to indicate no REPLY message has been sent

Computation Complexity

- Message Complexity: $3 * \sqrt{N}$
- Synchronization delay
 - $2 * (\text{max message transmission time})$
- Major problem: DEADLOCK possible
- Need three more types of messages (FAILED, INQUIRE, YIELD) to handle deadlock.
 - Message complexity can be $5 * \sqrt{N}$
- Important Issue:
 - How to build the request sets?

Raymond's Algorithm

- Forms a directed tree (logical) with the token token-holder as root
- Each node has variable "Holder" that points to its parent on the path to the root.
 - Root's Holder variable points to itself
- Each node P_i has a FIFO request queue Q_i

Raymond's Algorithm

- To request critical section:
 - Send REQUEST to parent on the tree, provided i does not hold the token currently and Q_i is empty. Then place request in Q_i
- When a non-root node j receives a request from k
 - place request in Q_j
 - send REQUEST to parent if no previous REQUEST sent

Raymond's Algorithm (contd)

When the root receives a REQUEST:

- send the token to the requesting node
- set Holder variable to point to that node

When a node receives the token:

- delete first entry from the queue
- send token to that node
- set Holder variable to point to that node
- if queue is non non-empty, send a REQUEST message to the parent (node pointed at by Holder variable)

Raymond's Algorithm (contd)

→ To execute critical section:

- enter if token is received and own entry is at the top of the queue; delete the entry from the queue

→ To release critical section:

- if queue is non non-empty, delete first entry from the queue, send token to that node and make Holder variable point to that node
- If queue is still non non-empty, send a REQUEST message to the parent (node pointed at by Holder variable)

Features of Raymond's Algo

- Average message complexity:
 - $O(\log n)$
- Sync. Delay
 - $(T \log n)/2$, where T = max. message delay

Summary

→ Recap: Distributed Mutual Exclusion Algorithms

→ Mutual Exclusion Problem

→ Basics of MutEx algorithms

→ Types of MutEx algorithms

→ Token-based Algorithms

→ Raymond's Tree based algorithm

→ Non-Token based Algorithms

→ Quorum based algorithm

→ Performance Metrics

Many more to come up ... ! Stay tuned in !!

Penalties



- Every Student is expected to strictly follow a fair Academic Code of Conduct to avoid penalties
- Penalties is heavy for those who involve in:
 - Copy and Pasting the code
 - Plagiarism (copied from your neighbor or friend - in this case, both will get "0" marks for that specific take home assignments)
 - If the candidate is unable to explain his own solution, it would be considered as a "copied case"!!
 - Any other unfair means of completing the assignments

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <https://www.iiits.ac.in/people/regular-faculty/dr-rajendra-prasath/>

(OR)

→ <http://rajendra.2power3.com>

Assistance

- You may post your questions to me at any time
- You may meet me in person on available time or with an appointment
- You may ask for one-to-one meeting

Best Approach

- You may leave me an email any time
(email is the best way to reach me faster)



Questions It's Your Time

