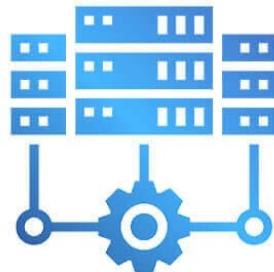


Spring 2023



Distributed Computing

- A Model of Distributed Executions



Dr. Rajendra Prasath

Indian Institute of Information Technology Sri City, Chittoor



> About this Lecture

What do we learn today?

- ▶ This Lecture covers the essential aspects (of Distributed Algorithms / Systems) that every serious programmer needs to know about

- ▶ **The Fundamentals of Distributed Algorithms**
- ▶ **Design principles** and **Analysis**

with an emphasis on certain properties of

- ▶ **The Scalable Application Development**

Let us **explore these topics** → → →

> **Distributed Computing?**

- How will you design a **Distributed Algorithm?**



- Learn to Solve using **Distributed Algorithms**

Recap: What do we learn?

Distributed Computing (DC)

- Core Theoretical Concepts
- Design Principles of DC
- Discrete Events Simulations
- Experimental Evaluations
- Designing Efficient Solution(s) ??
 - To Solve Some Interesting Problems !!

- An Overview of Distributed Computing
- → Simple to advanced?



Recap: Distributed Systems

A Distributed System:

- A collection of independent systems that appears to its users as a single coherent system
- A system in which hardware and software components of networked computers communicate and coordinate their activity only by passing messages
- A computing platform built with many computers that:
 - Operate concurrently
 - Are physically distributed (have their own failure modes)
 - Are linked by a network
 - Have independent clocks



Recap: Characteristics

- Concurrent execution of processes:
 - Non-determinism, Race Conditions, Synchronization, Deadlocks, and so on
- No global clock
 - Coordination is done by message exchange
 - No Single Global notion of the correct time
- No global state
 - No Process has a knowledge of the current global state of the system
- Units may fail independently
 - Network Faults may isolate computers that are still running
 - System Failures may not be immediately known



Recap: Challenges and Goals

→ Some Important aspects of DC:

- Reliable network
- Zero Latency
- Infinite Bandwidth
- Secure network
- Fixed Topology,
- Only one administrator
- Zero Transport cost
- Homogeneous Network

→ Remember these points while developing scalable applications



Recap: Challenges / Goals of a DS

What are the challenges / goals of distributed systems?

- Heterogeneity
- Openness
- Security
- Scalability
- Failure Handling
- Concurrency
- Transparency





FUNDAMENTALS OF DISTRIBUTED SYSTEMS

Can you parallelize all apps?

- What can you parallelize?
- Identify Instructions that could execute in parallel?
- If the proportion of an application that you can parallelize is x ($0 < x < 1$),
Maximum speed up = $1/(1-x)$
- Let us assume that a task needs t_s time to run on one machine. How much time does it take to run the same task on p CPUs?
 - Let the running time be t_p of a task running on p CPUs.
 - Now $t_p = t_o + t_s * (1 - x + x/p)$, where t_o is the overhead added due to parallelisation (communication, synchronization, etc.)

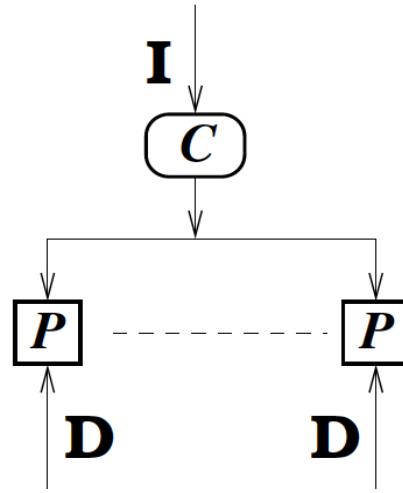


Flynn's Classification

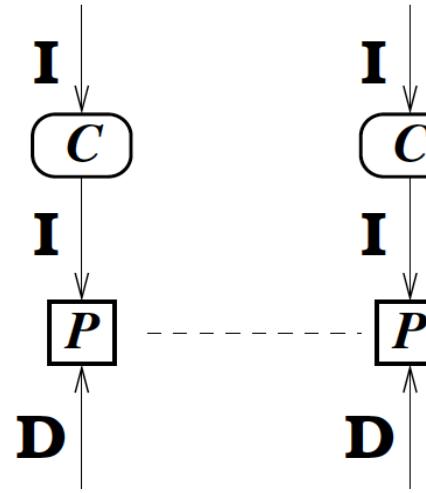
- Single Instruction Single Data (SISD) Stream
 - Traditional von Neumann architecture
- Single Instruction Multiple Data (SIMD) Stream
 - Scientific Applications, Vector Processors, array processors and so on
- Multiple Instruction Single Data (MISD) Stream
 - Visualization is an example
- Multiple Instruction Multiple Data (MIMD) Stream
 - Distributed Systems



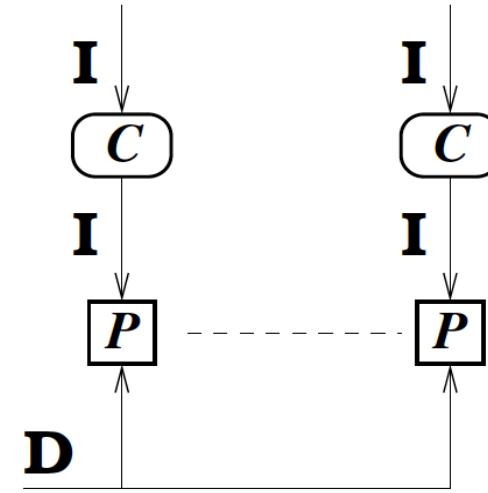
Flynn's Classification



(a) SIMD



(b) MIMD



(c) MISD

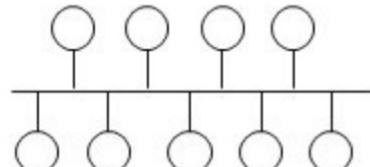
C Control Unit

P Processing Unit

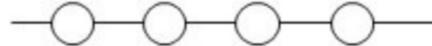
I instruction stream

D data stream

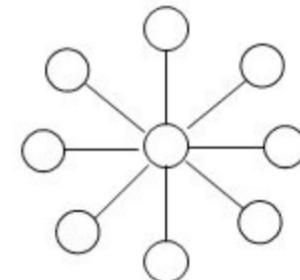
Interconnection Topologies



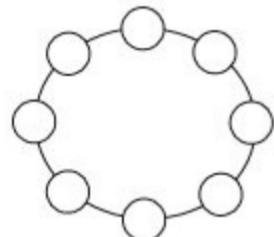
a) Bus



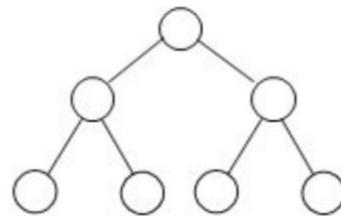
b) Linear array



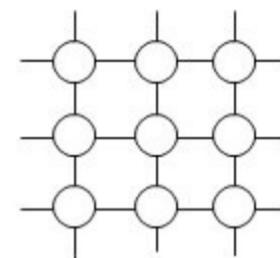
c) Star



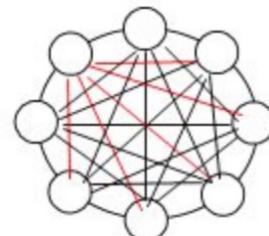
d) Ring



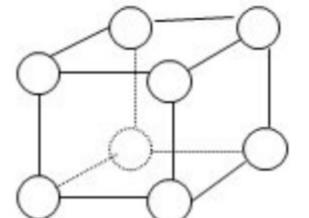
e) Tree



f) Near-neighbor mesh



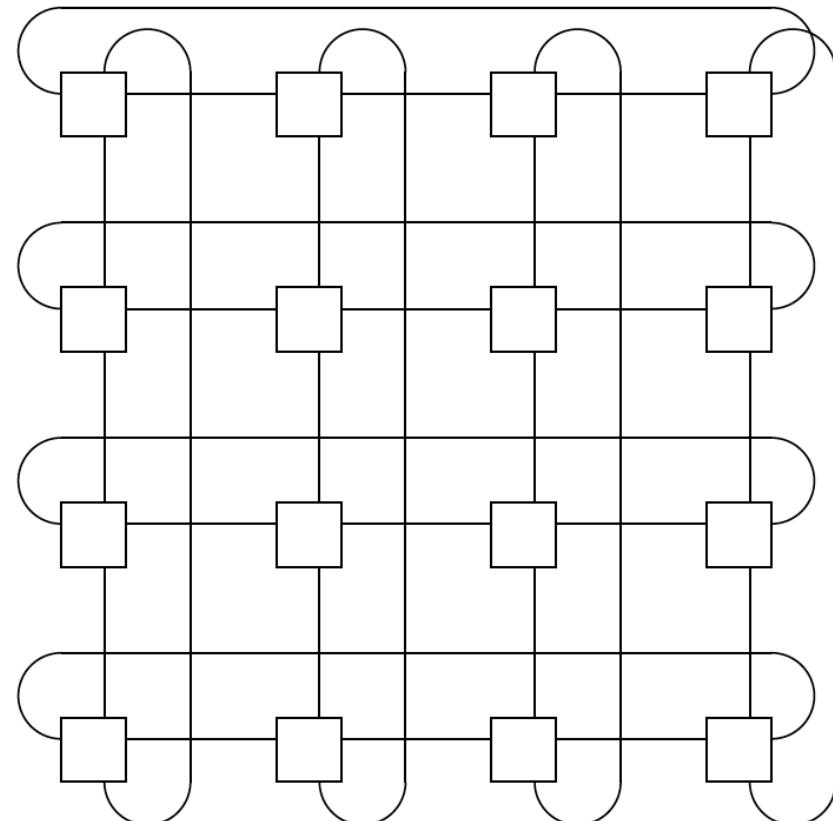
g) Completely connected



h) 3-cube (hypercube)

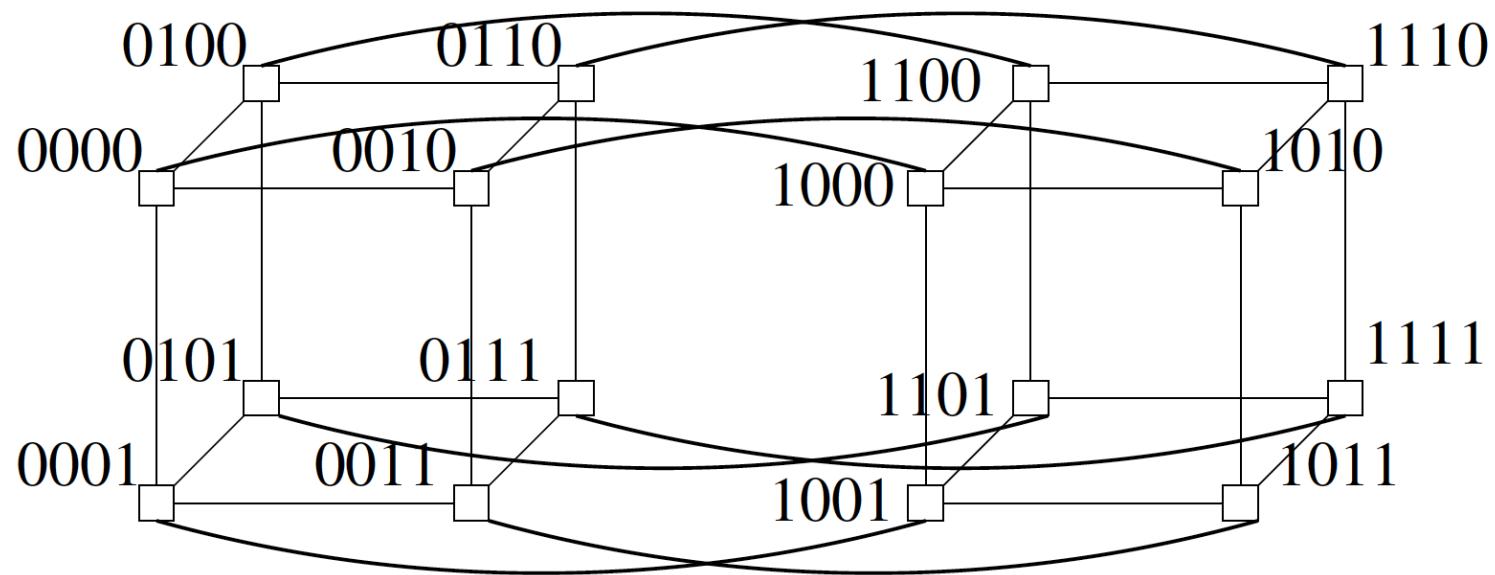
Interconnection Topologies (contd)

→ Wrap - around
mesh (torus)



k-ary d-cubes

→ Hypercube of Dimension 4



→ k-ary d-cubes (generalized version)

Message Passing Systems

Basic Primitive Operations

→ Send

→ send message from process A to Process B

A → B

→ Receive

→ receive message at Process B from Process A

B ← A

→ Compute at A and / or B

→ do the specific computations at A and / or B



Message Passing vs. Shared Memory

- Emulate MP on SM
 - Partition the shared address space
 - Send / Receive messages via shared address space
- Emulate SM on MP
 - Model each shared location as a separate process
 - Send: Write to the shared object by sending messages to owner process for the object
 - Receive: Read from shared object by sending query to the owner process of the object



Synchronous vs. Asynchronous

- **Synchronous (send / receive)**
 - Handshake between sender and receiver
 - *send* completes only when *receive* completes
 - *receive* completes only when copying of the data to the buffer is over
- **Asynchronous (send)**
 - No need for handshake between sender and receiver
 - Control returns to the invoking process when data copied out of user-specified buffer



Blocking vs. non-blocking

→ Blocking

- Control returns to the invoking process after the task completes

→ Non-blocking

- Control returns to the invoking process when data copied out of user-specified buffer
- *send* completes even before copying the data to the user buffer
- *receive* may happen even before data may have arrived from the sender
- How to order EVENTS?



A Distributed Program

- A distributed program is composed of a set of n asynchronous processes, $p_1, p_2, \dots, p_i, \dots, p_n$
- The processes do not share a global memory and communicate solely by passing messages
- The processes do not share a global clock that is instantaneously accessible to these processes
- Process execution and message transfer are asynchronous
- Without loss of generality, we assume that each process is running on a different processor
- Let C_{ij} denote the channel from process p_i to p_j and let m_{ij} denote a message sent by p_i to p_j
- The message transmission delay is finite and unpredictable



A Model of Distributed Executions

- The execution of a process consists of a sequential execution of its actions.
- The actions are atomic and modeled as three types of events: **internal events**, message **send events**, and **message receive events**
- Let e_i^x denote the x^{th} event at process p_i .
- For a message m , let $\text{send}(m)$ and $\text{receive}(m)$ denote send and receive events, respectively.
- The occurrence of events changes the states of respective processes and channels.
- Internal event → **changes state of the process**
- Send and Receive events change the **state of the process** that sends / receives the message & the **state of the channel** on which the message is sent / received respectively



A Model of Distributed Executions

- The events at a process are linearly ordered by their order of occurrence.
- The execution of process p_i produces a sequence of events $e_i^1, e_i^2, \dots, e_i^x, e_i^{x+1}, \dots$ and is denoted by H_i where

$$H_i = (h_i, \rightarrow i)$$

h_i is the set of events produced by p_i and binary relation $\rightarrow i$ defines a linear order on these events

- Linear Relation: Mathematically, the independent variable is multiplied by the slope coefficient, added by a constant, which determines the dependent variable
- Relation $\rightarrow i$ expresses causal dependencies among the events of p_i



A Model of Distributed Executions (contd)

- The send and the receive events signify the flow of information between processes and establish causal dependency from the sender process to the receiver process
- Define a relation \rightarrow_{msg} that captures the causal dependency due to message exchanges as follows:

For every message m that is exchanged between two processes, we have

$$send(m) \rightarrow_{msg} receive(m)$$

- Relation \rightarrow_{msg} defines causal dependencies between the pairs of corresponding send and receive events



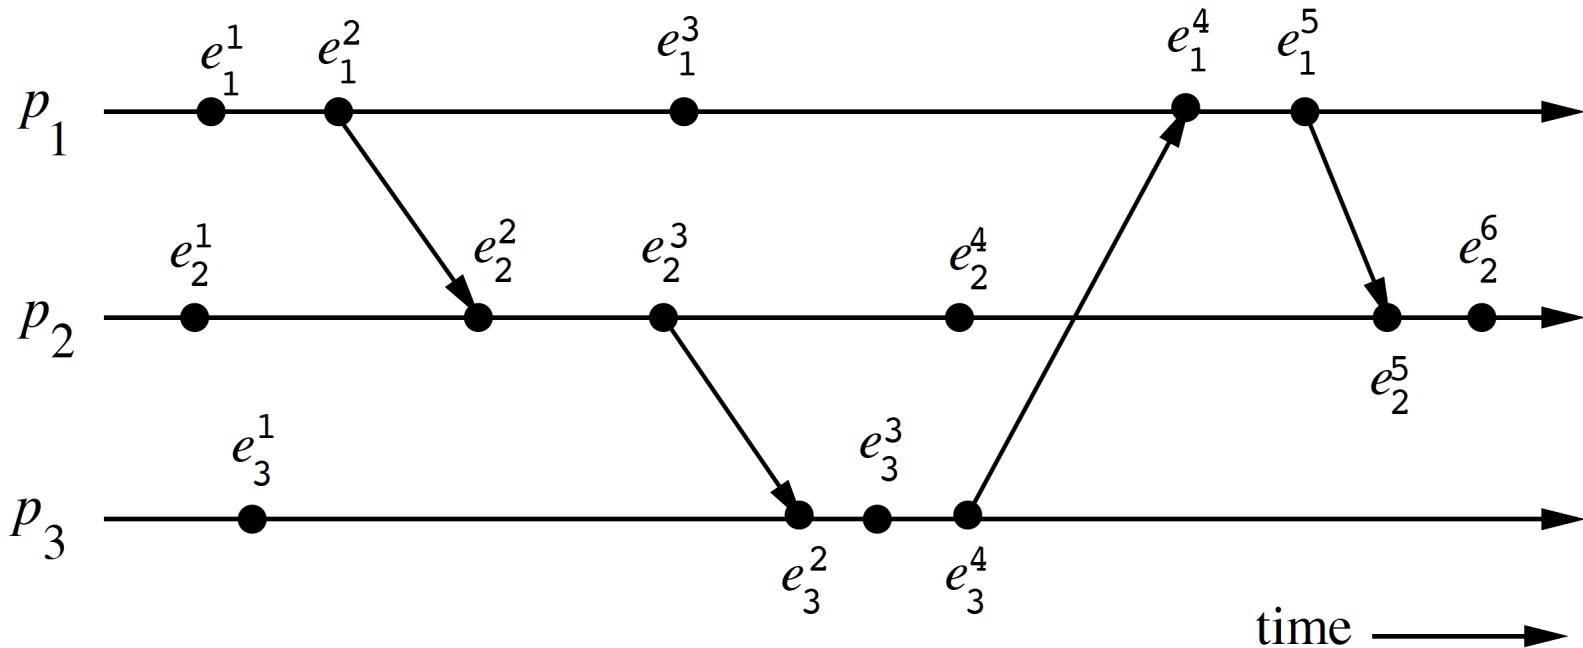
A State-Time diagram

- The evolution of a distributed execution is depicted by a space-time diagram
- A horizontal line represents the progress of a specific process
- A dot indicates an event
- A slant arrow indicates a message transfer

- Since an event execution is atomic (indivisible and instantaneous), it is justified to denote it as a dot on a process line



A State-Time diagram - An Example



- For Process p_1 :
- Second event is a message send event
 - First and Third events are internal events
 - Fourth event is a message receive event

Applications

- Mobile Systems
- Sensor networks
- Pervasive Computing
 - Smart workplace
 - Intelligent Home
- Peer-to-peer computing
- Distributed Agents
- Distributed Data Mining
- Grid Computing
- Security aspects in Distributed Systems



Summary

Focused on exploring the following:

- Goals and Challenges of DS
 - Fundamental aspects while building distributed applications
- A model of Distributed Computations
 - Primitives of Distributed Communications
 - Message Passing is the main focus
 - Properties of distributed Computations
 - Events and their ordering
 - How to handle Causal Precedence ?
 - Lamport's Logical Clocks ?
 - Many more to come up ... stay tuned in !!



Penalties



- Every Student is expected to strictly follow a fair Academic Code of Conduct to avoid penalties
- Penalties is heavy for those who involve in:
 - Copy and Pasting the code
 - Plagiarism (copied from your neighbor or friend - in this case, both will get "0" marks for that specific take home assignments)
 - If the candidate is unable to explain his own solution, it would be considered as a "copied case"!!
 - Any other unfair means of completing the assignments

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students



How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <https://www.iiits.ac.in/people/regular-faculty/dr-rajendra-prasath/>

(OR)

→ <http://rajendra.2power3.com>



Assistance

- You may post your questions to me at any time
- You may meet me in person on available time or with an appointment
- You may ask for one-to-one meeting

Best Approach

- You may leave me an email any time
(email is the best way to reach me faster)





Questions

It's Your Time



THANKS

