

Spring 2023



Distributed Computing

- Logical and Vector Clocks



Dr. Rajendra Prasath

Indian Institute of Information Technology Sri City, Chittoor



24th January 2023 (<http://rajendra.2power3.com>)

> Distributed Computing?

- How will you design a Distributed Algorithm?



- Learn to Solve using Distributed Algorithms

> About this Lecture

What do we learn today?

- This covers a model of distributed computations that every algorithm designer needs to know
 - **Leslie Lamport's Logical Clocks**
 - **Vector Clocks**
 - **Some Examples using Logical clocks and Vector Clocks**

Let us **explore** these topics ➔ ➔ ➔

Recap: Distributed Systems

A Distributed System:

- A collection of independent systems that appears to its users as a single coherent system
- A system in which hardware and software components of networked computers communicate and coordinate their activity only by passing messages
- A computing platform built with many computers that:
 - Operate concurrently
 - Are physically distributed (have their own failure modes)
 - Are linked by a network
 - Have independent clocks

Recap: Characteristics

- **Concurrent execution of processes:**
 - Non-determinism, Race Conditions, Synchronization, Deadlocks, and so on
- **No global clock**
 - Coordination is done by message exchange
 - No Single Global notion of the correct time
- **No global state**
 - No Process has a knowledge of the current global state of the system
- **Units may fail independently**
 - Network Faults may isolate computers that are still running
 - System Failures may not be immediately known

What did you learn so far?

- Goals / Challenges in Message Passing systems
- Distributed Sorting
- Space-Time diagram
- Partial Ordering / Total Ordering
- Causal Precedence Relation
 - Happens Before
- Concurrent Events
 - How to define Concurrent Events
 - Logical vs Physical Concurrency
- Causal Ordering
- Local State vs. Global State

Recap




Causal Ordering

- The “causal ordering” model is based on Lamport’s “happens before” relation
- A system that supports the causal ordering model satisfies the following property:

CO: For any two messages m_{ij} and m_{kj} ,
if $send(m_{ij}) \rightarrow send(m_{kj})$,
then $receive(m_{ij}) \rightarrow receive(m_{kj})$

- This property ensures that causally related messages destined to the same destination are delivered in an order that is consistent with their causality relation.
- Causally ordered delivery of messages implies FIFO message delivery. (Note that $CO \subset FIFO \subset Non-FIFO$.)
- Causal ordering model considerably simplifies the design of distributed algorithms because it provides a built-in synchronization.

Cuts of a Distributed Computation



Physical vs Logical clocks?

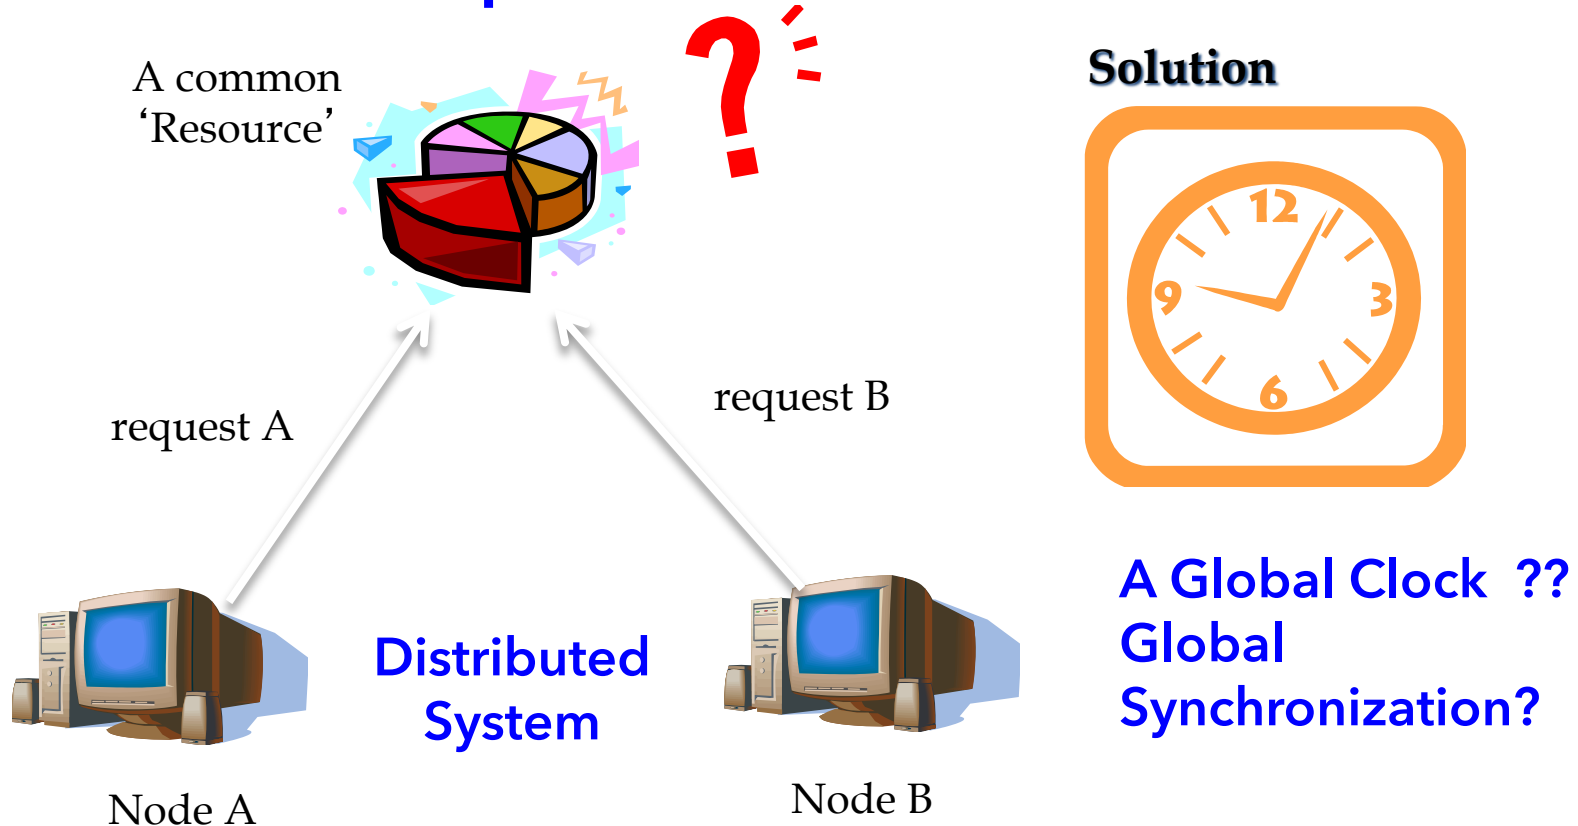
- Logical Clocks
 - Design and Implementation
- Three Different Ways
 - Scalar Time
 - Vector Time
 - Matrix Time
- Virtual Clocks
 - Time Wrap Mechanism
- Clock Synchronization
 - NTP Synchronization Protocol

Logical Clocks

- Logical Clocks (Lamport 1978)
 - Based on "Happens Before" concept
- Knowing the ordering of events is important (?!)
- not enough with physical time
- Two simple points [Lamport 1978]
 - the order of two events in the same process
 - the event of sending message always happens before the event of receiving the message


Events Ordering - An Example

→ Which request was made first?



Events Ordering - An Example (contd)

→ Which request was made first?




Solution

Individual
Clocks?

Are individual
clocks accurate,
precise?

One clock
might run
faster/slower?

Logical Clocks (Lamport 1978)



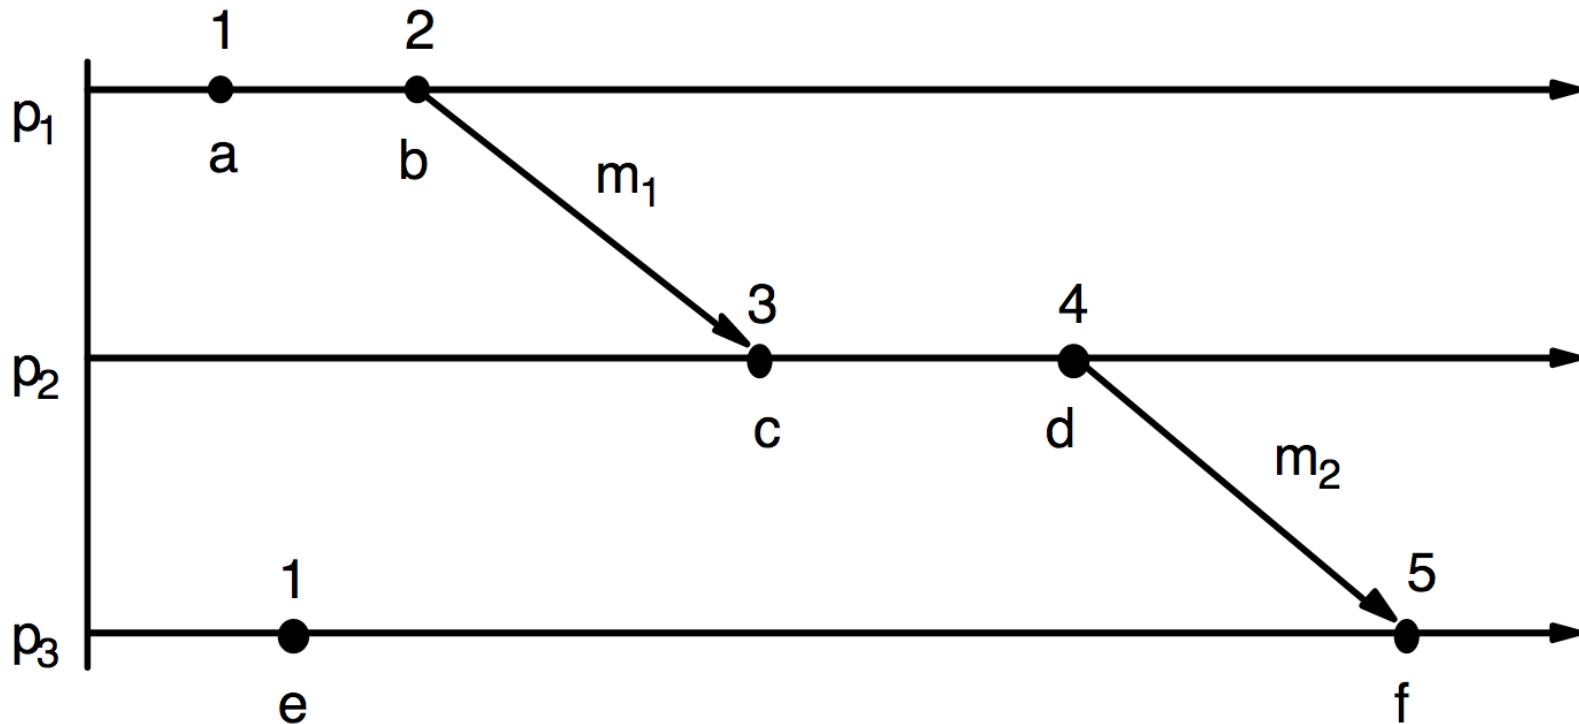
Synchronization in a Distributed System:

- Event Ordering:
 - Which event occurred first?
- How to sync the clocks across the nodes?
- Can we define the notion of **happened-before** without using physical clocks?

Lamport's Logical clocks

- A monotonically increasing software counter
- It does (need) not relate to a physical clock
- Each process p_i has a logical clock L_i
- LC_1 : L_i is incremented by 1 before each event at process p_i
- LC_2 :
 - A) when process p_i sends message m ,
it piggybacks $t = L_i$
 - B) when p_j receives (m, t) , it sets $L_j = \max(L_j, t)$ and
applies LC_1 before timestamping the event
 $receive(m)$

A Close Look




- $e \rightarrow e' \Rightarrow L(e) < L(e')$ but not vice versa
- Example: event b and event e


How to implement Lamport's clocks?

- When a message is transmitted from P1 to P2, P1 will encode the send time into the message.
- When P2 receives the message, it will record the time of receipt
- If P2 discovers that the time of receipt is before the send time, P2 will update its software clock to be one greater than the send time (1 milli second at least)
- If the time at P2 is already greater than the send time, then no action is required for P2
- With these actions the "happens-before" relationship of the message being sent and received is preserved

Correction of Clocks



(a)




(b)

An Illustration

$m1 \rightarrow m3$
 $\Rightarrow C(m1) < C(m3)$

$m2 \rightarrow m3$
 $\Rightarrow C(m2) < C(m3)$

Here which event,
either $m1$ or $m2$,
caused
 $m3$ to be sent?



Limitations

- Lamport's logical clocks lead to a situation where all events in a distributed system are totally ordered

If $a \rightarrow b$, then we can say $C(a) < C(b)$

- Unfortunately, with Lamport's clocks, nothing can be said about the actual time of a and b

If the logical clock says $a \rightarrow b$, that does not mean in reality that a actually happened before b in terms of **real time**

Issues with Lamport Clocks

- The problem with Lamport clocks is that they do not capture **causality**
- If we know that $a \rightarrow c$ and $b \rightarrow c$ we cannot say which action initiated c
- This kind of information can be important when trying to reply events in a distributed system (such as when trying to recover after a crash)
- If one node goes down, if we know the causal relationships between messages, then we can replay those messages and respect the causal relationship to get that node back up to the state it needs to be in

Vector Clocks

- Vector clocks allow causality to be captured
- Rules of Vector Clocks:
 - A vector clock $VC(a)$ is assigned to an event a
 - If $VC(a) < VC(b)$ for events a and b , then event a is known to causally precede b
- Each Process P_i maintains a vector VC_i with the following properties:
 - $VC_i[i]$ is the number of events that have occurred so far at P_i that is, $VC_i[i]$ is the **local logical clock** at process P_i
 - If $VC_i[j] = k$ then P_i knows that k events have occurred at P_j . It is thus P_i 's **knowledge of the local time at P_j**


Implementing Vector Clocks

- Increment $VC_i[i]$ at each new event at P_i
- **Updating Clocks:**
 - Before executing any event (sending a message or an internal event):

P_i executes $VC_i[i] \leftarrow VC_i[i] + 1$

- When process P_i sends a message m to P_j , it sets m 's (vector) timestamp $ts(m) = VC_i$
- Upon receiving a message m , process P_j adjusts its own vector by setting $VC_j[k] \leftarrow \max(VC_j[k], ts(m)[k])$ for each k

An Example



Understanding **Vector Clocks**

Meaning of $=$, \leq , $<$ for vector timestamps

(1) $VC = VC'$ iff $VC[j] = VC'[j]$ for $j = 1, 2, \dots, N$

(2) $VC \leq VC'$ iff $VC[j] \leq VC'[j]$ for $j = 1, 2, \dots, N$

(3) $VC < VC'$ iff $VC \leq VC'$ and $VC \neq VC'$


Examples:

$(1, 3, 2) < (1, 3, 3)$

$(1, 3, 2) \parallel (2, 3, 1)$

➔ **Note:** $e \rightarrow e'$ implies $VC(e) < VC(e')$ (The converse is also true)

An illustrative example



Less than or equal:

➔ $ts(a) \leq ts(b)$ if $ts(a)[i] \leq ts(b)[i]$ for all i
 $(2, 4, 2) \leq (3, 5, 2)$

➔ $ts(e11) = (1, 0, 0)$ and $ts(e22) = (2, 2, 0)$
 This implies $e11 \rightarrow e22$

Summary

- **A model of Distributed Computations**
 - Causal Precedence Relations
 - Global State and Cuts of a DS
 - PAST and FUTURE events
- **What about the ordering of events?**
 - How do we efficiently handle the ordering of events (discrete events)?
 - Lamport's Logical Clocks ?
 - Vector Clocks
- **Many more to come up ... stay tuned in !!**

Penalties



- Every Student is expected to strictly follow a fair Academic Code of Conduct to avoid penalties
- Penalties is heavy for those who involve in:
 - Copy and Pasting the code
 - Plagiarism (copied from your neighbor or friend - in this case, both will get "0" marks for that specific take home assignments)
 - If the candidate is unable to explain his own solution, it would be considered as a "copied case"!!
 - Any other unfair means of completing the assignments

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <https://www.iiits.ac.in/people/regular-faculty/dr-rajendra-prasath/>

(OR)

→ <http://rajendra.2power3.com>

Assistance

- You may post your questions to me at any time
- You may meet me in person on available time or with an appointment
- You may ask for one-to-one meeting

Best Approach

- You may leave me an email any time
(email is the best way to reach me faster)



Questions It's Your Time



Contact Information

Dr. Rajendra Prasath
IIIT Sri City, Chittoor