

Monsoon 2022



Distributed Computing

- Global Snapshot Algorithm



Dr. Rajendra Prasath

Indian Institute of Information Technology Sri City, Chittoor



20th September 2022 (<http://rajendra.2power3.com>)

> Distributed Computing?

- How will you design a Distributed Algorithm?



- Learn to Solve using Distributed Algorithms

Recap: Distributed Systems

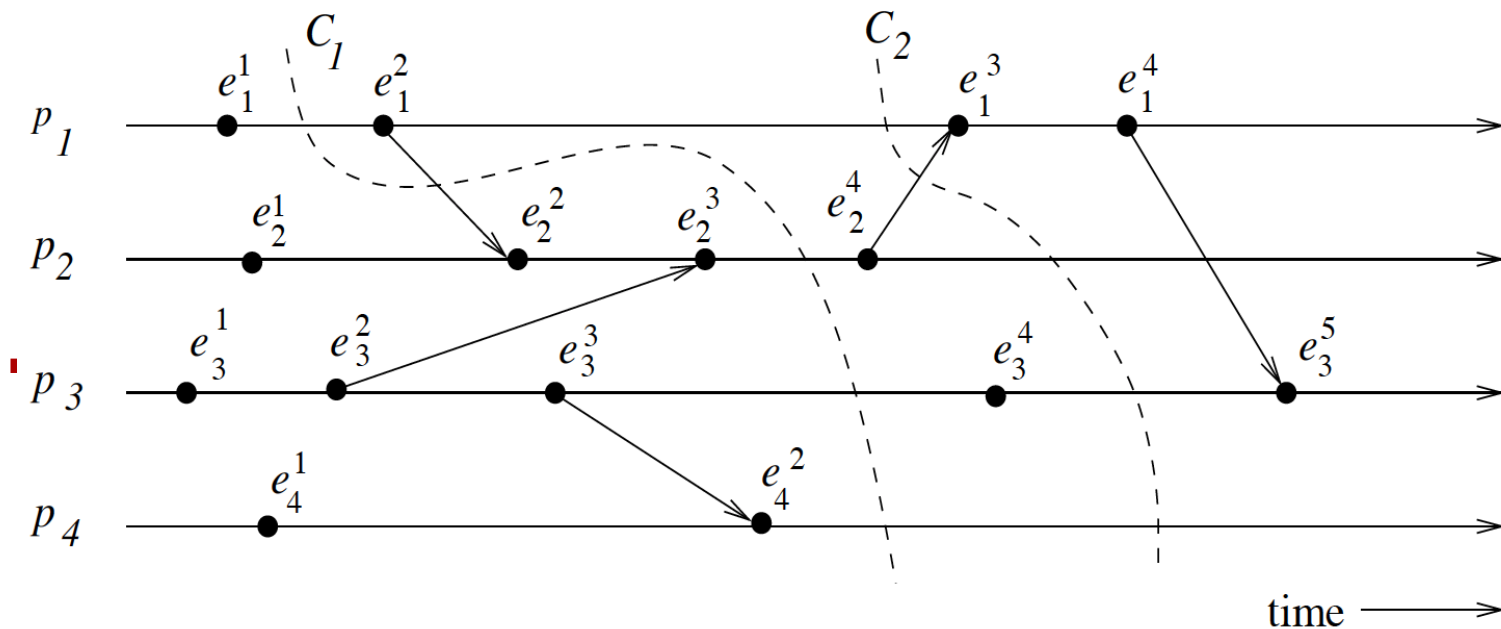
A Distributed System:

- A collection of independent systems that appears to its users as a single coherent system
- A system in which hardware and software components of networked computers communicate and coordinate their activity only by passing messages
- A computing platform built with many computers that:
 - Operate concurrently
 - Are physically distributed (have their own failure modes)
 - Are linked by a network
 - Have independent clocks

Recap: Characteristics

- **Concurrent execution of processes:**
 - Non-determinism, Race Conditions, Synchronization, Deadlocks, and so on
- **No global clock**
 - Coordination is done by message exchange
 - No Single Global notion of the correct time
- **No global state**
 - No Process has a knowledge of the current global state of the system
- **Units may fail independently**
 - Network Faults may isolate computers that are still running
 - System Failures may not be immediately known

Recap: Cuts - Distributed Computation



What did you learn so far?

- Goals / Challenges in Message Passing systems
- Distributed Sorting
- Space-Time diagram
- Partial Ordering / Total Ordering
- Causal Precedence Relation
 - Happens Before
- Concurrent Events
 - How to define Concurrent Events
 - Logical vs Physical Concurrency
- Causal Ordering
- Logical Clocks vs Physical Clocks

> About this Lecture

What do we learn today?

- This covers a model of distributed computations that every algorithm designer needs to know
- **Global States**
- **Global Snapshot algorithm**

Let us **explore** these topics → → →

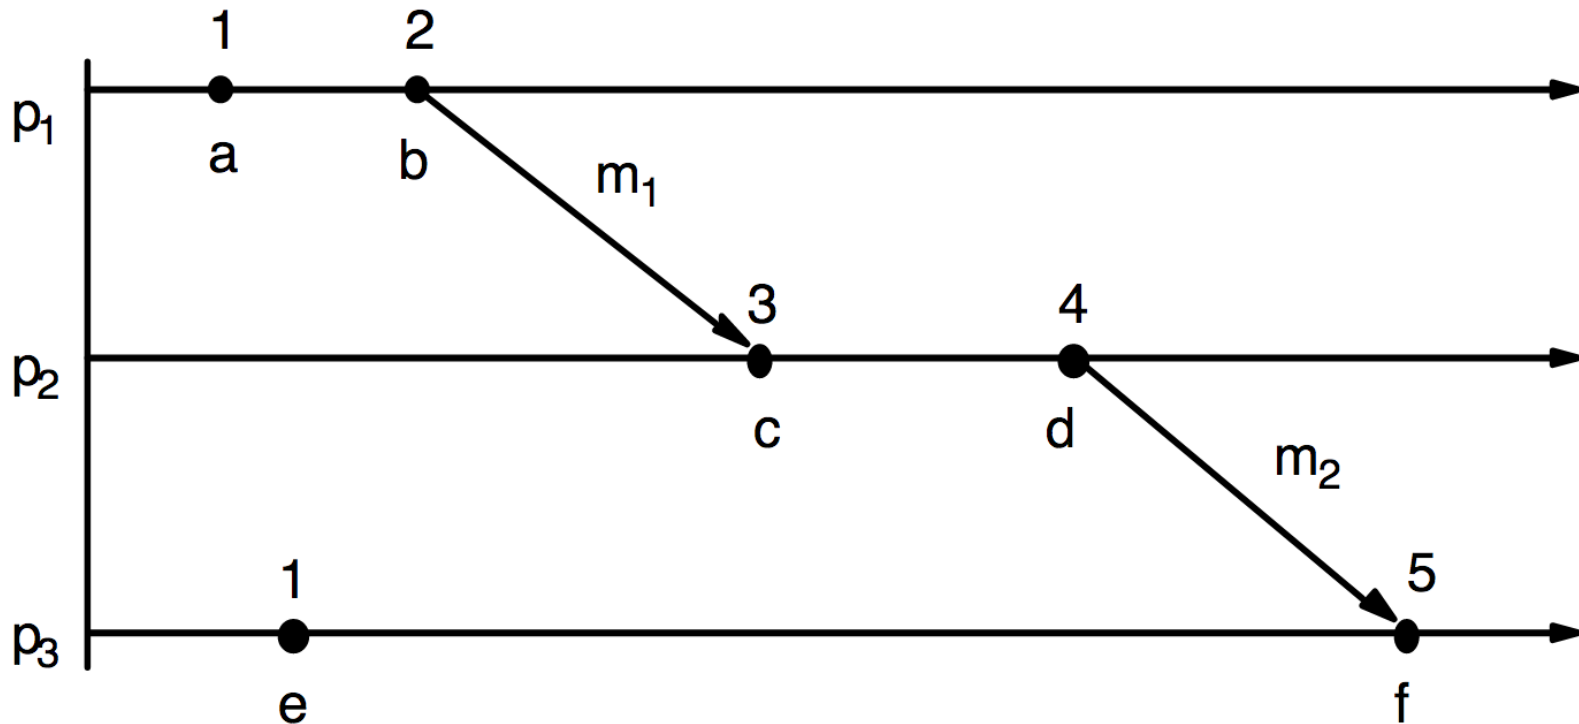
Causal Ordering

- The “causal ordering” model is based on Lamport’s “happens before” relation
- A system that supports the causal ordering model satisfies the following property:

CO: For any two messages m_{ij} and m_{kj}
if $send(m_{ij}) \rightarrow send(m_{kj})$,
then $receive(m_{ij}) \rightarrow receive(m_{kj})$

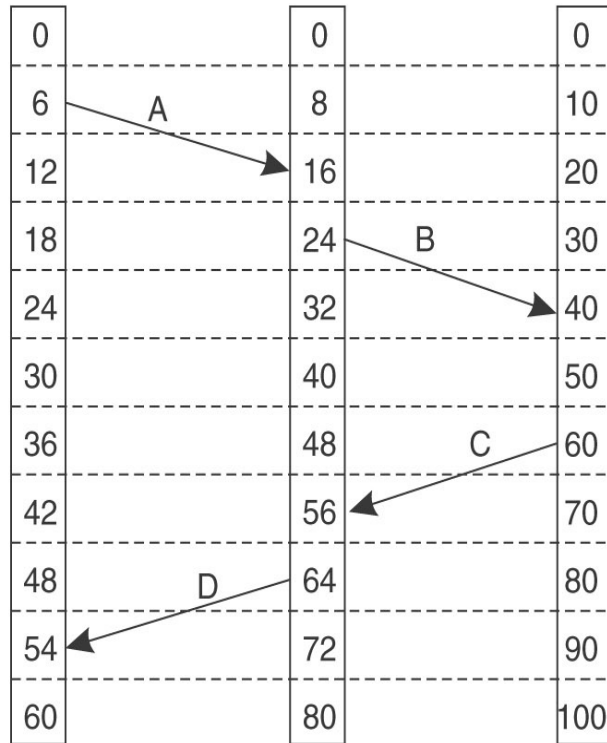
- This property ensures that causally related messages destined to the same destination are delivered in an order that is consistent with their causality relation.
- Causally ordered delivery of messages implies FIFO message delivery. (Note that $CO \subset FIFO \subset Non-FIFO$.)
- Causal ordering model considerably simplifies the design of distributed algorithms because it provides a built-in synchronization.

A Close Look

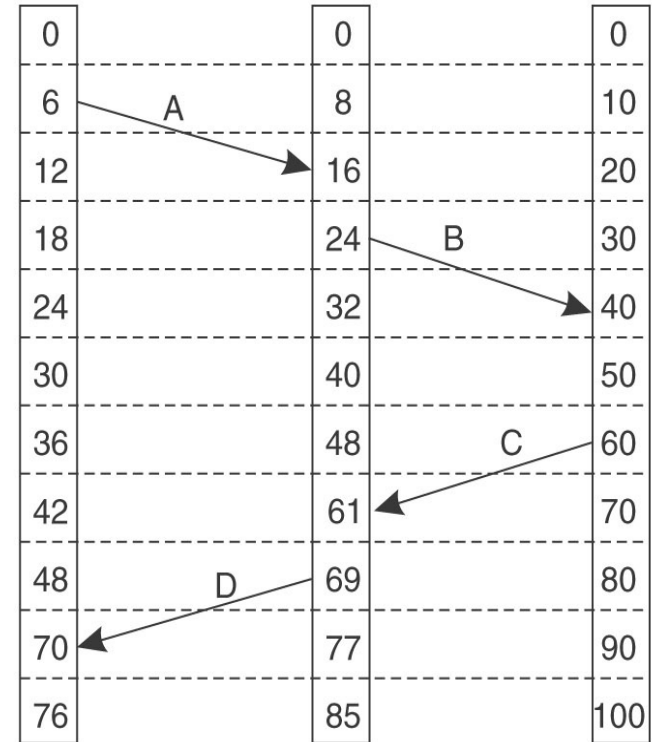


- $e \rightarrow e' \Rightarrow L(e) < L(e')$ but not vice versa
- Example: event b and event e

Correction of Clocks



(a)

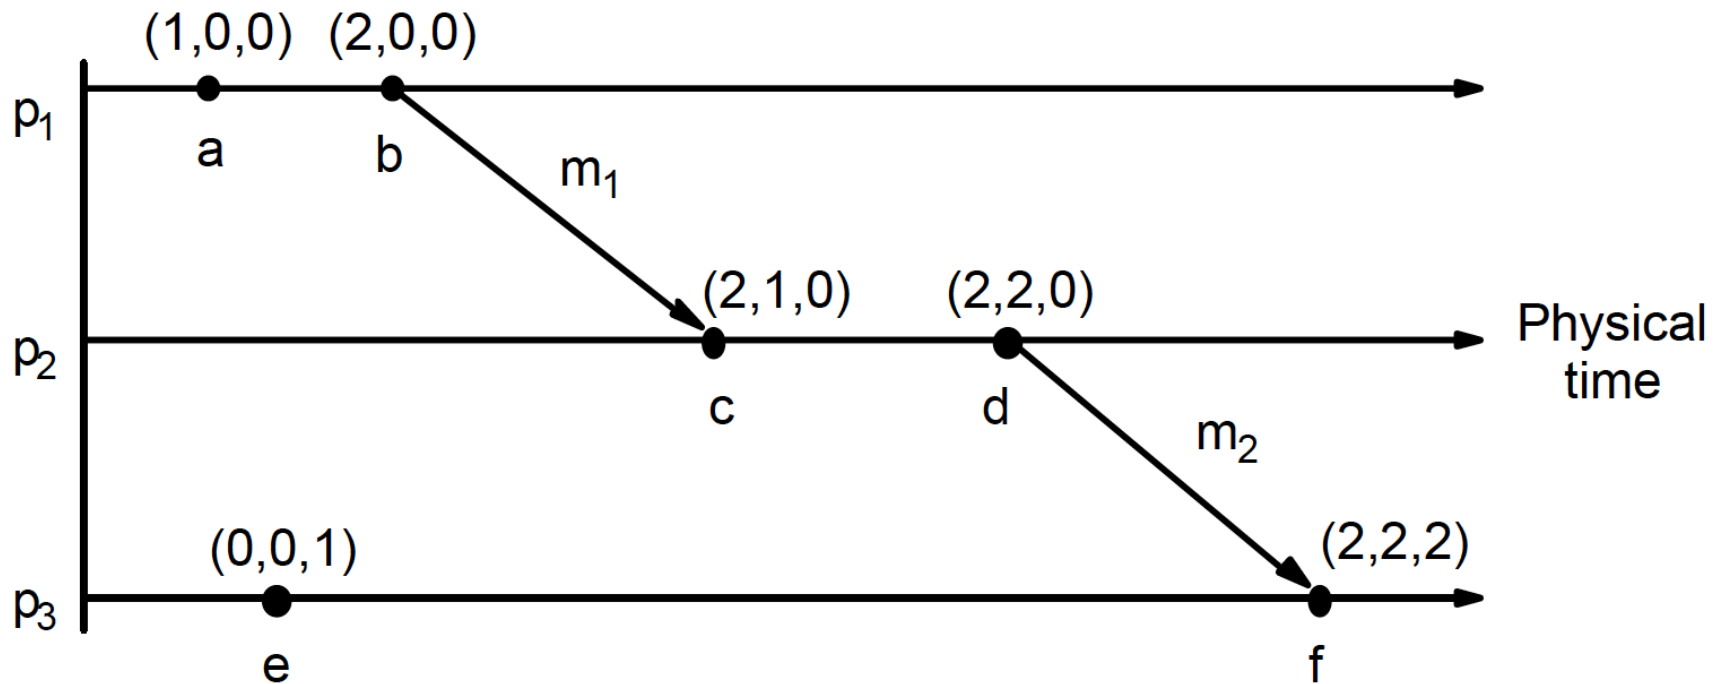


(b)

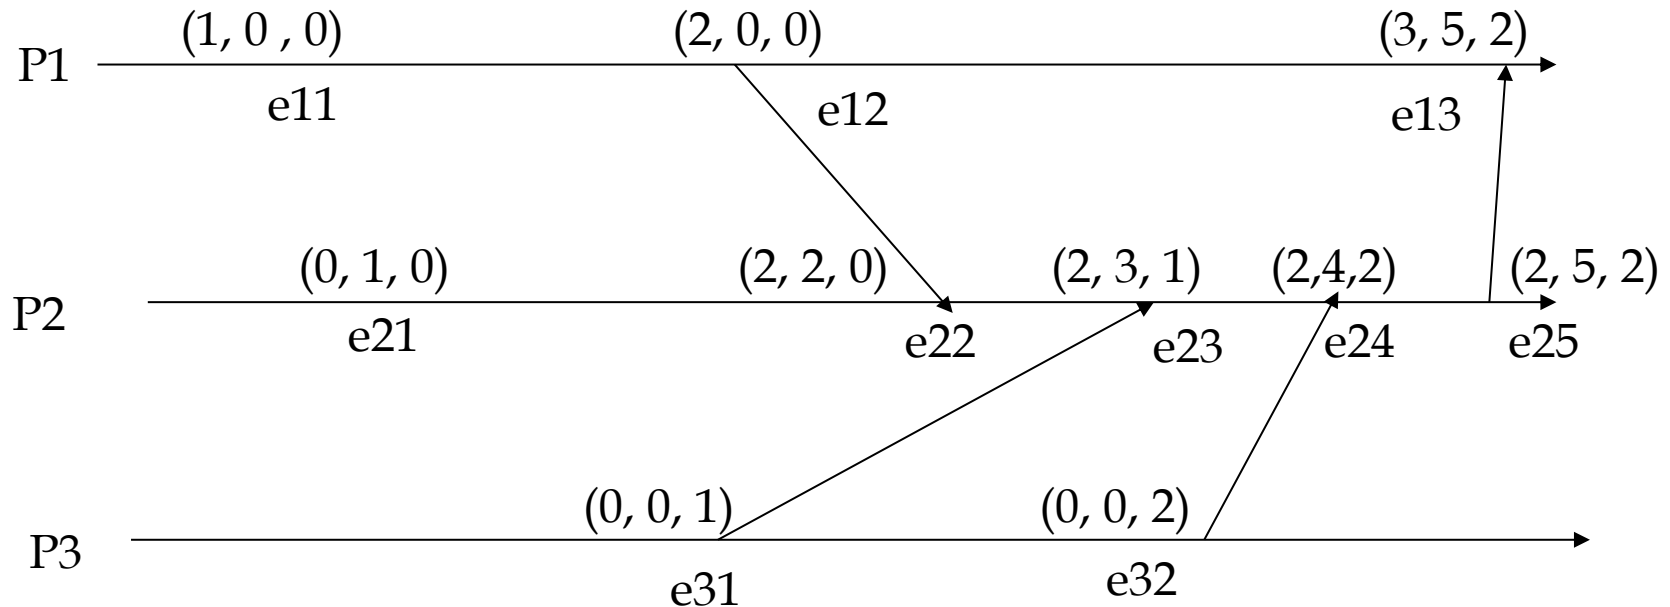
Vector Clocks

- Vector clocks allow causality to be captured
- **Rules of Vector Clocks:**
 - A vector clock $VC(a)$ is assigned to an event a
 - If $VC(a) < VC(b)$ for events a and b , then event a is known to causally precede b
- Each Process P_i maintains a vector VC_i with the following properties:
 - $VC_i[i]$ is the number of events that have occurred so far at P_i that is, $VC_i[i]$ is the **local logical clock** at process P_i
 - If $VC_i[j] = k$ then P_i knows that k events have occurred at P_j . It is thus **P_i 's knowledge of the local time at P_j**

An Example



An illustrative example



Less than or equal:

→ $ts(a) \leq ts(b)$ if $ts(a)[i] \leq ts(b)[i]$ for all i
 $(3, 3, 5) \leq (3, 4, 5)$

→ $ts(e11) = (1, 0, 0)$ and $ts(e22) = (2, 2, 0)$
 This implies $e11 \square e22$

A decorative graphic on the left side of the slide, consisting of a thick blue wavy line that curves upwards and then downwards, with a lighter blue translucent shadow behind it.

GLOBAL STATES IN A DISTRIBUTED SYSTEM

Partial / Total Ordering

A relation \leq is a total order on a set S (" \leq totally orders S ") if the following properties hold:

1. Reflexivity: $a \leq a$ for all a in S
2. Anti-Symmetry: $a \leq b$ and $b \leq a$ implies $a = b$
3. Transitivity: $a \leq b$ and $b \leq c$ implies $a \leq c$
4. Comparability (Trichotomy law):
For any a, b in S , either $a \leq b$ or $b \leq a$.

First 3 properties \rightarrow the axioms of a **partial ordering**
Adding **Trichotomy law** defines a **total ordering** $H=(H, \rightarrow)$



Global State

- A collection of the local states of its components:
 - The processes and the communication channels
- The state of a process is defined by the local contents of processor: **registers, stacks, local memory**
- The state of channel depends the **set of messages in transit in the channel**
- An internal event changes only state of the process
- A send event changes
 - state of the process that sends the message and
 - the state of the channel on which the message is sent.
- Similarly a receive event changes
 - the state of the process that receives the message and
 - the state of the channel on which the message is received

Global State (contd)

Notations

- LS_i^x denotes the state of p_i after occurrence of event e_i^x and before the event e_i^{x+1}
- LS_i^0 denotes the initial state of process p_i
- LS_i^x is a result of the execution of all the events executed by process p_i till e_i^x
- Let $send(m) \leq LS_i^x$ denote the fact:
$$\exists y, 1 \leq y \leq x \text{ s.t. } e_i^y = send(m)$$
- Let $rec(m) (not \leq) LS_i^x$ denote the fact:
$$\forall y, 1 \leq y \leq x \text{ s.t. } e_i^y \text{ (not equal to) } rec(m)$$

Global State (contd)

- The global state of a distributed system is a collection of the local states of the processes and the channels.

A global state GS is defined as,

$$GS = \{ \bigcup_i LS_i^{x_i}, \bigcup_{j,k} SC_{jk}^{y_j, z_k} \}$$

- For a global state to be meaningful, the states of all the components of the distributed system must be recorded at the same instant
- Two important situations (Impossible !!):
 - local clocks at processes were perfectly synchronized
 - there were a global system clock that can be instantaneously read by the processes

A Consistent Global State

Basic idea:

- A state should not violate causality - an effect should not be present without its cause
- A message cannot be received if it was not sent.
- Such states are called consistent global states and are meaningful global states.
- Inconsistent global states are not meaningful in the sense that a distributed system can never be in an inconsistent state

A Consistent Global State

Definition:

→ A global state is a consistent global state iff

$$\forall m_{ij} : \text{send}(m_{ij}) \not\leq LS_i^{x_i} \Leftrightarrow m_{ij} \notin SC_{ij}^{x_i, y_j} \wedge \text{rec}(m_{ij}) \not\leq LS_j^{y_j}$$

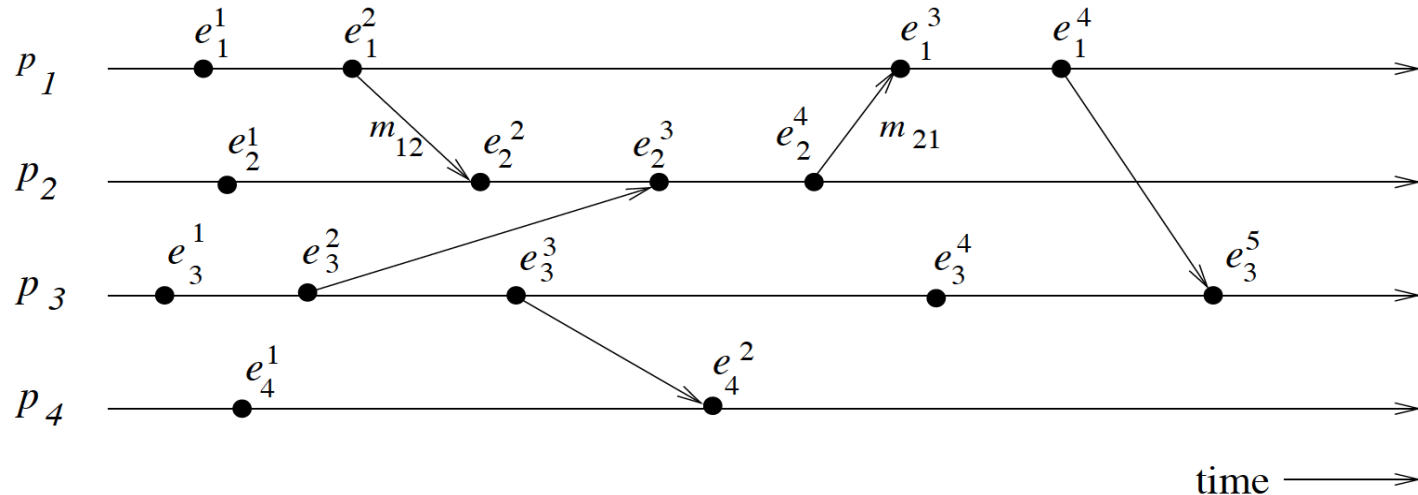
Where the global state is given by

$$GS = \{ \bigcup_i LS_i^{x_i}, \bigcup_{j,k} SC_{jk}^{y_j, z_k} \}$$

→ This implies that the channel state and process state must not include any message that process P_i sent after executing event

Consistent Global State - An Example

→ Space-Time Diagram



→ $GS_1 = \{LS_1^1, LS_2^3, LS_3^3, LS_4^2\}$ is **inconsistent**

→ The state of p_2 has recorded the receipt of m_{12}
however, the state of p_1 has not recorded its send

→ $GS_2 = \{LS_1^2, LS_2^4, LS_3^4, LS_4^2\}$ is **consistent**

→ All channels are empty except C_{21} that contains m_{21}

Consistent Global State - Details

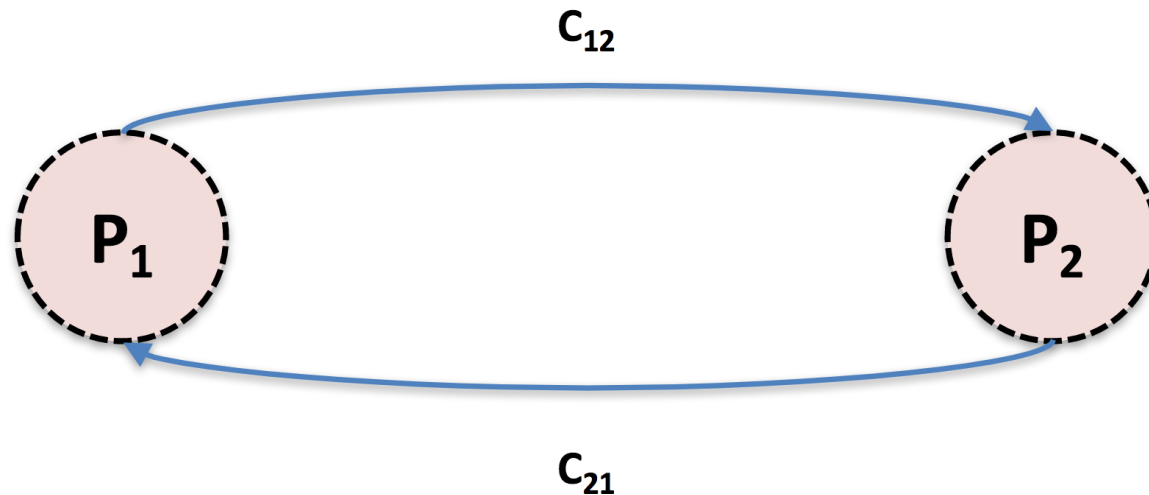
- A global state $GS1 = \{LS_1^1, LS_2^3, LS_3^3, LS_4^2\}$ is **inconsistent** because
 - the state of p_2 has recorded the receipt of message m_{12}
 - The state of p_1 has not recorded its send
- A global state $GS2$ consisting of local states $GS2 = \{LS_1^2, LS_2^4, LS_3^4, LS_4^2\}$ is **consistent**
 - all the channels are empty except C_{21} that contains message m_{21} .

Changes in Global State?

- The global state changes whenever an event happens
 - Process sends message
 - Process receives message
 - Process performs a local event
- Moving from state to state obeys causality

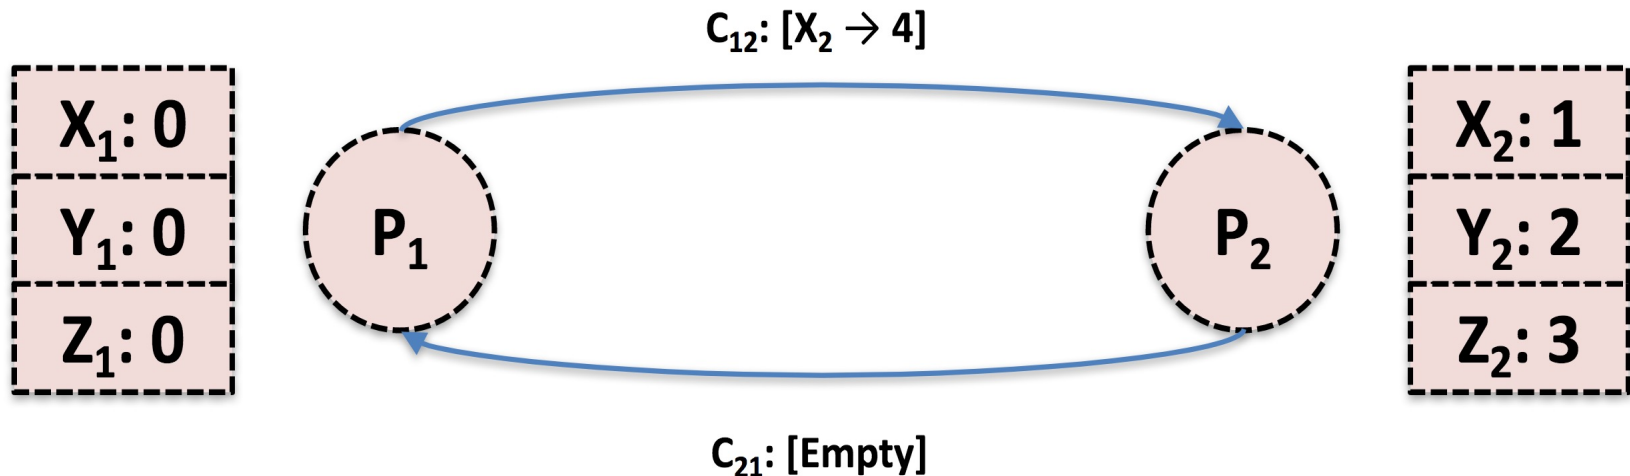
An Example Scenario

- There are two processes: P_1 and P_2
- Channel C_{12} from P_1 to P_2
- Channel C_{21} from P_2 to P_1



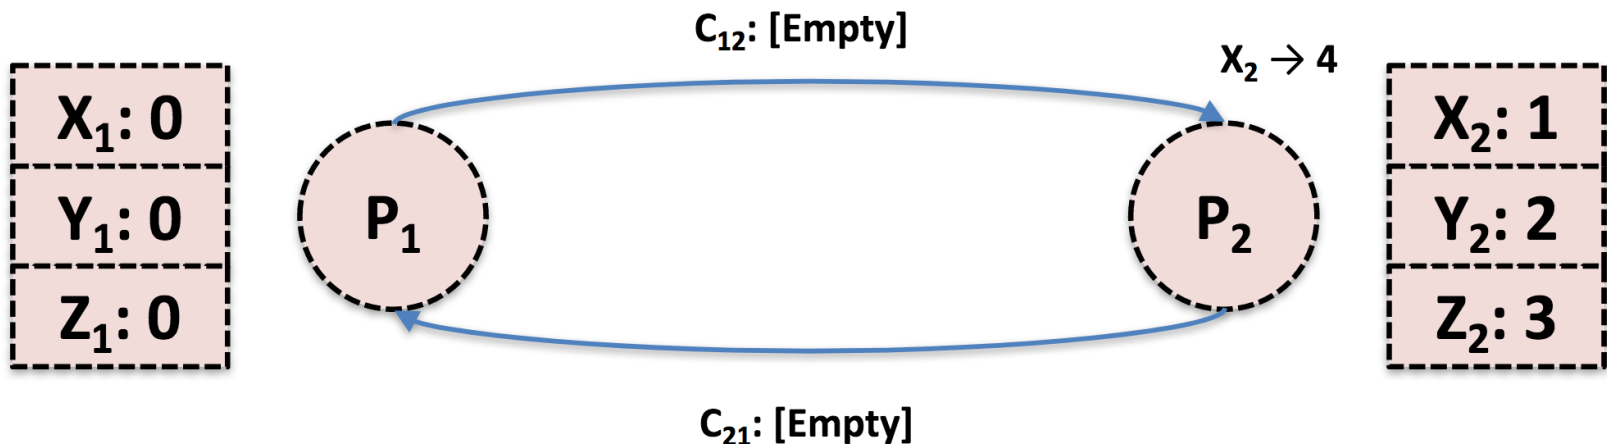
An Example (contd)

→ P_1 tells P_2 to change its state variable, X_2 , from 1 to 4



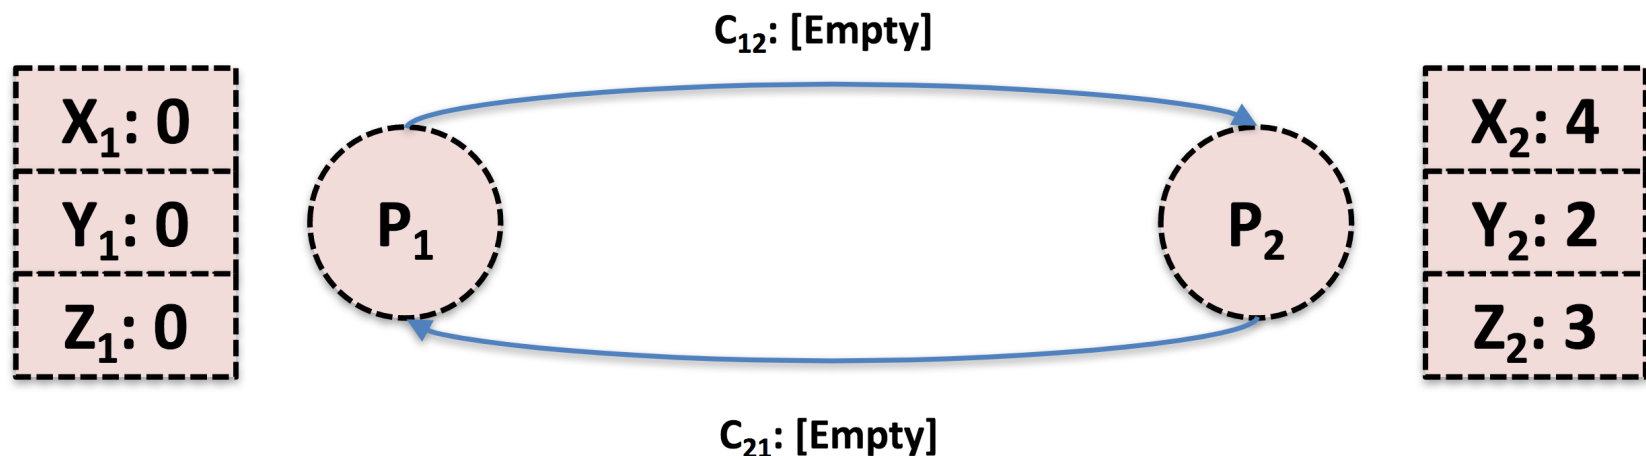
An Example (contd)

→ P_2 receives message from P_1 to change its state variable, X_2 from 1 to 4



An Example (contd)

→ P_2 changes its state variable, X_2 from 1 to 4



Why do we need Global Snapshot?

Checkpointing:

- restart if the application fails

Collecting garbage:

- remove objects that do not have any references

Detecting deadlocks:

- can examine the current application state

Other debugging:

- a little easier to work with than printf



CHANDY-LAMPORT'S GLOBAL SNAPSHOT RECORDING ALGORITHM

System Model (Lamport and Chandy)

Problem: How to record a **global snapshot** (state for each process and each channel)?

System Model (assumptions):

- N Processes in the system
- Communication channels are bi-directional between each ordered pair $(p_i, p_j): p_i \rightarrow p_j$ and $p_j \rightarrow p_i$
- Channels are FIFO: First In First Out
- No Failures
- All messages arrive intact and not duplicated

Requirements

- Snapshot should not interfere with normal application actions
- It should not require application to stop sending messages
- Each Process is able to record its own state:
 - Process State: Application-defined state or in the worst case
 - Its heaps, registers, program counters, code and other related application interfaces
- Global state is collected in a distributed manner (no centralized approach)
- Any Process may initiate the recording of a snapshot (there can be multiple snapshots, we focus on only one)

Initiating a snapshot

- Let's say process P_i initiates the snapshot
- P_i records its own state and prepares a special marker message (distinct from application messages)
- Send the marker message to all other processes (using $N-1$ outbound channels)
- Start recording all incoming messages from channels C_{ji} for j not equal to i

Propagating a snapshot

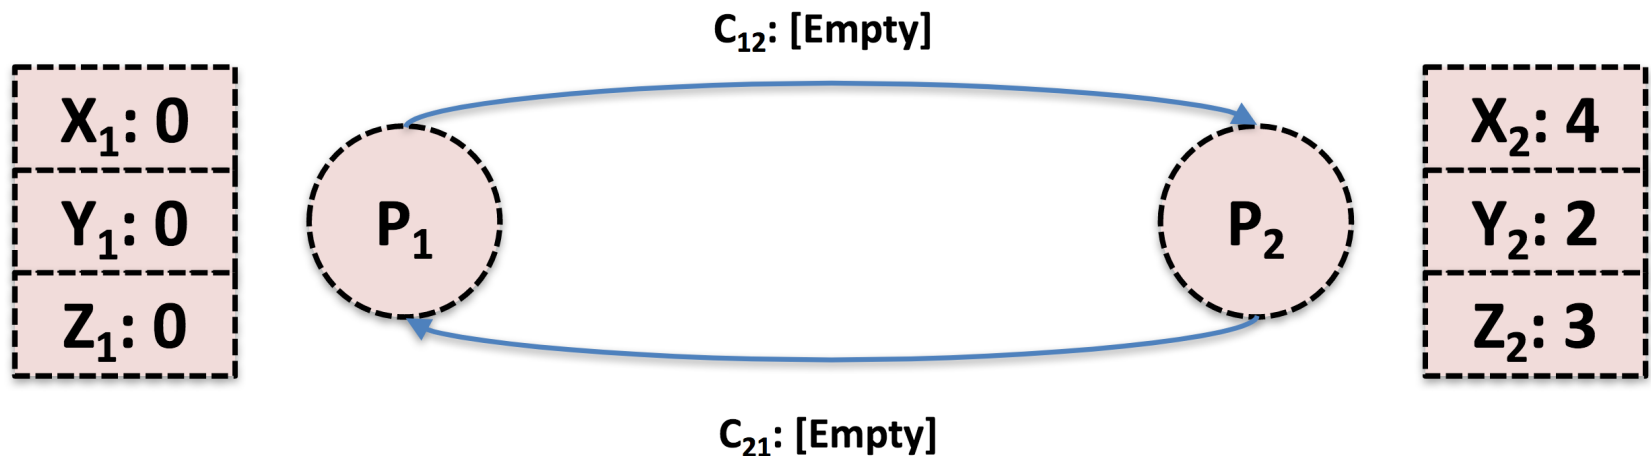
- For all processes P_i (including the initiator) consider a message on channel C_{kj}
- If P_j sees the marker message for the first time
 - P_j records own state and marks C_{kj} as empty
 - Send the marker message to all other processes (using $N-1$ outbound channels)
 - Start recording all incoming messages from channels C_{lj} for l not equal to j or k
- Otherwise
 - add all messages from inbound channels since we began recording their states

Terminating a snapshot

- All processes have received a marker (and recorded their own state)
- All processes have received a marker on all $N-1$ incoming channels (and recorded their states)
- Later, a central server can gather the partial state to build a global snapshot
(Do we require this? – Not Really!!)

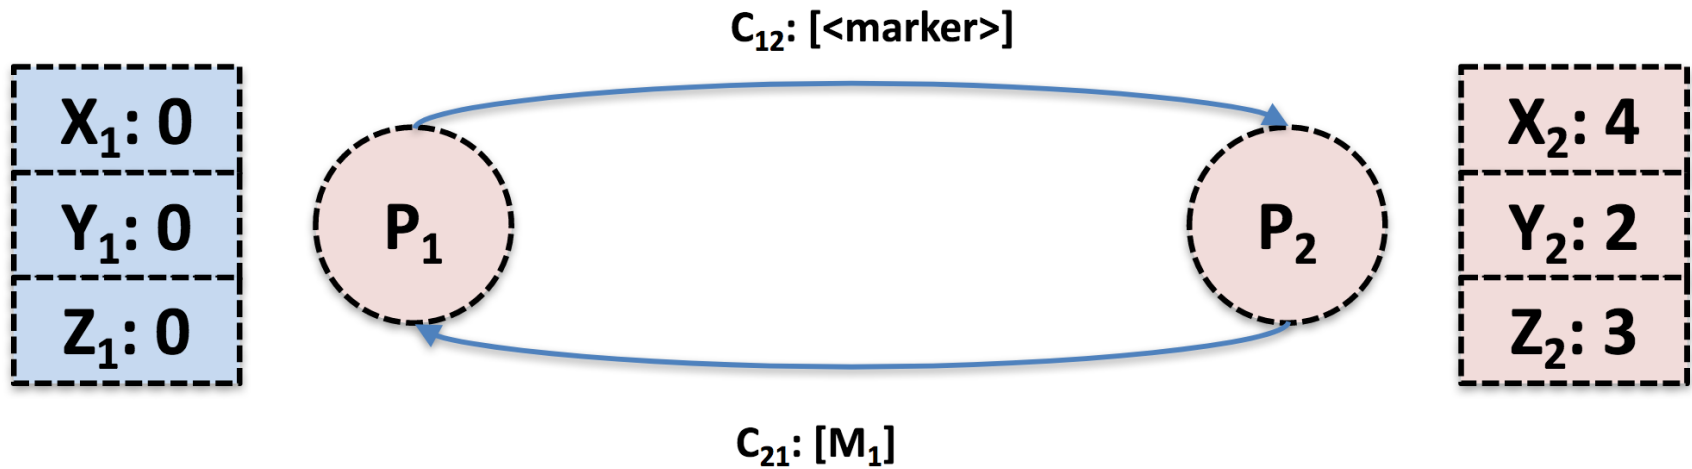
A Close Look of an Example

→ Initially all channels are empty!



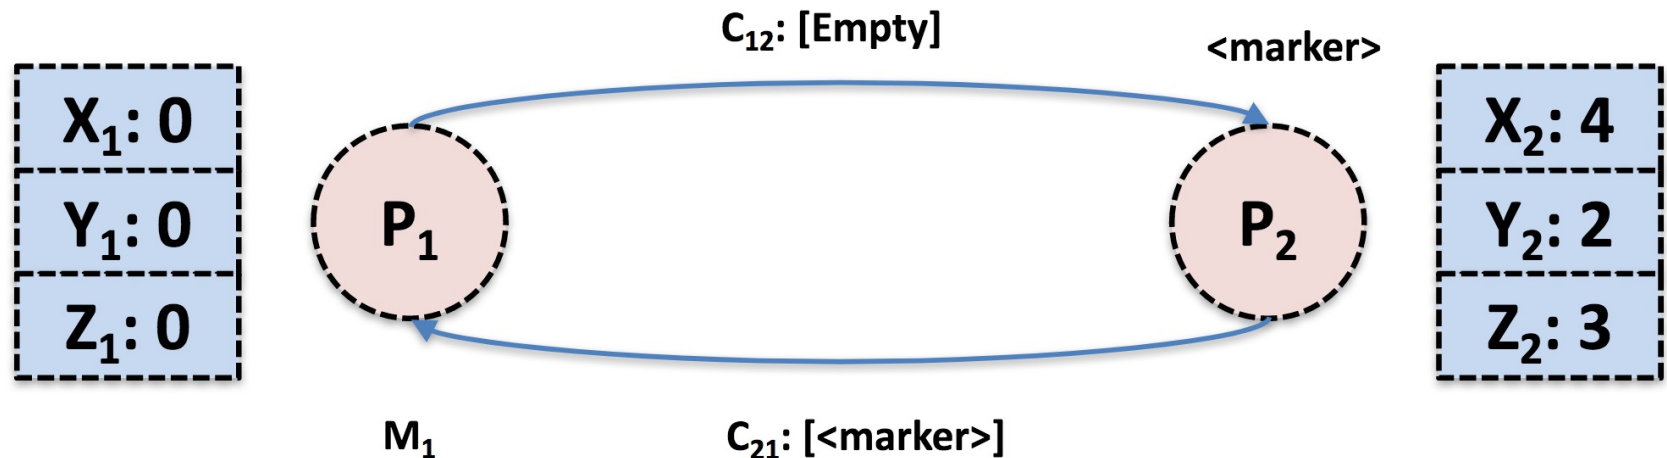
An Example (contd)

- Then, P_1 sends a marker message to P_2 and begins recording all messages on inbound channels
- Meanwhile, P_2 sent a message to P_1



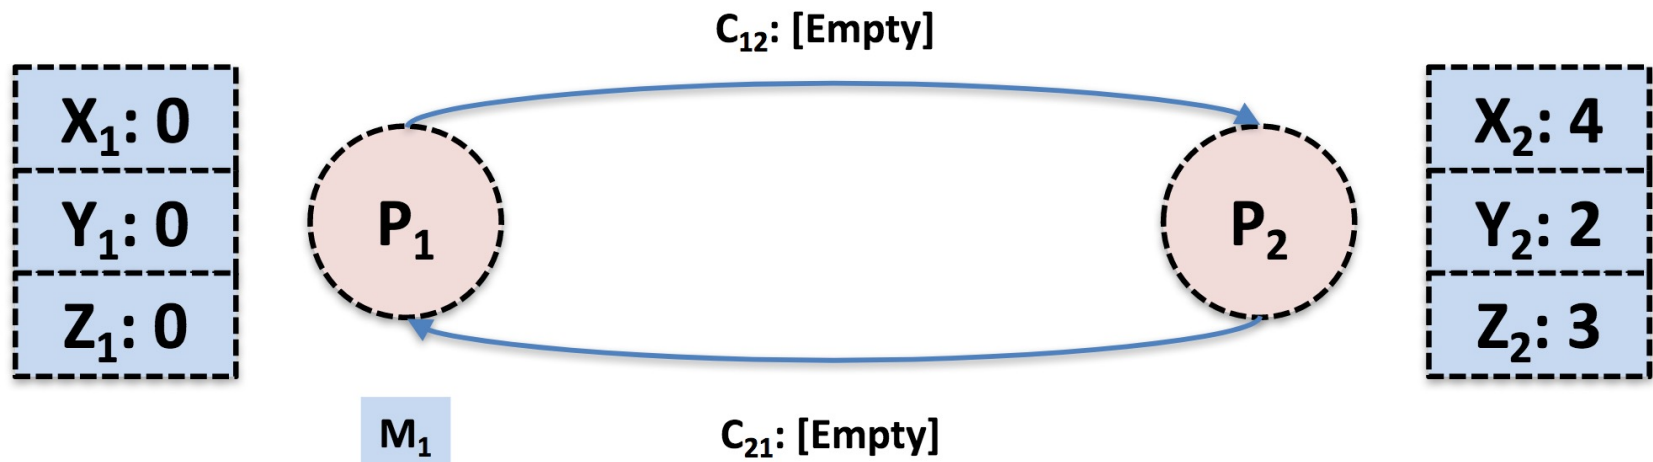
An Example (contd)

- P_2 receives a marker message for the first time, so records its state
- P_2 then sends a marker message to P_1



An Example (contd)

- P_1 has already sent a marker message, so it records all messages it received on inbound channels to the appropriate channel's state



Causal Consistency

- Related to Lamport's clock partial ordering
- An event is pre-snapshot if it occurs before the local snapshot on a process
- Post-snapshot if afterwards
- If an event A happens causally before an event B, and B is a pre-snapshot, then A is too

Summary

→ Global Snapshots

- Causal Precedence Relations
- Global State of a DS

→ Chandy - Lamport's Algorithm

- Global Snapshot Recording Algorithm
- Initiating Snapshot
- Propagating Snapshots
- Terminating the Snapshot Algorithm

→ Causal Consistency

→ Many more to come up ... stay tuned in !!

Penalties



- Every Student is expected to strictly follow a fair Academic Code of Conduct to avoid penalties
- Penalties is heavy for those who involve in:
 - Copy and Pasting the code
 - Plagiarism (copied from your neighbor or friend - in this case, both will get "0" marks for that specific take home assignments)
 - If the candidate is unable to explain his own solution, it would be considered as a "copied case"!!
 - Any other unfair means of completing the assignments

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <https://www.iiits.ac.in/people/regular-faculty/dr-rajendra-prasath/>

(OR)

→ <http://rajendra.2power3.com>

Assistance

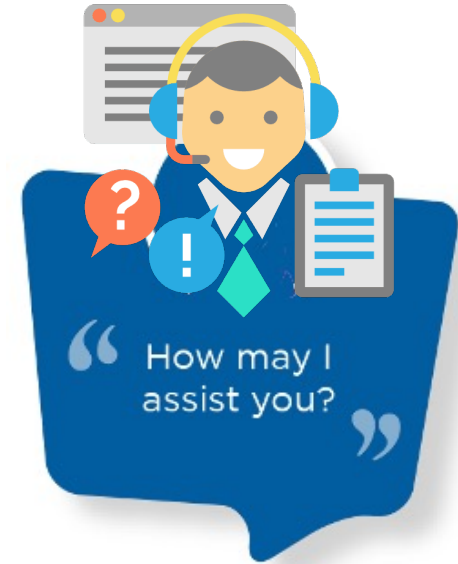
- You may post your questions to me at any time
- You may meet me in person on available time or with an appointment
- You may ask for one-to-one meeting

Best Approach

- You may leave me an email any time
(email is the best way to reach me faster)



Questions It's Your Time



Contact Information

Dr. Rajendra Prasath
IIIT Sri City, Chittoor